

ECE 1508S2: Applied Deep Learning

Chapter 2: Feedforward Neural Networks

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Winter 2024

Various Architectures for NNs

Let's abbreviate the term *Neural Network* from now on with *NN*

Now that we know *what they are* and *how to train them*, we are going through the famous architectures for NNs which are

① *Feedforward NNs* abbreviated as *FNNs*

↳ Some people call them also *Multi-Layer Perceptrons (MLPs)*: you may say that this is a *misnomer* and you are right! Check the [wikipedia page](#)


② *Convolutional NNs* abbreviated as *CNNs*

③ *Recurrent NNs* abbreviated as *RNNs*

In this chapter, we start with *FNNs* which are known to be *vanilla NNs*, i.e., *the most basic architecture we could think for a NN*

FNNs: Architecture

In FNNs, the *inputs flow in one direction*: each layer's output is connected to the next layers, and thus we *do not have any feedback*




Though it is not a must, we usually use *same activation for all neurons in a layer*

Fully-Connected FNNs


We start with the most straightforward FNNs: *fully-connected FNNs*

Fully-Connected FNNs

In a fully-connected FNN, each node is connected to all nodes in the next layer



Fully-Connected FNNs: *Few Definitions*




In this FNN, we have L hidden layers; thus, its depth is $L + 1$

Recall: this network is *Deep* if $L > 1$

Fully-Connected FNNs: *Few Definitions*

Width of a Layer

The width of layer ℓ is the number of neurons in layer ℓ



Some people call the largest width, the width of the network, i.e.,

$$\mathcal{W} = \max_{\ell \in \{1, \dots, L+1\}} \mathcal{W}_\ell$$

Fully-Connected FNNs: Looking as a Model

- + We should look at a fully-connected FNN as a model. Then, what are the *hyperparameters* and *learnable parameters*?
- I am glad that you ask! Let's take a look

Assume that someone tells us that we should use a fully-connected FNN with only ReLU activation. Now, we could say

- To write down the model, we need to know *the number of hidden layers L* and *width of each layer \mathcal{W}_ℓ* : these are the *hyperparameters*
- If we set the L and \mathcal{W}_ℓ , we can specify the *learnable parameters*
 - in *hidden layer 1*, we have \mathcal{W}_1 neurons each having N weights and a bias
 - in *hidden layer 2*, we have \mathcal{W}_2 neurons each having \mathcal{W}_1 weights and a bias
 - \vdots
 - in *output layer*, we have \mathcal{W}_{L+1} neurons each having \mathcal{W}_L weights and a bias

$$\# \text{ model parameters} = (N + 1) \mathcal{W}_1 + \sum_{\ell=1}^L (\mathcal{W}_\ell + 1) \mathcal{W}_{\ell+1}$$

Fully-Connected FNNs: *Forward Pass*

Let us first see how a given data-point propagates through the FNN: we want to write the **outputs** y_1, \dots, y_M when **inputs** x_1, \dots, x_N are given

this is called *forward propagation* through the network

or simply

the *forward pass*

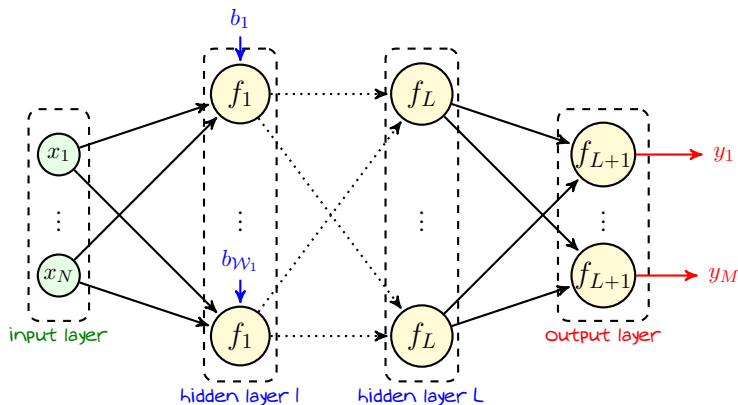
which tracks values *passed* through the NN from the **input** to **output layer**

To present the forward pass compactly, we need to *define some notations* and apply *some modifications* in the network

Fully-Connected FNNs: *Few Definitions*

We can get rid of biases by defining a new constant node in each layer


Let's look at the first layer: we have \mathcal{W}_1 *neurons* and each has *a bias*



Fully-Connected FNNs: *Few Definitions*

We can get rid of biases by defining a new constant node in each layer

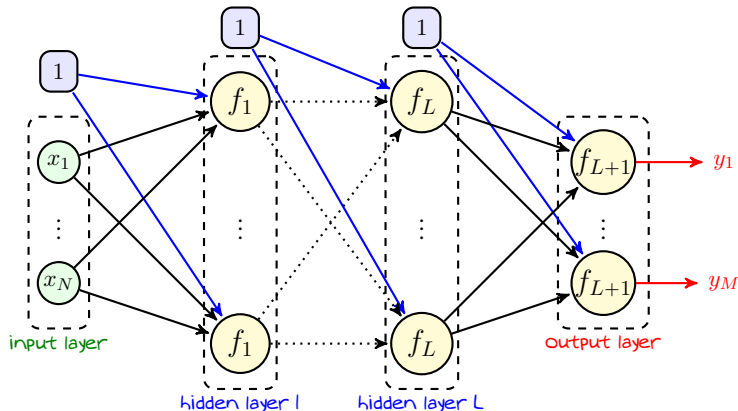
We introduce a constant input and let these biases being the weights of its links



Fully-Connected FNNs: *Few Definitions*

We can get rid of biases by defining a new constant node in each layer


We do the same in all layers: now neurons have no biases



Fully-Connected FNNs: *Few Definitions*

We next give an index to each layer each layer

- **Input layer** is layer 0
- **Hidden layer ℓ** is layer ℓ
- **Output layer** is layer $L + 1$





Fully-Connected FNNs: Few Definitions

We next give an index to each layer each layer

So, our layers are indexed by $\ell \in \{0, \dots, L + 1\}$

- We denote the width of layer ℓ with \mathcal{W}_ℓ
 - ↳ For $\ell \geq 1$ this is exactly the layer width \equiv # of neurons in the layer
 - ↳ For $\ell = 0$ this is the number of inputs, i.e., $\mathcal{W}_0 = N$
- In layer ℓ , we have $\mathcal{W}_\ell + 1$ nodes
 - ↳ \mathcal{W}_ℓ neurons (or inputs if $\ell = 0$)
 - ↳ One constant node that always returns 1




Fully-Connected FNNs: *Few Definitions*

We next index the nodes in each layer

In layer ℓ : we have $\mathcal{W}_\ell + 1$ nodes

↳ One **constant** node \equiv **node** $j = 0$



↳ \mathcal{W}_ℓ **neurons/inputs** \equiv **node** $j = 1, \dots, \mathcal{W}_\ell$



Fully-Connected FNNs: Forward Pass

We next give weights to the links

Weight of the link connecting **node i in layer $\ell - 1$** to **node j in layer ℓ** is denoted by $w_{\ell}[j, i]$



Fully-Connected FNNs: *Forward Pass*

We next give weights to the links

Weight of the link connecting *node i in layer $\ell - 1$* to *node j in layer ℓ* is denoted by $w_\ell[j, i]$

↳ The weights coming out of $i = 0$ are *biases*

$w_\ell[j, 0]$ is the bias of *neuron j in layer ℓ*


↳ There is no link from a node to a *constant node*: $w_\ell[j, i]$ exists for

- ▶ $i = 0, \dots, \mathcal{W}_{\ell-1}$
- ▶ $j = 1, \dots, \mathcal{W}_\ell$

This means that there exists *no such a weight* $w_\ell[0, i]$

Fully-Connected FNNs: *Forward Pass*

We finally specify the output of each node



Fully-Connected FNNs: *Forward Pass*

We finally specify the output of each node

We represent the output of *node j* in *layer ℓ* with $y_\ell[j]$

- ↳ Since $j = 0$ is the *constant node*: $y_\ell[0] = 1$ for $\ell = 0, \dots, L + 1$
- ↳ Since $\ell = 0$ is the *input layer*: $y_0[j] = x_j$ for $j = 1, \dots, N$
- ↳ For *neuron j* in layer ℓ we can write

$$y_\ell[j] = f_\ell(z_\ell[j])$$

where $z_\ell[j]$ is the *output of the affine function* in *neuron j*

Fully-Connected FNNs: *Forward Pass*

We finally specify the output of each node

We represent the output of **node j** in **layer ℓ** with $y_\ell[j]$

- ↳ Since **$j = 0$** is the **constant node**: $y_\ell[0] = 1$ for $\ell = 0, \dots, L + 1$
- ↳ Since $\ell = 0$ is the **input layer**: $y_0[j] = x_j$ for $j = 1, \dots, N$
- ↳ For **neuron j** in layer ℓ we can write


$$y_\ell[j] = f_\ell(z_\ell[j])$$

where $z_\ell[j]$ is the *output of the affine function* in **neuron j**

$$\begin{aligned} z_\ell[j] &= w_\ell[j, 0] + \sum_{i=1}^{W_{\ell-1}} w_\ell[j, i] y_{\ell-1}[i] \\ &= \sum_{i=0}^{W_{\ell-1}} w_\ell[j, i] y_{\ell-1}[i] \end{aligned}$$

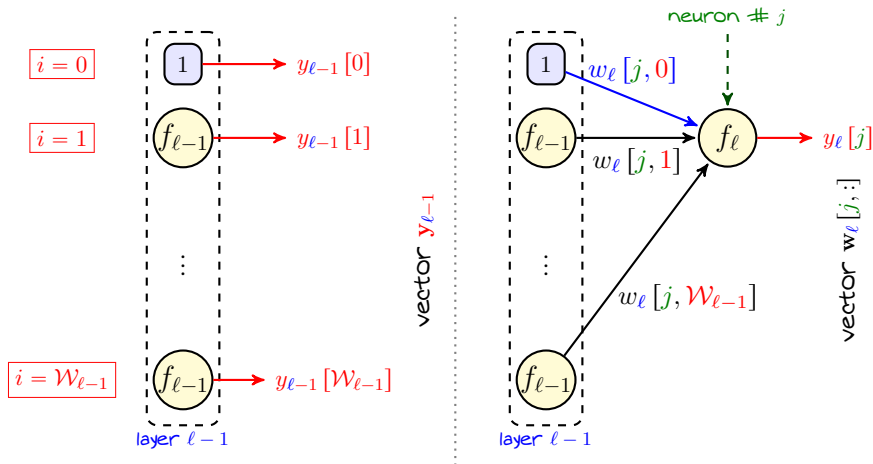
Fully-Connected FNNs: *Forward Pass*

We finally specify the output of each node



Fully-Connected FNNs: Forward Pass

We finally specify the output of each node



Fully-Connected FNNs: *Forward Pass*

We finally specify the output of each node

We can represent $z_j[\ell]$ more compactly via vectorized notation

$$\begin{aligned}
 z_\ell[j] &= \sum_{i=0}^{\mathcal{W}_{\ell-1}} w_\ell[i, j] y_{\ell-1}[i] \\
 &= \underbrace{\begin{bmatrix} w_\ell[j, 0] & w_\ell[j, 1] & \dots & w_\ell[j, \mathcal{W}_{\ell-1}] \end{bmatrix}}_{\mathbf{w}_\ell[j, :]^T} \underbrace{\begin{bmatrix} y_{\ell-1}[0] \\ y_{\ell-1}[1] \\ \vdots \\ y_{\ell-1}[\mathcal{W}_{\ell-1}] \end{bmatrix}}_{\mathbf{y}_{\ell-1}} \\
 &= \mathbf{w}_\ell[j, :]^T \mathbf{y}_{\ell-1}
 \end{aligned}$$

Fully-Connected FNNs: *Forward Pass*

We finally specify the output of each node

We can further extend vectorized notation by defining

$$\mathbf{z}_\ell = \begin{bmatrix} z_\ell [1] \\ \vdots \\ z_\ell [\mathcal{W}_\ell] \end{bmatrix}$$

and thus writing \mathbf{y}_ℓ as


$$\mathbf{y}_\ell = f_\ell(\mathbf{z}_\ell)$$

where $f_\ell(\cdot)$ is applied entry-wise

and don't forget to add the dummy input 1, i.e.,

$$\mathbf{y}_\ell \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_\ell \end{bmatrix}$$

Fully-Connected FNNs: Forward Pass



```

1: Initiate the output of the first layer as  $\mathbf{y}_0 = \mathbf{x}$ 
2: for  $\ell = 0, \dots, L$  do
3:   for  $j = 1, \dots, \mathcal{W}_\ell$  do
4:     Add  $\mathbf{y}_\ell[0] = 1$  and set  $\mathbf{z}_{\ell+1}[j] = \mathbf{w}_{\ell+1}[j, :]^T \mathbf{y}_\ell$  # affine function
5:   end for
6:   Compute  $\mathbf{y}_{\ell+1} = f_{\ell+1}(\mathbf{z}_{\ell+1})$  # activation
7: end for
8: for  $\ell = 1, \dots, L + 1$  do
9:   Return  $\mathbf{y}_\ell$  and  $\mathbf{z}_\ell$ 
10: end for

```


Fully-Connected FNNs: *Forward Pass*

We can present everything even more compactly: *let's write down \mathbf{z}_ℓ*

$$\mathbf{z}_\ell = \begin{bmatrix} z_\ell[1] \\ \vdots \\ z_\ell[\mathcal{W}_\ell] \end{bmatrix} = \begin{bmatrix} \mathbf{w}_\ell[1, :]^\top \mathbf{y}_{\ell-1} \\ \vdots \\ \mathbf{w}_\ell[\mathcal{W}_\ell, :]^\top \mathbf{y}_{\ell-1} \end{bmatrix} = \begin{bmatrix} \mathbf{w}_\ell[1, :]^\top \\ \vdots \\ \mathbf{w}_\ell[\mathcal{W}_\ell, :]^\top \end{bmatrix} \mathbf{y}_{\ell-1}$$

Now, we can define the matrix \mathbf{W}_ℓ as

$$\mathbf{W}_\ell = \begin{bmatrix} \mathbf{w}_\ell[1, :]^\top \\ \vdots \\ \mathbf{w}_\ell[\mathcal{W}_\ell, :]^\top \end{bmatrix} = \begin{bmatrix} w_\ell[1, 0] & \dots & w_\ell[1, \mathcal{W}_{\ell-1}] \\ \vdots & & \vdots \\ w_\ell[\mathcal{W}_\ell, 0] & \dots & w_\ell[\mathcal{W}_\ell, \mathcal{W}_{\ell-1}] \end{bmatrix}$$

This matrix collects all learning parameters of layer ℓ

Note that \mathbf{W}_ℓ has \mathcal{W}_ℓ rows and $\mathcal{W}_{\ell-1} + 1$ columns

Forward Propagation: *Pseudo Code*


So, we can compactly present the forward propagation algorithm as follow

```
ForwardProp():  
1: Initiate with  $\mathbf{y}_0 = \mathbf{x}$   
2: for  $\ell = 0, \dots, L$  do  
3:   Add  $\mathbf{y}_\ell[0] = 1$  and determine  $\mathbf{z}_{\ell+1} = \mathbf{W}_{\ell+1}\mathbf{y}_\ell$  # forward affine  
4:   Determine  $\mathbf{y}_{\ell+1} = f_{\ell+1}(\mathbf{z}_{\ell+1})$  # forward activation  
5: end for  
6: for  $\ell = 1, \dots, L + 1$  do  
7:   Return  $\mathbf{y}_\ell$  and  $\mathbf{z}_\ell$   
8: end for
```

After getting data-point x , we convert it to \mathbf{x} by adding an entry 1 at its index zero. We then pass it through a **linear layer** whose weights are **learnable** and a **nonlinear** transform that is specified by **activation**. The output of this layer passes forward to the next layer till we get to the output.


Forward Propagation: *Compact Diagram*

Inspired by forward propagation, we can represent the FNN




Forward Propagation: Compact Diagram


Inspired by forward propagation, we can represent the FNN




by the following compact diagram



Forward Propagation: Compact Diagram



Here, we compactly represent **layer ℓ** as



- The link W_ℓ represents the affine function of layer ℓ
- The block f_ℓ represents the **activation** of layer ℓ and **adding** $y_\ell[0] = 1$
 - ↳ The input of this block can be considered z_ℓ
 - ↳ The output of this block can be considered y_ℓ

This compact diagram will come in handy when we derive backpropagation!

Fully-Connected FNNs: *Training*

- + *So, what do we do with forward pass?*
- Say the **FNN is fixed** and all **the weights and biases** are given. Forward pass determines the **label** of a given **data-point x** .
- + *Well! But we need to train the network! Right?!*
- Yes! We define the **loss** and then **train** it via **gradient descent**
- + *Then, we need to determine the **gradient**! It sounds complicated!*
- Well! there is an efficient algorithm for that called **backpropagation**

Let's see what **backpropagation** is!

Fully-Connected FNNs: *Training*

Let's recall how we train the network: in our FNN, we considered M outputs; thus, we could assume that the dataset is of the form

$$\mathbb{D} = \{(\mathbf{x}_b, \mathbf{v}_b) \text{ for } b = 1, \dots, B\}$$

Here, we have denoted the true labels by $\mathbf{v}_b = [v_{b,1}, \dots, v_{b,M}]^T$ to avoid confusion with the FNN's outputs.

Fully-Connected FNNs: *Training*

Let's recall how we train the network: in our FNN, we considered M outputs; thus, we could assume that the dataset is of the form

$$\mathbb{D} = \{(\mathbf{x}_b, \mathbf{v}_b) \text{ for } b = 1, \dots, B\}$$

Here, we have denoted the **true labels** by $\mathbf{v}_b = [v_{b,1}, \dots, v_{b,M}]^T$ to avoid confusion with the FNN's outputs. Now, let's denote the forward pass by $\text{PassF}(\mathbf{x}_b | \mathbf{w})$ with \mathbf{w} is a vector collecting $\{\mathbf{W}_\ell\}$ for $\ell = 1, \dots, L + 1$

Given **data-point** \mathbf{x}_b , by forward pass we get output $\text{PassF}(\mathbf{x}_b | \mathbf{w})$ from FNN with weights \mathbf{w} . This output is desired to be the **true label** \mathbf{v}_b

How do we do the training?

$$\mathbf{w}^\star = \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\text{PassF}(\mathbf{x}_b | \mathbf{w}), \mathbf{v}_b) \quad (\text{Training})$$

Fully-Connected FNNs: *Training*

Let's recall how we train the network

$$\mathbf{w}^{\star} = \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\operatorname{PassF}(\mathbf{x}_b | \mathbf{w}), \mathbf{v}_b) \quad (\text{Training})$$

Also, we recall the *gradient descent* algorithm

- 1: Initiate at some $\mathbf{w}^{(0)} \in \mathbb{R}^D$ and deviation $\Delta = +\infty$
- 2: Choose some small ϵ and η , and set $t = 1$
- 3: **while** $\Delta > \epsilon$ **do**
- 4: Update weights as $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla \hat{R}(\mathbf{w}^{(t-1)})$
- 5: Update the deviation $\Delta = |\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^{(t-1)})|$
- 6: **end while**

In this algorithm, the main challenge is to calculate $\nabla \hat{R}(\mathbf{w}^{(t-1)})$

Fully-Connected FNNs: Training

The main challenge is to calculate $\nabla \hat{R}(\mathbf{w})$

First, let's see what are the entries of \mathbf{w} : \mathbf{w} contains all weights and biases. Following our notations, we can say

$$\mathbf{w} = \begin{bmatrix} \mathbf{w}_1 [1, :] \\ \vdots \\ \mathbf{w}_1 [\mathcal{W}_1, :] \\ \vdots \\ \mathbf{w}_{L+1} [1, :] \\ \vdots \\ \mathbf{w}_{L+1} [\mathcal{W}_{L+1}, :] \end{bmatrix} \begin{matrix} \text{layer } \ell = 1 \\ \\ \\ \\ \text{layer } \ell = L + 1 \end{matrix} \longrightarrow \mathbf{w}_\ell [j, :] = \begin{bmatrix} w_\ell [j, 0] \\ w_\ell [j, 1] \\ \vdots \\ w_\ell [j, \mathcal{W}_{\ell-1}] \end{bmatrix}$$

So, the entries of \mathbf{w} are $w_\ell [j, i]$ for different choices of i, j and ℓ

Fully-Connected FNNs: *Training*

The main challenge is to calculate $\nabla \hat{R}(\mathbf{w})$

Let's try to open up the gradient: we need partial derivatives of $\hat{R}(\mathbf{w})$ with respect to $w_\ell [j, i]$ for $i = 0, \dots, \mathcal{W}_{\ell-1}$ and $j = 1, \dots, \mathcal{W}_\ell$ as ℓ runs over $\ell = 1, \dots, L + 1$

$$\begin{aligned} \frac{\partial}{\partial w_\ell [j, i]} \hat{R}(\mathbf{w}) &= \frac{\partial}{\partial w_\ell [j, i]} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\text{PassF}(\mathbf{x}_b | \mathbf{w}), \mathbf{v}_b) \\ &= \frac{1}{B} \sum_{b=1}^B \boxed{\frac{\partial}{\partial w_{i,j} [\ell]} \mathcal{L}(\text{PassF}(\mathbf{x}_b | \mathbf{w}), \mathbf{v}_b)} \end{aligned}$$

So, it's enough to develop an algorithm that determined the partial derivative for a single *data-point*. The partial derivative is then the average of these point-wise derivatives.

Fully-Connected FNNs: *Training*

Let's make an agreement: we consider a *data-point* \mathbf{x} with *label* \mathbf{v} and write

$$\frac{\partial}{\partial w_{\ell} [j, i]} \mathcal{L} (\text{PassF} (\mathbf{x} | \mathbf{w}), \mathbf{v}) = \frac{\partial}{\partial w_{\ell} [j, i]} \mathcal{L} (\mathbf{y}, \mathbf{v})$$

while keeping in mind that \mathbf{y} is a function of \mathbf{w}

To determine the partial derivatives, we note that

\mathbf{y} is a nested function of \mathbf{w}

so, we can determine the derivative via *chain rule*. Let's recall the *chain rule* and see how we can *apply it on a graph*

Review: Chain Rule

Assume $z = g(x)$ and $y = f(z)$: y is a nested function of x , as we can write

$$y = f(g(x))$$

Intuitively, we can say: if at point x we move with tiny step dx , z varies as

$$dz = g'(x)dx$$

This variation also varies y : moving from $z = g(x)$ with tiny step dz leads to

$$dy = f'(z)dz$$

Review: Chain Rule

Assume $z = g(x)$ and $y = f(z)$: y is a nested function of x , as we can write

$$y = f(g(x))$$

Intuitively, we can say: if at point x we move with tiny step dx , z varies as

$$dz = \dot{g}(x)dx$$

This variation also varies y : moving from $z = g(x)$ with tiny step dz leads to

$$dy = \dot{f}(z)dz$$

So, we have

$$dy = \dot{f}(z)\dot{g}(x)dx$$

Review: Chain Rule

We have concluded that by moving x with dx , we get

$$dy = \dot{f}(z)\dot{g}(x)dx$$

On the other hand, we know that

$$dy = \frac{d}{dx}f(g(x))dx$$

This concludes the *chain rule*

Review: Chain Rule

We have concluded that by moving x with dx , we get

$$dy = f'(z)g'(x)dx$$

On the other hand, we know that

$$dy = \frac{d}{dx} f(g(x))dx$$

This concludes the *chain rule*

Chain Rule: Scalar Form

The derivative of nested function $y = f(g(x))$ with respect to x is given by

$$\frac{dy}{dx} = \frac{d}{dx} f(g(x)) = f'(z)g'(x) = \frac{dy}{dz} \frac{dz}{dx}$$

Computation Graph

We can extend this idea to deeper nested functions:

Let $z_1 = g_1(x)$ and $z_{n+1} = g_{n+1}(z_n)$ for $n = 1, \dots, N$; then, derivative of $y = f(z_N)$ with respect to x is given by

$$\frac{dy}{dz_j} = \frac{dy}{dz_N} \left(\prod_{n=1}^{N-1} \frac{dz_{n+1}}{dz_n} \right) \frac{dz_1}{dx} = \dot{f}(z_N) \left(\prod_{n=1}^{N-1} \dot{g}_{n+1}(z_n) \right) \dot{g}_1(x)$$


Computation Graph

We can extend this idea to deeper nested functions:

Let $z_1 = g_1(x)$ and $z_{n+1} = g_{n+1}(z_n)$ for $n = 1, \dots, N$; then, derivative of $y = f(z_N)$ with respect to x is given by

$$\frac{dy}{dz_j} = \frac{dy}{dz_N} \left(\prod_{n=1}^{N-1} \frac{dz_{n+1}}{dz_n} \right) \frac{dz_1}{dx} = \dot{f}(z_N) \left(\prod_{n=1}^{N-1} \dot{g}_{n+1}(z_n) \right) \dot{g}_1(x)$$


We can represent the chain rule, using a *computation graph*: for the deep nested function given above, the computation graph is given by



In this graph, we start from x and pass forward to $z_1 \rightarrow z_2 \rightarrow \dots$ until we get to y . In each pass, we determine next variable via the function on the link

Computation Graph

The derivative of y with respect to any variable on this graph is determined by a *backward pass* from y towards the variable




Let's start from the last node

$$\frac{dy}{dz_N} = \dot{f}(z_N)$$

Computation Graph

The derivative of y with respect to any variable on this graph is determined by a *backward pass* from y towards the variable




Let's start from the last node

$$\frac{dy}{dz_N} = \dot{f}(z_N)$$

$$\frac{dy}{dz_{N-1}} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} = \dot{f}(z_N) \dot{g}_N(z_{N-1})$$

Computation Graph

The derivative of y with respect to any variable on this graph is determined by a *backward pass* from y towards the variable



Let's start from the last node


$$\frac{dy}{dz_N} = \dot{f}(z_N)$$

$$\frac{dy}{dz_{N-1}} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} = \dot{f}(z_N) \dot{g}_N(z_{N-1})$$

$$\frac{dy}{dz_1} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} \dots \frac{dz_2}{dz_1} = \dot{f}(z_N) \dot{g}_N(z_{N-1}) \dots \dot{g}_2(z_2)$$

Computation Graph

The derivative of y with respect to any variable on this graph is determined by a *backward pass* from y towards the variable



Let's start from the last node

$$\frac{dy}{dz_N} = \dot{f}(z_N)$$

$$\frac{dy}{dz_{N-1}} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} = \dot{f}(z_N) \dot{g}_N(z_{N-1})$$

$$\frac{dy}{dz_1} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} \cdots \frac{dz_2}{dz_1} = \dot{f}(z_N) \dot{g}_N(z_{N-1}) \cdots \dot{g}_2(z_2)$$

$$\frac{dy}{dx} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} \cdots \frac{dz_2}{dz_1} \frac{dz_1}{dx} = \dot{f}(z_N) \dot{g}_N(z_{N-1}) \cdots \dot{g}_2(z_2) \dot{g}_1(x)$$

Computation Graph: *Example*

Example: y is a nested function of x through the following chain of functions:

$$z_1 = 2x \quad z_2 = z_1^2 \quad z_3 = e^{z_2} \quad y = \log z_3$$

Determine the derivative of y with respect to x at $x = 0.5$.


Computation Graph: Example

Example: y is a nested function of x through the following chain of functions:

$$z_1 = 2x \quad z_2 = z_1^2 \quad z_3 = e^{z_2} \quad y = \log z_3$$


Determine the derivative of y with respect to x at $x = 0.5$.

Let's first plot the computation graph




During the *forward pass* we get

$$z_1 = 2 \times 0.5 = 1 \rightarrow z_2 = 1^2 = 1 \rightarrow z_3 = e^1 = e \rightarrow y = \log e = 1$$




Computation Graph: *Example*

Now, we *pass backward* to determine the derivative




Computation Graph: *Example*

Now, we *pass backward* to determine the derivative




Let's first enter the values into the backward links



We now *navigate backward* to each variable that we want to determine the derivative y with respect to it


Computation Graph: *Example*



The derivatives are easily determined *recursively*

$$\frac{dy}{dz_3} = \frac{1}{e}$$

Computation Graph: *Example*




The derivatives are easily determined *recursively*

$$\frac{dy}{dz_3} = \frac{1}{e}$$

$$\frac{dy}{dz_2} = \frac{dy}{dz_3} \frac{dz_3}{dz_2} = \frac{e}{e} = 1$$


Computation Graph: *Example*



The derivatives are easily determined *recursively*

$$\begin{aligned}\frac{dy}{dz_3} &= \frac{1}{e} \\ \frac{dy}{dz_2} &= \frac{dy}{dz_3} \frac{dz_3}{dz_2} = \frac{e}{e} = 1 \\ \frac{dy}{dz_1} &= \frac{dy}{dz_2} \frac{dz_2}{dz_1} = 1 \times 2 = 2\end{aligned}$$


Computation Graph: *Example*



The derivatives are easily determined *recursively*

$$\begin{aligned}\frac{dy}{dz_3} &= \frac{1}{e} \\ \frac{dy}{dz_2} &= \frac{dy}{dz_3} \frac{dz_3}{dz_2} = \frac{e}{e} = 1 \\ \frac{dy}{dz_1} &= \frac{dy}{dz_2} \frac{dz_2}{dz_1} = 1 \times 2 = 2 \\ \frac{dy}{dx} &= \frac{dy}{dz_1} \frac{dz_1}{dx} = 2 \times 2 = 4\end{aligned}$$

Computation Graph



In the example, we had to *first determine the value of each variable*, in order to be able to determine the *values of the backward links*.

This is an important fact that we should remember

Backward pass is only possible if we have already taken the forward pass

Review: Chain Rule

The nested function can be a multivariate: *assume for* $n = 1, \dots, N$

$$z_n = g_n(x)$$

and let the nested function be

$$y = f(z_1, \dots, z_N)$$

Let's follow the same logic: *starting from point* x , *we move with tiny step* dx .
This leads to

$$dz_n = \dot{g}_n(x) dx$$

These variations lead to variation dy *in the nested function*

$$dy = \nabla f(\mathbf{z})^\top d\mathbf{z} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} dz_n = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) dz_n$$

Review: Chain Rule

We can hence write

$$dy = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) dz_n = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x) dx$$

This concludes the *vector form* of the chain rule

Review: Chain Rule

We can hence write

$$dy = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) dz_n = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x) dx$$

This concludes the *vector form* of the chain rule


Chain Rule: Vector Form

Let $\mathbf{z} = [z_1, \dots, z_N]^T$ and $z_n = g_n(x)$. The derivative of nested function $y = f(\mathbf{z})$ with respect to x is given by

$$\frac{dy}{dx} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \frac{dz_n}{dx} = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x)$$

Computation Graph


We can again represent the vector form via its *computation graph*



In this graph, we start from x and pass forward to $\mathbf{z} = [z_1, z_2, \dots, z_N]$. We then pass forward \mathbf{z} to y . In each pass, we determine next variable via the function on the link.

Computation Graph

The derivative with respect to any node is then given by *backward pass* towards the node on the computation graph




We add all backward passes towards x to determine the derivative

$$\frac{dy}{dx} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \frac{dz_n}{dx} = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x)$$

Computation Graph: *Single Neuron*


Let's now plot the *computation graph* of a single neuron and determine the gradient of the loss by *backward pass*



After passing the data-point \mathbf{x} through the neuron, we get \mathbf{y} and we calculate the loss for the *true label* v as $\mathcal{L}(\mathbf{y}, v)$

Computation Graph: *Single Neuron*

The *computation graph* is hence given by




Here, the computation nodes are the *weights and bias of the neuron*

once we fix them, we can pass forward and get to the loss \hat{R}

Computation Graph: Single Neuron

Once passed forward, we can move backward to determine the derivatives




For a particular weight w_n , we can write (we drop arguments whenever clear)

$$\frac{\partial \hat{R}}{\partial w_n} = \frac{d\hat{R}}{dy} \frac{dy}{dz} \frac{\partial z}{\partial w_n} = \dot{\mathcal{L}} f' x_n$$

We can extend it to a deeper and wider network

Computation Graph: *Multivariate Form*

Let's get back to the following *computation graph*



We define *vector-valued functions*, and show the graph *compactly*: let's define

$$\mathbf{g}(x) = \begin{bmatrix} g_1(x) \\ \vdots \\ g_N(x) \end{bmatrix}$$

Computation Graph: *Multivariate Form*

Function $\mathbf{g}(\cdot)$ gets x as the input and returns all z_n 's in a **vector** \mathbf{z} , i.e.,


$$\mathbf{g}(x) = \begin{bmatrix} g_1(x) \\ \vdots \\ g_N(x) \end{bmatrix} = \begin{bmatrix} z_1 \\ \vdots \\ z_N \end{bmatrix} = \mathbf{z}$$

Computation Graph: *Multivariate Form*

Function $g(\cdot)$ gets x as the input and returns all z_n 's in a **vector** \mathbf{z} , i.e.,

$$\mathbf{g}(x) = \begin{bmatrix} g_1(x) \\ \vdots \\ g_N(x) \end{bmatrix} = \begin{bmatrix} z_1 \\ \vdots \\ z_N \end{bmatrix} = \mathbf{z}$$

We now use this vectorized notation to simplify the *computation graph* as



The **forward pass** on this graph is exactly the same: we give x to the **vectorized function** $g(\cdot)$ to get \mathbf{z} which is then **passed forward** to $f(\cdot)$ to get y


How does the **backward pass** look like then?

Computation Graph: *Multivariate Form*

We could define the derivative $\dot{\mathbf{g}}(\cdot)$ as the vector of derivatives $\dot{g}(\cdot)$

$$\dot{\mathbf{g}}(x) = \begin{bmatrix} \dot{g}_1(x) \\ \vdots \\ \dot{g}_N(x) \end{bmatrix} = \begin{bmatrix} \frac{dz_1}{dx} \\ \vdots \\ \frac{dz_N}{dx} \end{bmatrix} = \frac{d\mathbf{z}}{dx}$$

Let's show this **vectorized derivative** and **gradient** of f on the backward links




Computation Graph: *Multivariate Form*

We could define the derivative $\dot{\mathbf{g}}(\cdot)$ as the vector of derivatives $\dot{g}(\cdot)$

$$\dot{\mathbf{g}}(x) = \begin{bmatrix} \dot{g}_1(x) \\ \vdots \\ \dot{g}_N(x) \end{bmatrix} = \begin{bmatrix} \frac{dz_1}{dx} \\ \vdots \\ \frac{dz_N}{dx} \end{bmatrix} = \frac{d\mathbf{z}}{dx}$$

Let's show this **vectorized derivative** and **gradient** of f on the backward links



Well, we can **pass backward** as follows


$$\frac{dy}{dx} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \frac{dz_n}{dx} = \nabla f(\mathbf{z})^T \dot{\mathbf{g}}(x)$$

Computation Graph: *Multivariate Form*

- + What can we conclude then?
- We can sketch the computation graph very compactly using *vectorized derivatives* and *gradients*
- + Does it mean that we should then *pass backward* exactly the same as in a computation graph with *scalar* variables and derivatives?
- Pretty much Yes! Only one delicate detail: we should *know how to multiply those gradients and vectorized derivatives!*

Computation Graph: Multivariate Form

- + What can we conclude then?
- We can sketch the computation graph very compactly using **vectorized derivatives** and **gradients**
- + Does it mean that we should then **pass backward** exactly the same as in a computation graph with **scalar** variables and derivatives?
- Pretty much Yes! Only one delicate detail: we should **know how to multiply those gradients and vectorized derivatives!**




In our example, we determined *the inner product*

$$\frac{dy}{dx} = \nabla f^T \dot{g}$$

Computation Graph: *Multivariate Form*

- + How do we know *which type of product* we should use?
- Well! If you were in doubt, we could always do it by *expanding in terms of entries*; however, we are going to practice *all key functions* that appear in NN computation graphs!

Before we start with all key functions, let's get back to a *single neuron*



Computation Graph: *Multivariate Form*


Let's define, as we did earlier, the following vectors

$$\mathbf{x} = \begin{bmatrix} x_0 = 1 \\ x_1 \\ \vdots \\ x_N \end{bmatrix} \quad \text{and} \quad \mathbf{w} = \begin{bmatrix} w_0 = b \\ w_1 \\ \vdots \\ w_N \end{bmatrix}$$

Recall that **output** of the **neuron** is determined as $y = f(z)$ for


$$z = \mathbf{x}^T \mathbf{w}$$

So, we can show the **computation graph** compactly as



Computation Graph: *Multivariate Form*


Let's look at each link carefully: we pass *backward*, so we start with *last link*



- \hat{R} is a scalar function of scalar y , i.e., $\hat{R} = \mathcal{L}(y, v)$
 ↳ the backward link contains the *scalar derivative* $\dot{\mathcal{L}}$

Computation Graph: *Multivariate Form*


Let's look at each link carefully: we pass *backward*, so we start with *last* link



- \hat{R} is a scalar function of scalar y , i.e., $\hat{R} = \mathcal{L}(y, v)$
 ↳ the backward link contains the *scalar derivative* $\dot{\hat{R}}$
- y is a scalar function of scalar z , i.e., $y = f(z)$
 ↳ the backward link contains the *scalar derivative* \dot{f}

Computation Graph: *Multivariate Form*


Let's look at each link carefully: we pass *backward*, so we start with *last* link



- \hat{R} is a scalar function of scalar y , i.e., $\hat{R} = \mathcal{L}(y, v)$
 ↳ the backward link contains the *scalar derivative* $\dot{\hat{R}}$
- y is a scalar function of scalar z , i.e., $y = f(z)$
 ↳ the backward link contains the *scalar derivative* \dot{f}
- z is a scalar function of vector \mathbf{w} , i.e., $z = \mathbf{x}^T \mathbf{w}$
 ↳ the backward link contains the *gradient* ∇z


Computation Graph: *Multivariate Form*

Let's look at each link carefully: we pass *backward*, so we start with *last* link




- \hat{R} is a scalar function of scalar y , i.e., $\hat{R} = \mathcal{L}(y, v)$
 ↳ the backward link contains the *scalar derivative* $\dot{\mathcal{L}}$
- y is a scalar function of scalar z , i.e., $y = f(z)$
 ↳ the backward link contains the *scalar derivative* \dot{f}
- z is a scalar function of vector \mathbf{w} , i.e., $z = \mathbf{x}^T \mathbf{w}$
 ↳ the backward link contains the *gradient* ∇z

So, the graph with the *backward* links looks like




Computation Graph: *Multivariate Form*



We are almost complete; only we need to *calculate* ∇z


Computation Graph: *Multivariate Form*



We are almost complete; only we need to *calculate* ∇z

$$z = w_0 + w_1 x_1 + \dots + w_N x_N \rightsquigarrow \nabla z = \begin{bmatrix} \partial z / \partial w_0 \\ \partial z / \partial w_1 \\ \vdots \\ \partial z / \partial w_N \end{bmatrix} = \begin{bmatrix} 1 = x_0 \\ x_1 \\ \vdots \\ x_N \end{bmatrix} = \mathbf{x}$$


Computation Graph: Multivariate Form



We are almost complete; only we need to *calculate* ∇z


$$z = w_0 + w_1 x_1 + \dots + w_N x_N \rightsquigarrow \nabla z = \begin{bmatrix} \partial z / \partial w_0 \\ \partial z / \partial w_1 \\ \vdots \\ \partial z / \partial w_N \end{bmatrix} = \begin{bmatrix} 1 = x_0 \\ x_1 \\ \vdots \\ x_N \end{bmatrix} = \mathbf{x}$$

So, we are complete! Here is the *vectorized computation graph* of the *neuron*




Computation Graph: *Multivariate Form*

Now, how do we pass *backward* on this graph?



Computation Graph: Multivariate Form

Now, how do we pass *backward* on this graph?




We arrived at y at the end of forward pass: at this point, we can determine

$$\frac{d\hat{R}}{dy} = \dot{\mathcal{L}}(y, v) = \dot{\mathcal{L}}$$

and we are at the *computing node* y . We then pass *backward* $\dot{\mathcal{L}}$.

Computation Graph: Multivariate Form

Now, how do we pass **backward** on this graph?



We arrived at y at the **end of forward pass**: at this point, we can determine

$$\frac{d\hat{R}}{dy} = \dot{\mathcal{L}}(y, v) = \dot{\mathcal{L}}$$


and we are at the **computing node** y . We then pass **backward** $\dot{\mathcal{L}}$. At node z , we can compute $\dot{f}(z)$, and use what we received from y to compute

$$\frac{d\hat{R}}{dz} = \frac{d\mathcal{L}}{dy} \frac{dy}{dz} = \dot{f} \dot{\mathcal{L}}$$

and pass it **backward**

Computation Graph: Multivariate Form

Now, how do we pass *backward* on this graph?



Arriving at w , we can determine $\nabla z = \mathbf{x}$ and use what we received from z to compute *what we want*

$$\nabla \hat{R} \text{ w.r.t. } \mathbf{w} \equiv \nabla_{\mathbf{w}} \hat{R} = \frac{d\hat{R}}{dz} \nabla z = \dot{f} \dot{\mathcal{L}} \mathbf{x}$$

- + Well! That seems easier!
- Right! Let's now try some **important cases**

Backpropagation: *Local Operations*

Let's consider a general problem: an *objective scalar* \hat{R} is a function of *K -dimensional vector* $\mathbf{y} \in \mathbb{R}^K$. Clearly in this case, we have a *gradient*


$$\nabla_{\mathbf{y}} \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial y_1 \\ \vdots \\ \partial \hat{R} / \partial y_K \end{bmatrix}$$

Backpropagation: Local Operations

Let's consider a general problem: an *objective scalar* \hat{R} is a function of *K -dimensional vector* $\mathbf{y} \in \mathbb{R}^K$. Clearly in this case, we have a *gradient*

$$\nabla_{\mathbf{y}} \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial y_1 \\ \vdots \\ \partial \hat{R} / \partial y_K \end{bmatrix}$$

Assume that we know this gradient. The *vector* \mathbf{y} is also function of *another variable*. We want to compute *gradient of* \hat{R} with respect to this *other variable*




We now consider different cases for the *other variable* and its link to \mathbf{y}

Backpropagation: *Local Operation 1*

Entry-wise Functional Operation


$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{z} \in \mathbb{R}^K$ as $\mathbf{y} = f(\mathbf{z})$ with $f(\cdot)$ operating entry-wise



Backpropagation: *Local Operation 1*

Entry-wise Functional Operation

$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{z} \in \mathbb{R}^K$ as $\mathbf{y} = f(\mathbf{z})$ with $f(\cdot)$ operating entry-wise



For this case, we note that y_k is only a function of z_k ; thus we have

$$\frac{\partial \hat{R}}{\partial z_k} = \frac{\partial \hat{R}}{\partial y_k} \frac{\partial y_k}{\partial z_k} = \frac{\partial \hat{R}}{\partial y_k} f'(z_k)$$

So, we can use *entry-wise product* \odot to get from $\nabla_{\mathbf{y}} \hat{R}$ to $\nabla_{\mathbf{z}} \hat{R}$


$$\nabla_{\mathbf{z}} \hat{R} = \nabla_{\mathbf{y}} \hat{R} \odot \dot{f}(\mathbf{z})$$

Backpropagation: Local Operation 1

Reminder: Entry-wise product of two vectors of the same size is

$$\mathbf{z} \odot \mathbf{y} \begin{bmatrix} z_1 \\ \vdots \\ z_K \end{bmatrix} \odot \begin{bmatrix} y_1 \\ \vdots \\ y_K \end{bmatrix} = \begin{bmatrix} y_1 z_1 \\ \vdots \\ y_K z_K \end{bmatrix}$$

So, we can compactly perform this local operation as follows




with the backward step

$$\nabla_{\mathbf{z}} \hat{R} = \nabla_{\mathbf{y}} \hat{R} \odot \dot{f}(\mathbf{z})$$

Backpropagation: *Local Operation 2*

Linear Vector-to-Vector Operation


$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{z} \in \mathbb{R}^N$ as $\mathbf{y} = \mathbf{A}\mathbf{z}$ with $\mathbf{A} \in \mathbb{R}^{K \times N}$



Backpropagation: Local Operation 2

Linear Vector-to-Vector Operation

$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{z} \in \mathbb{R}^N$ as $\mathbf{y} = \mathbf{A}\mathbf{z}$ with $\mathbf{A} \in \mathbb{R}^{K \times N}$



Here, y_k is a linear function of z_1, \dots, z_N

$$y_k = \sum_{n=1}^N \mathbf{A}[k, n] z_n$$

where $\mathbf{A}[k, n]$ is entry of \mathbf{A} at row k and column n . We thus can write

$$\frac{\partial \hat{R}}{\partial z_n} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \frac{\partial y_k}{\partial z_n} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \mathbf{A}[k, n]$$

Backpropagation: *Local Operation 2*

Let's denote column n of \mathbf{A} by notation $\mathbf{A}[:, n]$; so, we can write

$$\frac{\partial \hat{R}}{\partial z_{\mathbf{n}}} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \mathbf{A}[k, n] = \nabla_{\mathbf{y}} \hat{R}^T \mathbf{A}[:, n] = \mathbf{A}[:, n]^T \nabla_{\mathbf{y}} \hat{R}$$

Backpropagation: *Local Operation 2*

Let's denote column n of \mathbf{A} by notation $\mathbf{A}[:, n]$; so, we can write

$$\frac{\partial \hat{R}}{\partial z_n} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \mathbf{A}[k, n] = \nabla_{\mathbf{y}} \hat{R}^T \mathbf{A}[:, n] = \mathbf{A}[:, n]^T \nabla_{\mathbf{y}} \hat{R}$$


Now, if we collect them in a vector form we get

$$\nabla_{\mathbf{z}} \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial z_1 \\ \vdots \\ \partial \hat{R} / \partial z_N \end{bmatrix} = \begin{bmatrix} \mathbf{A}[:, 1]^T \\ \vdots \\ \mathbf{A}[:, N]^T \end{bmatrix} \nabla_{\mathbf{y}} \hat{R} = \mathbf{A}^T \nabla_{\mathbf{y}} \hat{R}$$

*This makes sense! Since we are **changing dimensions from K to N** , we need a product that **does such dimensionality change** for us*

Backpropagation: *Local Operation 2*

Long story short . . .




with backward step

$$\nabla_z \hat{R} = A^T \nabla_y \hat{R}$$

Backpropagation: *Local Operation 3*

Linear Matrix-to-Vector Operation

$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{A} \in \mathbb{R}^{K \times N}$ as $\mathbf{y} = \mathbf{A}\mathbf{x}$ with $\mathbf{x} \in \mathbb{R}^N$



- + Wait a moment! The *other variable* is a matrix! How do we define $\nabla_{\mathbf{A}} \hat{R}$?
- Right! Let's first extend the definition

Backpropagation: Local Operation 3

Assume scalar \hat{R} is a function of matrix $\mathbf{A} \in \mathbb{R}^{K \times N}$, we define

$$\nabla_{\mathbf{A}} \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial \mathbf{A} [1, 1] & \dots & \partial \hat{R} / \partial \mathbf{A} [1, N] \\ \vdots & & \vdots \\ \partial \hat{R} / \partial \mathbf{A} [K, 1] & \dots & \partial \hat{R} / \partial \mathbf{A} [K, N] \end{bmatrix}$$

with $\mathbf{A} [k, n]$ being the entry of \mathbf{A} at row k and column n


It is worth to also think of **gradient descent** in this case: assume we are minimizing \hat{R} over \mathbf{A} using **gradient descent** with **learning rate** η . At iteration t we got point $\mathbf{A}^{(t)}$; now, in the next iteration we can readily write

$$\mathbf{A}^{(t+1)} = \mathbf{A}^{(t)} - \nabla_{\mathbf{A}} \hat{R} |_{\mathbf{A}=\mathbf{A}^{(t)}}$$

so apparently everything is as before!

Backpropagation: *Local Operation 3*

Back to our problem, we can write



Entry k of y is a linear function of the k -th row of \mathbf{A} , i.e.,

$$y_k = \sum_{n=1}^N x_n \mathbf{A}[k, n]$$

So, we can write

$$\frac{\partial \hat{R}}{\partial \mathbf{A}[j, n]} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \frac{\partial y_k}{\partial \mathbf{A}[j, n]} = \frac{\partial \hat{R}}{\partial y_k} x_n$$

Backpropagation: *Local Operation 3*

Let's now put them in a matrix


$$\nabla_{\mathbf{A}} \hat{R} = \begin{bmatrix} \frac{\partial \hat{R}}{\partial y_1} x_1 & \dots & \frac{\partial \hat{R}}{\partial y_1} x_N \\ \vdots & & \vdots \\ \frac{\partial \hat{R}}{\partial y_K} x_1 & \dots & \frac{\partial \hat{R}}{\partial y_K} x_N \end{bmatrix} = \nabla_{\mathbf{y}} \hat{R} \mathbf{x}^\top$$

So, we should now *apply outer product*!

This again makes sense! We have a K -dimensional gradient $\nabla_{\mathbf{y}} \hat{R}$ and an N -dimensional vector \mathbf{x} , we need an outer product to convert it into the $K \times N$ matrix $\nabla_{\mathbf{A}} \hat{R}$

Backpropagation: *Local Operation 3*

So, we could conclude




with the backward step

$$\nabla_{\mathbf{A}} \hat{R} = \nabla_y \hat{R} x^\top$$

Now, we are ready to “*backpropagate*” over an FNN


Backpropagation: *Algorithm*

Let's recall the compact diagram of an FNN with L hidden layers




Backpropagation: Algorithm

Let's recall the compact diagram of an FNN with L hidden layers




We can easily expand it into a *computation graph*



Backpropagation: Algorithm


Let's recall the compact diagram of an FNN with L hidden layers



We can easily expand it into a *computation graph*




Our objective is *the empirical risk*; so let's include it also in the graph



Backpropagation: Algorithm

Given data-point \mathbf{x} and its true label \mathbf{v} , we once complete a forward pass




At the end of forward pass,

we know the value of all variables, i.e., \mathbf{z}_ℓ and \mathbf{y}_ℓ for all ℓ

Now, let's assume we want to find $\nabla_{\mathbf{W}_{L+1}} \hat{R}$

Backpropagation: Algorithm

We now cut the graph at the link \mathbf{W}_{L+1}



Let's recall ...

- + *what is the variable here?*
- It's \mathbf{W}_{L+1}
- + *Can we modify the graph such that it becomes a node?*
- Sure! We note that $\mathbf{z}_{L+1} = \mathbf{W}_{L+1}\mathbf{y}_L$. We can look at it as a linear *matrix*-to-vector operation; so, we could modify the graph as

Backpropagation: Algorithm

Let's now move backward to \mathbf{W}_{L+1}


- 1 We have \mathbf{y}_{L+1} , so we compute $\nabla_{\mathbf{y}_{L+1}} \hat{R}$, and pass it to \mathbf{z}_{L+1}
- 2 We have \mathbf{z}_{L+1} , so we compute $\dot{f}_{L+1}(\mathbf{z}_{L+1})$, and then we get

$$\nabla_{\mathbf{z}_{L+1}} \hat{R} = \nabla_{\mathbf{y}_{L+1}} \hat{R} \odot \dot{f}_{L+1}$$

We now pass $\nabla_{\mathbf{z}_{L+1}} \hat{R}$ to \mathbf{W}_{L+1}

- 3 We have \mathbf{y}_L , so we compute $\nabla_{\mathbf{W}_{L+1}} \hat{R}$ from the last pass as

$$\nabla_{\mathbf{W}_{L+1}} \hat{R} = \nabla_{\mathbf{z}_{L+1}} \hat{R} \mathbf{y}_L^T$$




Backpropagation: Algorithm

We can *propagate* backward deeper and deeper

- We cut at the link that *we want to compute the gradient with respect to*
- We exchange the linear vector-to-vector function *at that particular link* to a linear *matrix*-to-vector function
- We move backwards till we get to the source of this graph

Let's see the example for \mathbf{W}_L



Backpropagation: Algorithm


We already have passed backward messages till \mathbf{z}_{L+1}

① We now pass $\nabla_{\mathbf{z}_{L+1}} \hat{R}$ to \mathbf{y}_L : $\nabla_{\mathbf{y}_L} \hat{R} = \mathbf{W}_{L+1}^\top \nabla_{\mathbf{z}_{L+1}} \hat{R}$

② We then pass $\nabla_{\mathbf{y}_L} \hat{R}$ to \mathbf{z}_L : $\nabla_{\mathbf{z}_L} \hat{R} = \nabla_{\mathbf{y}_L} \hat{R} \odot \dot{f}_L$

↳ We should remove entry of $\nabla_{\mathbf{y}_L} \hat{R}$ at index 0: we don't want $\partial \hat{R} / \partial y_L[0]$

③ We finally pass $\nabla_{\mathbf{z}_L} \hat{R}$ to \mathbf{W}_L : $\nabla_{\mathbf{W}_L} \hat{R} = \nabla_{\mathbf{z}_L} \hat{R} \mathbf{y}_{L-1}^\top$



Backpropagation: *Few Notations*

To formally present backpropagation, *let us define a few notations*

For $\ell = 1, \dots, L + 1$, we define

$$\overleftarrow{\mathbf{y}}_{\ell} = \nabla_{\mathbf{y}_{\ell}} \hat{R}$$

$$\overleftarrow{\mathbf{z}}_{\ell} = \nabla_{\mathbf{z}_{\ell}} \hat{R}$$

and keep in mind that

- \mathbf{y}_{ℓ} and $\overleftarrow{\mathbf{y}}_{\ell}$ are totally different things
- \mathbf{z}_{ℓ} and $\overleftarrow{\mathbf{z}}_{\ell}$ are totally different things

Backpropagation: Pseudo Code

```

1: Initiate with  $\overleftarrow{\mathbf{y}}_{L+1} = \nabla \mathcal{L}(\mathbf{y}_{L+1}, \mathbf{v})$  and  $\overleftarrow{\mathbf{z}}_{L+1} = \overleftarrow{\mathbf{y}}_{L+1} \odot \dot{f}_{L+1}(\mathbf{z}_{L+1})$ 
2: for  $\ell = L, \dots, 1$  do
3:   Determine  $\overleftarrow{\mathbf{y}}_{\ell} = \mathbf{W}_{\ell+1}^T \overleftarrow{\mathbf{z}}_{\ell+1}$  and drop  $\overleftarrow{\mathbf{y}}_{\ell}[0]$  # backward affine
4:   Determine  $\overleftarrow{\mathbf{z}}_{\ell} = \dot{f}_{\ell}(\mathbf{z}_{\ell}) \odot \overleftarrow{\mathbf{y}}_{\ell}$  # backward activation
5: end for
6: for  $\ell = 1, \dots, L + 1$  do
7:   Return  $\nabla_{\mathbf{w}_{\ell}} \hat{R} = \overleftarrow{\mathbf{z}}_{\ell} \mathbf{y}_{\ell-1}^T$ 
8: end for

```

- + This looks very similar to forward propagation! Right?!
- Yeah! Just we go **backward**! That's the whole point of backpropagation

You need to go once **forth** and then **back** to determine all gradients

Let's put them next to each other

Backpropagation: Pseudo Code

ForwardProp() :

```

1: Initiate with  $\mathbf{y}_0 = \mathbf{x}$ 
2: for  $\ell = 0, \dots, L$  do
3:   Add  $\mathbf{y}_\ell[0] = 1$  and determine  $\mathbf{z}_{\ell+1} = \mathbf{W}_{\ell+1}\mathbf{y}_\ell$            # forward affine
4:   Determine  $\mathbf{y}_{\ell+1} = f_{\ell+1}(\mathbf{z}_{\ell+1})$                              # forward activation
5: end for
6: for  $\ell = 1, \dots, L + 1$  do
7:   Return  $\mathbf{y}_\ell$  and  $\mathbf{z}_\ell$ 
8: end for

```

BackProp() :

```

1: Initiate with  $\bar{\mathbf{y}}_{L+1} = \nabla \mathcal{L}(\mathbf{y}_{L+1}, \mathbf{v})$  and  $\bar{\mathbf{z}}_{L+1} = \bar{\mathbf{y}}_{L+1} \odot \dot{f}_{L+1}(\mathbf{z}_{L+1})$ 
2: for  $\ell = L, \dots, 1$  do
3:   Determine  $\bar{\mathbf{y}}_\ell = \mathbf{W}_{\ell+1}^\top \bar{\mathbf{z}}_{\ell+1}$  and drop  $\bar{\mathbf{y}}_\ell[0]$            # backward affine
4:   Determine  $\bar{\mathbf{z}}_\ell = \dot{f}_\ell(\mathbf{z}_\ell) \odot \bar{\mathbf{y}}_\ell$                      # backward activation
5: end for
6: for  $\ell = 1, \dots, L + 1$  do
7:   Return  $\nabla_{\mathbf{W}_\ell} \hat{R} = \bar{\mathbf{z}}_\ell \mathbf{y}_{\ell-1}^\top$ 
8: end for

```

Complete Training Loop via *Gradient Descent*

- + Say we use *backpropagation*; then, how does *gradient descent* look?
- Well! We should go back and forth for all data-points

Say we have dataset

$$\mathbb{D} = \{(\mathbf{x}_b, \mathbf{v}_b) \text{ for } b = 1, \dots, B\}$$

GradientDescent():

```

1: Initiate with some initial values  $\{\mathbf{W}_\ell^{(0)}\}$  and set a learning rate  $\eta$ 
2: while weights not converged do
3:   for  $b = 1, \dots, B$  do
4:     NN.values  $\leftarrow$  ForwardProp ( $\mathbf{x}_b, \{\mathbf{W}_\ell^{(t)}\}$ )           # forward
5:      $\{\nabla_{\mathbf{W}_\ell^{(t)}} \hat{R}_b\} \leftarrow$  BackProp ( $\mathbf{x}_b, \mathbf{v}_b, \{\mathbf{W}_\ell^{(t)}\}, \text{NN.values}$ )   # backward
6:   end for
7:   for  $\ell = 1, \dots, L + 1$  do
8:      $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \text{ mean}(\nabla_{\mathbf{W}_\ell^{(t)}} \hat{R}_1, \dots, \nabla_{\mathbf{W}_\ell^{(t)}} \hat{R}_B)$ 
9:   end for
10: end while

```

Binary Classification via FNN

Let's now design a deep FNN for *binary classification*

We have a set of *images* of *hand-written numbers*, something like this^a



We intend to train a fully-connected FNN that given a new hand-written image, *it finds out whether it is “2” or not*

^aSource: Wikipedia

This is a *binary classification*!

Binary Classification via FNN: *Data*

Let's get clear about the data: *our dataset looks like*

$$\mathbb{D} = \{(\mathbf{x}_b, v_b) \text{ for } b = 1, \dots, B\}$$

where in this set each component is defined as follows:

- B is the number of **images** we have
 - ↳ we also call the **set** of images: a **batch** of images
- $\mathbf{x}_b \in \mathbb{R}^N$ is the *pixel vector* of image b
 - ↳ N is the number of pixels in the image
- $v_b \in \{0, 1\}$ is a binary label indicating *whether it is "2" or not*
 - ↳ if the image is a hand-written "2" we set $v_b = 1$
 - ↳ if the image is **not** a hand-written "2" we set $v_b = 0$

Binary Classification via FNN: Model

Let's now set the **model**: we use a *fully-connected FNN*

What are the **hyperparameters**?


- We want it to be **deep**; so, we consider **2 hidden layers**
 - ↳ the **depth** is hence **3**
- We specify the width of each hidden layer
 - ↳ first hidden layer has **width K**
 - ↳ second hidden layer has **width J**
- All hidden neurons use **ReLU activation**
 - ↳ $f_1(\cdot) = f_2(\cdot) = \text{ReLU}(\cdot)$: let's **show ReLU by r** , i.e.,

$$r(x) = \text{ReLU}(x)$$

- **Output layer** has a single **perceptron**

We can now write down the model!

Binary Classification via FNN: Model




Binary Classification via FNN: Model

Let's now set the **model**: we use a *fully-connected FNN*

What are the *learnable parameters*?

- Layer 1 has $(N + 1)K$ links
 - ↳ NK of them are *weights*
 - ↳ K of them are *biases* \equiv *weights* of *dummy node* $x_0 = 1$
- Layer 2 has $(K + 1)J$ links
 - ↳ KJ of them are *weights*
 - ↳ J of them are *biases* \equiv *weights* of *dummy node* $y_1 [0] = 1$
- Output layer has $J + 1$ links
 - ↳ J of them are *weights*
 - ↳ one is *bias* \equiv *weight* of *dummy node* $y_2 [0] = 1$

Binary Classification via FNN: Model



Binary Classification via FNN: Loss

How to calculate the loss? *Let's do what we did before*


We use the error indicator as the loss function

$$\mathcal{L}(y_b, v_b) = \mathbb{1}\{y_b \neq v_b\} = \begin{cases} 1 & y_b \neq v_b \\ 0 & y_b = v_b \end{cases}$$

- + *Wait a moment! Didn't you say that this was a **bad choice**?*
- Yeah! So said I also for the **perceptron's** activation! Let's try it out to find out really **why they are bad**! We should be able to understand it now

Binary Classification via FNN: *Training*

Let's look at the computation graph: *for a given data-point* (\mathbf{x}_b, v_b) , we have




Here, we have 3 *linear* operations

- First operation is $\mathbf{z}_1 = \mathbf{W}_1 \mathbf{x}$ with $\mathbf{W}_1 \in \mathbb{R}^{K \times (N+1)}$
 \hookrightarrow first column of \mathbf{W}_1 is *bias* and the remaining columns are weights
- Second operation is $\mathbf{z}_2 = \mathbf{W}_2 \mathbf{y}_1$ with $\mathbf{W}_2 \in \mathbb{R}^{J \times (K+1)}$
 \hookrightarrow first column of \mathbf{W}_2 is *bias* and the remaining columns are weights
- Last operation is $\mathbf{z}_3 = \mathbf{w}_3^T \mathbf{x}$ with $\mathbf{w}_3 \in \mathbb{R}^{J+1}$
 \hookrightarrow first entry of \mathbf{w}_3 is *bias* and the remaining entries are weights

Binary Classification via FNN: *Training*

Let's look at the computation graph: *for a given data-point* (\mathbf{x}_b, v_b) , we have



We have 3 *functional* operations

- The first two are $\mathbf{y}_1 = r(\mathbf{z}_1)$ and $\mathbf{y}_2 = r(\mathbf{z}_2)$
- The last one is $\mathbf{y}_3 = s(\mathbf{z}_3)$, and recall that $s(\cdot)$ is the **step function**

$$y_3 = s(z_3) = \begin{cases} 1 & z_3 \geq 0 \\ 0 & z_3 < 0 \end{cases}$$

Binary Classification via FNN: *Training*

Let's write **gradient descent** for training of our model

GradientDescent() :

```

1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{w}_3^{(0)}\}$  and set a learning rate  $\eta$ 
2: while weights not converged do
3:   for  $b = 1, \dots, B$  do
4:      $\text{NN.values} \leftarrow \text{ForwardProp}(\mathbf{x}_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{w}_3^{(t)}\})$ 
5:      $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{g}_{3,b}\} \leftarrow \text{BackProp}(\mathbf{x}_b, \mathbf{v}_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{w}_3^{(t)}\}, \text{NN.values})$ 
6:   end for
7:   Update

```

$$\mathbf{W}_1^{(t+1)} \leftarrow \mathbf{W}_1^{(t)} - \eta \text{ mean}(\mathbf{G}_{1,1}, \dots, \mathbf{G}_{1,B})$$

$$\mathbf{W}_2^{(t+1)} \leftarrow \mathbf{W}_2^{(t)} - \eta \text{ mean}(\mathbf{G}_{2,1}, \dots, \mathbf{G}_{2,B})$$

$$\mathbf{w}_3^{(t+1)} \leftarrow \mathbf{w}_3^{(t)} - \eta \text{ mean}(\mathbf{g}_{3,1}, \dots, \mathbf{g}_{3,B})$$

8: **end while**

Let's look at *forward* and *backward* propagation!

Binary Classification via FNN: *Forward Pass*

Forward pass is very straightforward: say we are at iteration t

- 1 For each pixel vector \mathbf{x}_b , we determine \mathbf{z}_1 as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The output of first layer is then given by $\mathbf{y}_1 = r(\mathbf{z}_1)$: $r(\cdot)$ is ReLU, so

we keep positive entries of \mathbf{z}_1 and replace negative ones with zero

Binary Classification via FNN: *Forward Pass*

Forward pass is very straightforward: say we are at iteration t

- 1 For each pixel vector \mathbf{x}_b , we determine \mathbf{z}_1 as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The output of first layer is then given by $\mathbf{y}_1 = r(\mathbf{z}_1)$: $r(\cdot)$ is ReLU, so

we keep positive entries of \mathbf{z}_1 and replace negative ones with zero

- 2 We add 1 at index 0 of \mathbf{y}_1 and determine \mathbf{z}_2 as

$$\mathbf{y}_1 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_1 \end{bmatrix} \rightsquigarrow \mathbf{z}_2 = \mathbf{W}_2^{(t)} \mathbf{y}_1$$

The output of second layer is given by $\mathbf{y}_2 = r(\mathbf{z}_2)$

Binary Classification via FNN: *Forward Pass*

- ③ We add 1 at index 0 of \mathbf{y}_2 and determine z_3 as

$$\mathbf{y}_2 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_2 \end{bmatrix} \rightsquigarrow z_3 = \mathbf{w}_3^{(t)\top} \mathbf{y}_2$$

The network **output** is given by $y_3 = s(z_3)$: $s(\cdot)$ is step function, so

it's 0 if z_3 is negative, and 1 if it is not negative

Binary Classification via FNN: *Forward Pass*

- ③ We add 1 at index 0 of \mathbf{y}_2 and determine z_3 as

$$\mathbf{y}_2 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_2 \end{bmatrix} \rightsquigarrow z_3 = \mathbf{w}_3^{(t)\top} \mathbf{y}_2$$

The network **output** is given by $y_3 = s(z_3)$: $s(\cdot)$ is step function, so

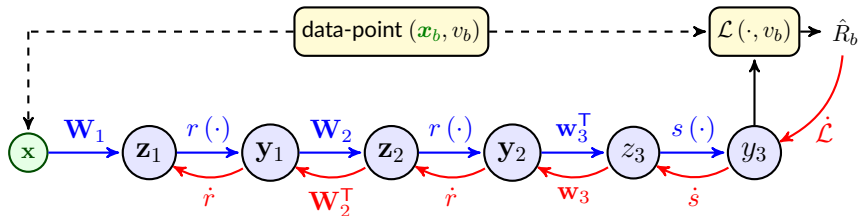
it's 0 if z_3 is negative, and 1 if it is not negative

At this point, we have all that we need, i.e.,

\mathbf{x} , \mathbf{z}_1 , \mathbf{y}_1 , \mathbf{z}_2 , \mathbf{y}_2 , z_3 and y_3

Binary Classification via FNN: *Training*

How does the graph look like on the backward pass?



Let's move backward!

Binary Classification via FNN: *Backward Pass*

We know all the derivatives, i.e.,

$$\dot{\mathcal{L}}(\textcolor{red}{y}, v_b) = \frac{d}{d\textcolor{red}{y}} \mathbb{1}\{\textcolor{red}{y} \neq v_b\} \quad \dot{s}(z) = \frac{d}{dz} s(z) \quad \dot{r}(z) = \frac{d}{dz} r(z)$$

Binary Classification via FNN: *Backward Pass*

We know all the derivatives, i.e.,

$$\dot{\mathcal{L}}(\textcolor{red}{y}, v_b) = \frac{d}{d\textcolor{red}{y}} \mathbb{1}\{\textcolor{red}{y} \neq v_b\} \quad \dot{s}(z) = \frac{d}{dz} s(z) \quad \dot{r}(z) = \frac{d}{dz} r(z)$$

For backward pass we start at node $\textcolor{red}{y}_3$:

- 1 We find derivative w.r.t. output $\overleftarrow{\textcolor{red}{y}}_3 = \dot{\mathcal{L}}(\textcolor{red}{y}_3, v_b)$ and set

$$\overleftarrow{z}_3 = \overleftarrow{\textcolor{red}{y}}_3 \dot{s}(z_3)$$

Binary Classification via FNN: *Backward Pass*

We know all the derivatives, i.e.,

$$\dot{\mathcal{L}}(\textcolor{red}{y}, v_b) = \frac{d}{d\textcolor{red}{y}} \mathbb{1}\{\textcolor{red}{y} \neq v_b\} \quad \dot{s}(z) = \frac{d}{dz} s(z) \quad \dot{r}(z) = \frac{d}{dz} r(z)$$

For backward pass we start at node $\textcolor{red}{y}_3$:

- 1 We find derivative w.r.t. output $\overleftarrow{\textcolor{red}{y}}_3 = \dot{\mathcal{L}}(\textcolor{red}{y}_3, v_b)$ and set

$$\overleftarrow{z}_3 = \overleftarrow{\textcolor{red}{y}}_3 \dot{s}(z_3)$$

- 2 We compute $\overleftarrow{\textcolor{red}{y}}_2 = \mathbf{w}_3 \overleftarrow{\textcolor{green}{z}}_3$ and drop its first entry; then, compute

$$\overleftarrow{\textcolor{red}{y}}_2 \leftarrow \begin{bmatrix} \overleftarrow{\textcolor{red}{y}}_2[0] \\ \overleftarrow{\textcolor{red}{y}}_2[1:] \end{bmatrix} \rightsquigarrow \overleftarrow{\textcolor{green}{z}}_2 = \dot{r}(\mathbf{z}_2) \odot \overleftarrow{\textcolor{red}{y}}_2$$

Binary Classification via FNN: *Backward Pass*

- ③ We compute $\hat{\mathbf{y}}_1 = \mathbf{W}_2^T \hat{\mathbf{z}}_2$ and drop its first entry; then, compute

$$\hat{\mathbf{y}}_1 \leftarrow \begin{bmatrix} \hat{\mathbf{y}}_1[0] \\ \hat{\mathbf{y}}_1[1:] \end{bmatrix} \rightsquigarrow \hat{\mathbf{z}}_1 = \dot{r}(\mathbf{z}_1) \odot \hat{\mathbf{y}}_1$$

Binary Classification via FNN: *Backward Pass*

- ③ We compute $\hat{\mathbf{y}}_1 = \mathbf{W}_2^T \hat{\mathbf{z}}_2$ and drop its first entry; then, compute

$$\hat{\mathbf{y}}_1 \leftarrow \begin{bmatrix} \hat{\mathbf{y}}_1[0] \\ \hat{\mathbf{y}}_1[1:] \end{bmatrix} \rightsquigarrow \hat{\mathbf{z}}_1 = \dot{r}(\mathbf{z}_1) \odot \hat{\mathbf{y}}_1$$

At this point, we can calculate all gradients

$$\mathbf{G}_{1,b} = \nabla_{\mathbf{w}_1} \hat{R}_b = \hat{\mathbf{z}}_1 \mathbf{y}_0^T = \hat{\mathbf{z}}_1 \mathbf{x}^T$$

$$\mathbf{G}_{2,b} = \nabla_{\mathbf{w}_2} \hat{R}_b = \hat{\mathbf{z}}_2 \mathbf{y}_1^T$$

$$\mathbf{g}_{3,b}^T = \nabla_{\mathbf{w}_3^T} \hat{R}_b = \hat{\mathbf{z}}_3 \mathbf{y}_2^T$$

All done! We repeat it for *every image* in the *batch* and then average gradients. Finally, we move one step in *gradient descent* and find the weights of the next iteration

Binary Classification via FNN: *Differentiability Issue*

- + Where is then the issue with *perceptron* and *indicator error*?
- $\dot{\mathcal{L}}(y, v_b)$ and $\dot{s}(z)$ are not well-defined!
 - ↳ Recall that they are *discontinuous*

In fact, the empirical risk is not a *smooth function* of the weights and biases; therefore, using *gradient descent* we do not end up with a well-trained network

- + How can we get over it?
- Well! There is a very well-established trick!


Binary Classification via FNN: *Differentiability Issue*

We first replace the **perceptron** with a neuron whose **activation** is a good approximation of *step function* and *differentiable*¹

We already have seen the **sigmoid** function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

which looks pretty close to step function



Using sigmoid instead of step function **resolves** the **differentiability** issue

¹Or at least, we can easily calculate a sub-gradient for it

Binary Classification via FNN: *Differentiability Issue*

But, replacing *perceptron* by *sigmoid-activated neuron* makes a new problem

The *output* of the network is now *not binary*!

How can we address this problem?

We now interpret the *output* as *probability*, i.e.,

y_3 is the *probability* of the *label being 1*

Binary Classification via FNN: *Differentiability Issue*

But, replacing *perceptron* by *sigmoid-activated neuron* makes a new problem

The *output* of the network is now *not binary*!

How can we address this problem?

We now interpret the *output* as *probability*, i.e.,

y_3 is the *probability* of the *label being 1*

- + OK! But how can we define the *loss* now?
- Well! We could look at the *true label* from the same point of view

Say $v \in \{0, 1\}$ is *true label*: if $v = 1$ then the *true label is 1* with probability 1; if $v = 0$ then the *true label is 1* with probability 0. So, we could say

the *true label is 1* with probability v

Binary Classification via FNN: *Differentiability Issue*

true label is 1 with probability v \longleftrightarrow y_3 is probability of the *label being 1*

Apparently, v and y_3 are of the same nature: we can still define a *loss* that evaluates the difference between y_3 and v

- + What should be the loss then?
- Definitely **not the indicator error!**

Binary Classification via FNN: *Differentiability Issue*

true label is 1 with probability v \longleftrightarrow y_3 is probability of the *label being 1*

Apparently, v and y_3 are of the same nature: we can still define a *loss* that evaluates the difference between y_3 and v

- + What should be the loss then?
- Definitely **not the indicator error!**

Indicator error is not suitable because

- 1 we already now that it is **not differentiable**
- 2 more importantly, with **sigmoid activation** becomes useless

$$\mathbb{1}\{\sigma(z_3) \neq 1\} = \mathbb{1}\{\sigma(z_3) \neq 0\} = 1$$

Binary Classification via FNN: MSE

One may suggest that we use the *squared error*, i.e.,

$$\mathcal{L}(y_3, v) = (y_3 - v)^2$$

in this case the empirical risk is called

Mean Squared Error (MSE)

This loss is *differentiable*

$$\dot{\mathcal{L}}(y_3, v) = 2(y_3 - v)$$

and proportional to the distance between y_3 and v

*It's a *good* choice but *not best**

Binary Classification via FNN: Cross-Entropy

A better choice is to determine the **cross-entropy loss**

$$\begin{aligned}\mathcal{L}(y_3, v) &= \text{CE}(y_3, v) = -v \log y_3 - (1 - v) \log (1 - y_3) \\ &= \begin{cases} \log \frac{1}{y_3} & v = 1 \\ \log \frac{1}{(1 - y_3)} & v = 0 \end{cases}\end{aligned}$$

This loss function is sometimes **wrongly** called **KL-divergence**: it is proportional to the **Kullback-Leibler divergence** but it's **different**


This loss is again **differentiable**

$$\dot{\mathcal{L}}(y_3, v) = \dot{\text{CE}}(y_3, v) = -\frac{v}{y_3} + \frac{1 - v}{1 - y_3}$$

Note: The logarithm is usually in natural base, i.e., $\log x = \ln x$

Binary Classification via FNN: Cross-Entropy

- + But why *cross-entropy* is a *better loss*?
- It pushes y_3 *more* towards the *edges* of interval $[0, 1]$



Cross entropy returns *much higher loss* when y_3 is different from v

Binary Classification via FNN: *Training with Cross-Entropy*


- + *What changes in the training loop in this case?*
- Pretty much *nothing!* Just replace
 - $\mathcal{L}(y_3, v)$ with $\text{CE}(y_3, v)$
 - $\dot{\mathcal{L}}(y_3, v)$ with $\dot{\text{CE}}(y_3, v)$
 - $s(z_3)$ with $\sigma(z_3)$
 - $\dot{s}(z_3)$ with $\dot{\sigma}(z_3)$

Binary Classification via FNN: Training with Cross-Entropy

+ What changes in the training loop in this case?

- Pretty much *nothing!* Just replace

- $\mathcal{L}(\mathbf{y}_3, v)$ with $\text{CE}(\mathbf{y}_3, v)$
- $\dot{\mathcal{L}}(\mathbf{y}_3, v)$ with $\dot{\text{CE}}(\mathbf{y}_3, v)$
- $s(z_3)$ with $\sigma(z_3)$
- $\dot{s}(z_3)$ with $\dot{\sigma}(z_3)$



Binary Classification via FNN: Training with Cross-Entropy

- + How do we use the output of network then, when we give a new *image* to it for classification? *It's not binary!*
- Just follow the *interpretation*

y_3 gives the *probability* of the *image* being *hand-written "2"*; therefore, $(1 - y_3)$ gives the *probability* of image being *any other hand-written number*. So, we select the outcome with *higher chance*, i.e.,

- if $y_3 \geq 0.5$, we label the new *image* as a hand-written "2"
- if $y_3 < 0.5$, we label the new *image* as *not being* a hand-written "2"

- + Can't we classify more classes? Like hand-written "0", "1", ..., "9"?
- Now that we have this nice interpretation: Yes! We can!

Multiclass Classification

We initially saw that any multiclass classification can be seen as a *sequence of binary classifications*; however, for that, we need **multiple NNs**!


- + *Why not follow the same idea and determine the probability of input belonging to each class?*
- Yes! That's actually the effective way!

Let's get back to our *image recognition*, but now with *multiple* classes!

We have *images* of *hand-written numbers from "0" to "9"* and want to train a NN that *recognizes any hand-written number*

We first draw our earlier FNN

Multiclass Classification via FNN




In this FNN, y_3 is interpreted as a *probability of label being 1*

probability of label being 0 is hence $1 - y_3$

This was done by a standard *single-output neuron*, since we had only *2 classes*

Multiclass Classification via FNN



With C classes, we need a module that computes probabilities of all C classes
this module can be seen as a **neuron** with **vector output**

Multiclass Classification: Vector-Activated Neuron

Vector-Activated Neuron

A **vector-activated neuron** is an artificial neuron with **multivariate activation function**: let $\mathbf{x} \in \mathbb{R}^N$ be the input to this neuron and C be its output dimension; then, the output vector $\mathbf{y} \in \mathbb{R}^C$ is given by

$$\mathbf{y} = F\left(\tilde{\mathbf{W}}\mathbf{x} + \mathbf{b}\right)$$

for weight matrix $\tilde{\mathbf{W}} \in \mathbb{R}^{C \times N}$, bias $\mathbf{b} \in \mathbb{R}^C$ and **activation** $F(\cdot) : \mathbb{R}^C \mapsto \mathbb{R}^C$

Multiclass Classification: Vector-Activated Neuron

Vector-Activated Neuron

A **vector-activated neuron** is an artificial neuron with **multivariate activation function**: let $\mathbf{x} \in \mathbb{R}^N$ be the input to this neuron and C be its output dimension; then, the output vector $\mathbf{y} \in \mathbb{R}^C$ is given by

$$\mathbf{y} = F(\tilde{\mathbf{W}}\mathbf{x} + \mathbf{b})$$


for weight matrix $\tilde{\mathbf{W}} \in \mathbb{R}^{C \times N}$, bias $\mathbf{b} \in \mathbb{R}^C$ and **activation** $F(\cdot) : \mathbb{R}^C \mapsto \mathbb{R}^C$

First thing first: let's get rid of the bias before we go on


$$\mathbf{y} = F\left(\begin{bmatrix} \mathbf{b} & \tilde{\mathbf{W}} \end{bmatrix} \begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix}\right) = F(\mathbf{W}\mathbf{x})$$

So, we keep on with our dummy 1 input here as well

Multiclass Classification: Vector-Activated Neuron



Next, let's see how its computation graph looks




This looks exactly like a standard layer with a **minor difference**

$F(\cdot)$ does **not necessarily** perform **entry-wise**

Multiclass Classification: *Softmax*

A very well-known example of **vector activation** is *softmax*



Softmax Function

For $\mathbf{z} \in \mathbb{R}^C$, softmax function returns $\text{Soft}_{\max}(\mathbf{z}) = \mathbf{y} \in \mathbb{R}^C$ whose entry c is

$$y[c] = \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}}$$

Multiclass Classification: *Softmax*

Softmax always returns a probability distribution on the set of classes

$$\sum_{c=1}^C y[c] =$$

Multiclass Classification: *Softmax*

Softmax always returns a probability distribution on the set of classes

$$\sum_{c=1}^C y[c] = \sum_{c=1}^C \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}}$$

Multiclass Classification: *Softmax*

Softmax always returns a probability distribution on the set of classes

$$\sum_{c=1}^C y[c] = \sum_{c=1}^C \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = \frac{\sum_{c=1}^C e^{z[c]}}{\sum_{j=1}^C e^{z[j]}}$$

Multiclass Classification: *Softmax*

Softmax always returns a probability distribution on the set of classes

$$\sum_{c=1}^C y[c] = \sum_{c=1}^C \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = \frac{\sum_{c=1}^C e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = 1$$

Multiclass Classification: *Softmax*


Softmax always returns a probability distribution on the set of classes

$$\sum_{c=1}^C y[c] = \sum_{c=1}^C \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = \frac{\sum_{c=1}^C e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = 1$$

We can hence use it to extend our FNN to a *multiclass classifier*

We replace *layer 3* with a *softmax-activated multivariate neuron* and treat its outcome as the *chance of input belonging to each class*; then, we select *the class with highest chance*


Multiclass Classification via FNN: *Softmax Activation*



Let's try again the forward pass

Multiclass Classification via FNN: *Softmax Activation*

Let's look at the computation graph: *for a given data-point (\mathbf{x}_b, v_b) , we have*




Note that **the output layer** has been changed

- We now have a **vector** $\mathbf{z}_3 \in \mathbb{R}^C$
- We now have a **vector** $\mathbf{y}_3 \in \mathbb{R}^C$
- We get from \mathbf{z}_3 to \mathbf{y}_3 via **softmax**

Multiclass Classification via FNN: *Softmax Activation*

Let's look at the computation graph: *for a given data-point (\mathbf{x}_b, v_b) , we have*




Note that **the output layer** has been changed

- We now have a **vector** $\mathbf{z}_3 \in \mathbb{R}^C$
 ↳ So we have a **matrix of weights** $\mathbf{W}_3 \in \mathbb{R}^{C \times (J+1)}$
- We now have a **vector** $\mathbf{y}_3 \in \mathbb{R}^C$
- We get from \mathbf{z}_3 to \mathbf{y}_3 via **softmax**

Multiclass Classification via FNN: *Softmax Activation*

Let's look at the computation graph: *for a given data-point (\mathbf{x}_b, v_b) , we have*



Note that **the output layer** has been changed

- We now have a **vector** $\mathbf{z}_3 \in \mathbb{R}^C$
 ↳ So we have a **matrix of weights** $\mathbf{W}_3 \in \mathbb{R}^{C \times (J+1)}$
- We now have a **vector** $\mathbf{y}_3 \in \mathbb{R}^C$
- We get from \mathbf{z}_3 to \mathbf{y}_3 via **softmax**
 ↳ This is **not** an entry-wise activation anymore!

Multiclass Classification via FNN: *Softmax Activation*

- 1 For each pixel vector \mathbf{x}_b , we determine \mathbf{z}_1 as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The *output* of first layer is then given by $\mathbf{y}_1 = r(\mathbf{z}_1)$

Multiclass Classification via FNN: *Softmax Activation*

- ① For each pixel vector \mathbf{x}_b , we determine \mathbf{z}_1 as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The *output* of first layer is then given by $\mathbf{y}_1 = r(\mathbf{z}_1)$

- ② We add 1 at index 0 of \mathbf{y}_1 and determine \mathbf{z}_2 as

$$\mathbf{y}_1 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_1 \end{bmatrix} \rightsquigarrow \mathbf{z}_2 = \mathbf{W}_2^{(t)} \mathbf{y}_1$$

The output of second layer is given by $\mathbf{y}_2 = r(\mathbf{z}_2)$

Multiclass Classification via FNN: *Softmax Activation*

- 1 For each **pixel vector** \mathbf{x}_b , we determine \mathbf{z}_1 as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The **output** of first layer is then given by $\mathbf{y}_1 = r(\mathbf{z}_1)$

- 2 We add 1 at index 0 of \mathbf{y}_1 and determine \mathbf{z}_2 as

$$\mathbf{y}_1 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_1 \end{bmatrix} \rightsquigarrow \mathbf{z}_2 = \mathbf{W}_2^{(t)} \mathbf{y}_1$$

The output of second layer is given by $\mathbf{y}_2 = r(\mathbf{z}_2)$

- 3 We add 1 at index 0 of \mathbf{y}_2 and determine \mathbf{z}_3 as

$$\mathbf{y}_2 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_2 \end{bmatrix} \rightsquigarrow \mathbf{z}_3 = \mathbf{W}_3^{(t)\top} \mathbf{y}_2$$

The network **output** is given by $\mathbf{y}_3 = \text{Soft}_{\max}(\mathbf{z}_3)$

Multiclass Classification: Loss

- + *How can we define the loss now? On one side we have a vector of probabilities; on the other side an integer label!*
- Again we need to convert *true labels* to *probabilities*

Multiclass Classification: Loss

- + How can we define the loss now? On one side we have a vector of probabilities; on the other side an integer label!
- Again we need to convert *true labels* to *probabilities*

Let's say we have C classes: the *vector* of *probabilities* contains C entries

$$\mathbf{p} = \begin{bmatrix} p_1 \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{ \text{image belongs to class } 1 \} \\ \vdots \\ \Pr \{ \text{image belongs to class } C \} \end{bmatrix}$$

Multiclass Classification: Loss

- + How can we define the loss now? On one side we have a vector of probabilities; on the other side an integer label!
- Again we need to convert *true labels* to *probabilities*

Let's say we have C classes: the *vector* of *probabilities* contains C entries

$$\mathbf{p} = \begin{bmatrix} p_1 \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{ \text{image belongs to class } 1 \} \\ \vdots \\ \Pr \{ \text{image belongs to class } C \} \end{bmatrix}$$

If we know that the *image* b belongs to class v_b , we could say that

$$\mathbf{p} \text{ of image } b = \begin{bmatrix} p_1 \\ \vdots \\ p_{v_b} \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{ \text{image } b \text{ belongs to class } 1 \} \\ \vdots \\ \Pr \{ \text{image } b \text{ belongs to class } v_b \} \\ \vdots \\ \Pr \{ \text{image } b \text{ belongs to class } C \} \end{bmatrix}$$

Multiclass Classification: Loss

- + How can we define the loss now? On one side we have a vector of probabilities; on the other side an integer label!
- Again we need to convert *true labels* to *probabilities*

Let's say we have C classes: the *vector* of *probabilities* contains C entries

$$\mathbf{p} = \begin{bmatrix} p_1 \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{\text{image belongs to class } 1\} \\ \vdots \\ \Pr \{\text{image belongs to class } C\} \end{bmatrix}$$

If we know that the *image* b belongs to class v_b , we could say that

$$\mathbf{p} \text{ of image } b = \begin{bmatrix} p_1 \\ \vdots \\ p_{v_b} \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{\text{image } b \text{ belongs to class } 1\} = 0 \\ \vdots \\ \Pr \{\text{image } b \text{ belongs to class } v_b\} = 1 \\ \vdots \\ \Pr \{\text{image } b \text{ belongs to class } C\} = 0 \end{bmatrix} = \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix}$$

Multiclass Classification: Loss

So, we could say that *label v is corresponding to a vector of size C whose entry v is 1 and the remaining entries are all 0*: this vector is called a *one-hot vector*

Multiclass Classification: Loss

So, we could say that *label v is corresponding to a vector of size C whose entry v is 1 and the remaining entries are all 0*: this vector is called a *one-hot vector*

One-hot Vector

The one-hot vector $\mathbf{1}_v \in \{0, 1\}^C$ is a C -dimensional vector whose entry v is 1 and all remaining entries are 0

For instance: say $C = 3$; then, we have

$$\mathbf{1}_1 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \mathbf{1}_2 = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \mathbf{1}_3 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

Multiclass Classification: Loss

So, we could say that *label v is corresponding to a vector of size C whose entry v is 1 and the remaining entries are all 0*: this vector is called a *one-hot vector*

One-hot Vector

The one-hot vector $\mathbf{1}_v \in \{0, 1\}^C$ is a C -dimensional vector whose entry v is 1 and all remaining entries are 0

For instance: say $C = 3$; then, we have

$$\mathbf{1}_1 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \mathbf{1}_2 = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \mathbf{1}_3 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

Moral of Story

*We can interpret *true label v* as a probability vector $\mathbf{1}_v$*

Multiclass Classification: Loss

We can interpret *true label* v as a probability vector $\mathbf{1}_v$

Now for *image* b with *label* v_b we compare network's output \mathbf{y}_3 to $\mathbf{1}_{v_b}$

$$\hat{R}_b = \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$$

for *loss* $\mathcal{L}(\cdot)$ that determines distance between two *probability vectors*

Multiclass Classification: Loss

We can interpret *true label* v as a probability vector $\mathbf{1}_v$

Now for *image* b with *label* v_b we compare network's output \mathbf{y}_3 to $\mathbf{1}_{v_b}$

$$\hat{R}_b = \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$$

for *loss* $\mathcal{L}(\cdot)$ that determines distance between two *probability vectors*

- + What kind of *loss functions* do we use usually?
- Like binary case: squared error is *good*, cross-entropy is the *best*
- + How do we define them in this case?
- Just extend them to multi-dimensional vectors

Multiclass Classification: Loss

We can extend squared error to vector form as

$$\begin{aligned}\mathcal{L}(\mathbf{y}, \mathbf{1}_v) &= \|\mathbf{y} - \mathbf{1}_v\|^2 = \sum_{c=1}^C (y[c] - \mathbb{1}\{c=v\})^2 \\ &= \sum_{c=1, c \neq v}^C y[c]^2 + (y[v] - 1)^2\end{aligned}$$

This gradient of this loss is

$$\nabla \mathcal{L}(\mathbf{y}, \mathbf{1}_v) = 2 \begin{bmatrix} y[1] \\ \vdots \\ y[v] - 1 \\ \vdots \\ y[C] \end{bmatrix} = 2(\mathbf{y} - \mathbf{1}_v)$$

Multiclass Classification: Loss

Cross entropy can also be extended as follows


$$\begin{aligned}\mathcal{L}(\mathbf{y}, \mathbf{1}_v) &= \text{CE}(\mathbf{y}, \mathbf{1}_v) = - \sum_{c=1}^C \mathbb{1}\{c = v\} \log y(c) \\ &= -\log y[v]\end{aligned}$$

The gradient of this loss is

$$\nabla \mathcal{L}(\mathbf{y}, \mathbf{1}_v) = \nabla \text{CE}(\mathbf{y}, \mathbf{1}_v) = \begin{bmatrix} 0 \\ \vdots \\ -1/y[v] \\ \vdots \\ 0 \end{bmatrix} = -\frac{1}{y[v]} \mathbf{1}_v$$

Backpropagation Through Vector Activation


How can we backpropagate through this neural network?



- 1 Compute $\nabla_{\mathbf{y}_3} \hat{R}_b = \nabla_{\mathbf{y}_3} \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$

Backpropagation Through Vector Activation


How can we backpropagate through this neural network?



- 1 Compute $\nabla_{\mathbf{y}_3} \hat{R}_b = \nabla_{\mathbf{y}_3} \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$
- 2 Compute $\nabla_{\mathbf{z}_3} \hat{R}_b$ from $\nabla_{\mathbf{y}_3} \hat{R}_b$ by its *backward* link that we don't know yet

Backpropagation Through Vector Activation


How can we backpropagate through this neural network?



- 1 Compute $\nabla_{\mathbf{y}_3} \hat{R}_b = \nabla_{\mathbf{y}_3} \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$
- 2 Compute $\nabla_{\mathbf{z}_3} \hat{R}_b$ from $\nabla_{\mathbf{y}_3} \hat{R}_b$ by its **backward** link that we don't know yet
- 3 Compute $\nabla_{\mathbf{y}_2} \hat{R}_b = \mathbf{W}_3^T \nabla_{\mathbf{z}_3} \hat{R}_b$

Backpropagation Through Vector Activation


How can we backpropagate through this neural network?



- 1 Compute $\nabla_{\mathbf{y}_3} \hat{R}_b = \nabla_{\mathbf{y}_3} \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$
- 2 Compute $\nabla_{\mathbf{z}_3} \hat{R}_b$ from $\nabla_{\mathbf{y}_3} \hat{R}_b$ by its **backward** link that we don't know yet
- 3 Compute $\nabla_{\mathbf{y}_2} \hat{R}_b = \mathbf{W}_3^T \nabla_{\mathbf{z}_3} \hat{R}_b$
- 4 Remove entry at index 0 of $\nabla_{\mathbf{y}_2} \hat{R}_b$ and compute $\nabla_{\mathbf{z}_2} \hat{R}_b = \dot{r}(\mathbf{z}_2) \odot \nabla_{\mathbf{y}_2} \hat{R}_b$

Backpropagation Through Vector Activation


How can we backpropagate through this neural network?



- 1 Compute $\nabla_{y_3} \hat{R}_b = \nabla_{y_3} \mathcal{L}(y_3, \mathbf{1}_{v_b})$
- 2 Compute $\nabla_{z_3} \hat{R}_b$ from $\nabla_{y_3} \hat{R}_b$ by its **backward** link that we don't know yet
- 3 Compute $\nabla_{y_2} \hat{R}_b = W_3^T \nabla_{z_3} \hat{R}_b$
- 4 Remove entry at index 0 of $\nabla_{y_2} \hat{R}_b$ and compute $\nabla_{z_2} \hat{R}_b = \dot{r}(z_2) \odot \nabla_{y_2} \hat{R}_b$
- 5 Compute $\nabla_{y_1} \hat{R}_b = W_2^T \nabla_{z_2} \hat{R}_b$

Backpropagation Through Vector Activation

How can we backpropagate through this neural network?




- 1 Compute $\nabla_{y_3} \hat{R}_b = \nabla_{y_3} \mathcal{L}(y_3, \mathbf{1}_{v_b})$
- 2 Compute $\nabla_{z_3} \hat{R}_b$ from $\nabla_{y_3} \hat{R}_b$ by its **backward** link that we don't know yet
- 3 Compute $\nabla_{y_2} \hat{R}_b = W_3^T \nabla_{z_3} \hat{R}_b$
- 4 Remove entry at index 0 of $\nabla_{y_2} \hat{R}_b$ and compute $\nabla_{z_2} \hat{R}_b = \dot{r}(z_2) \odot \nabla_{y_2} \hat{R}_b$
- 5 Compute $\nabla_{y_1} \hat{R}_b = W_2^T \nabla_{z_2} \hat{R}_b$
- 6 Remove entry at index 0 of $\nabla_{y_1} \hat{R}_b$ and compute $\nabla_{z_1} \hat{R}_b = \dot{r}(z_1) \odot \nabla_{y_1} \hat{R}_b$

Backpropagation Through Vector Activation

How is $\nabla_{\mathbf{z}_3} \hat{R}_b$ related to $\nabla_{\mathbf{y}_3} \hat{R}_b$? Let's do what we did before

In this graph, $F(\cdot)$ is a **vector activation**. We know $\nabla_{\mathbf{y}} \hat{R}$



We want to find $\nabla_{\mathbf{z}} \hat{R}$

As mentioned before: we can always *extend things entry-wise*

With **vector activation**, we need to use the notion of **Jacobian**

Recap: Jacobian Matrix

Consider **vector activation** $F(\cdot)$ that maps C -dimensional $\mapsto C$ -dimensional

$$\begin{bmatrix} y_1 \\ \vdots \\ y_C \end{bmatrix} = F\left(\begin{bmatrix} z_1 \\ \vdots \\ z_C \end{bmatrix}\right)$$

When we use this function, we can say

Any entry y_j is function of all^a z_1, \dots, z_C , so we have

$$\nabla_{\mathbf{z}} y_j = \begin{bmatrix} \partial y_j / \partial z_1 \\ \vdots \\ \partial y_j / \partial z_C \end{bmatrix}$$

^aIt is not any more an entry-wise functional operation

Recap: Jacobian Matrix

Consider **vector activation** $F(\cdot)$ that maps C -dimensional $\mapsto C$ -dimensional

$$\begin{bmatrix} y_1 \\ \vdots \\ y_C \end{bmatrix} = F\left(\begin{bmatrix} z_1 \\ \vdots \\ z_C \end{bmatrix}\right)$$

When we use this function, we can say

We can **collect all these gradients** into a **matrix**


$$\mathbf{J}_{\mathbf{z}}\mathbf{y} = \mathbf{J}_{\mathbf{z}}F = \begin{bmatrix} \nabla_{\mathbf{z}} y_1^\top \\ \vdots \\ \nabla_{\mathbf{z}} y_C^\top \end{bmatrix} = \begin{bmatrix} \partial y_1 / \partial z_1 & \dots & \partial y_1 / \partial z_C \\ \vdots & & \vdots \\ \partial y_C / \partial z_1 & \dots & \partial y_C / \partial z_C \end{bmatrix}$$

and we call it the **Jacobian matrix**

Backpropagation Through Vector Activation

Now, let's get back to our problem

In this graph, $F(\cdot)$ is a **vector activation**. We know $\nabla_{\mathbf{y}} \hat{R}$




We want to find $\nabla_{\mathbf{z}} \hat{R}$: let's write down a partial derivative \hat{R} w.r.t. z_c

$$\frac{\partial \hat{R}}{\partial z_c} = \sum_{j=1}^C \frac{\partial \hat{R}}{\partial y_j} \frac{\partial y_j}{\partial z_c} =$$

Backpropagation Through Vector Activation

Now, let's get back to our problem

In this graph, $F(\cdot)$ is a **vector activation**. We know $\nabla_{\mathbf{y}} \hat{R}$



We want to find $\nabla_{\mathbf{z}} \hat{R}$: let's write down a partial derivative \hat{R} w.r.t. z_c


$$\frac{\partial \hat{R}}{\partial z_c} = \sum_{j=1}^C \frac{\partial \hat{R}}{\partial y_j} \frac{\partial y_j}{\partial z_c} = \underbrace{\begin{bmatrix} \frac{\partial y_1}{\partial z_c} & \dots & \frac{\partial y_C}{\partial z_c} \end{bmatrix}}_{\text{transpose of column } c \text{ of } \mathbf{J}_{\mathbf{z}\mathbf{y}} \equiv \text{row } c \text{ of } \mathbf{J}_{\mathbf{z}\mathbf{y}}^T} \nabla_{\mathbf{y}} \hat{R}$$

Backpropagation Through Vector Activation

So, the gradient of \hat{R} w.r.t. \mathbf{z} is given by

$$\nabla_{\mathbf{z}} \hat{R} \begin{bmatrix} \partial \hat{R} / \partial z_1 \\ \vdots \\ \partial \hat{R} / \partial z_C \end{bmatrix} = \begin{bmatrix} \text{row 1 of } \mathbf{J}_{\mathbf{z}}^T \nabla_{\mathbf{y}} \hat{R} \\ \vdots \\ \text{row } C \text{ of } \mathbf{J}_{\mathbf{z}}^T \nabla_{\mathbf{y}} \hat{R} \end{bmatrix} = (\mathbf{J}_{\mathbf{z}}^T \mathbf{y}) \nabla_{\mathbf{y}} \hat{R}$$

So, we can complete the computation graph as follows




Backpropagation Through Vector Activation

So, the gradient of \hat{R} w.r.t. \mathbf{z} is given by

$$\nabla_{\mathbf{z}} \hat{R} \begin{bmatrix} \partial \hat{R} / \partial z_1 \\ \vdots \\ \partial \hat{R} / \partial z_C \end{bmatrix} = \begin{bmatrix} \text{row 1 of } \mathbf{J}_{\mathbf{z}}^T \nabla_{\mathbf{y}} \hat{R} \\ \vdots \\ \text{row } C \text{ of } \mathbf{J}_{\mathbf{z}}^T \nabla_{\mathbf{y}} \hat{R} \end{bmatrix} = (\mathbf{J}_{\mathbf{z}}^T \mathbf{y}) \nabla_{\mathbf{y}} \hat{R}$$

So, we can complete the computation graph as follows




Backward Pass of Vector Activation

To pass backward on a vector activation, we use the **transpose** of its **Jacobian**

Backpropagation Through Vector Activation

How can we backpropagate through this neural network? **Let's complete**



- 1 Compute $\nabla_{y_3} \hat{R}_b = \nabla_{y_3} \mathcal{L}(y_3, 1_{v_b})$
- 2 Compute $\nabla_{z_3} \hat{R}_b = (J\text{Soft}_{\max})^T \nabla_{y_3} \hat{R}_b$
- 3 Compute $\nabla_{y_2} \hat{R}_b = W_3^T \nabla_{z_3} \hat{R}_b$
- 4 Remove entry at index 0 of $\nabla_{y_2} \hat{R}_b$ and compute $\nabla_{z_2} \hat{R}_b = \dot{r}(z_2) \odot \nabla_{y_2} \hat{R}_b$
- 5 Compute $\nabla_{y_1} \hat{R}_b = W_2^T \nabla_{z_2} \hat{R}_b$
- 6 Remove entry at index 0 of $\nabla_{y_1} \hat{R}_b$ and compute $\nabla_{z_1} \hat{R}_b = \dot{r}(z_1) \odot \nabla_{y_1} \hat{R}_b$

Multiclass Classification via FNN: *Training*

Let's now recall **gradient descent** for training of multiclass classifier

GradientDescent():

```

1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}\}$  and set a learning rate  $\eta$ 
2: while weights not converged do
3:   for  $b = 1, \dots, B$  do
4:      $\text{NN.values} \leftarrow \text{ForwardProp}(\mathbf{x}_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\})$ 
5:      $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow \text{BackProp}(\mathbf{x}_b, \mathbf{u}_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}, \text{NN.values})$ 
6:   end for
7:   Update

```

$$\mathbf{W}_1^{(t+1)} \leftarrow \mathbf{W}_1^{(t)} - \eta \text{ mean}(\mathbf{G}_{1,1}, \dots, \mathbf{G}_{1,B})$$

$$\mathbf{W}_2^{(t+1)} \leftarrow \mathbf{W}_2^{(t)} - \eta \text{ mean}(\mathbf{G}_{2,1}, \dots, \mathbf{G}_{2,B})$$

$$\mathbf{W}_3^{(t+1)} \leftarrow \mathbf{W}_3^{(t)} - \eta \text{ mean}(\mathbf{G}_{3,1}, \dots, \mathbf{G}_{3,B})$$

8: **end while**

We call this form of gradient descent **full-batch**

Full-Batch Training

Batch \equiv *the dataset reserved for training*

In full-batch training, we calculate the gradient for all data-points in the **batch**: so, we need to wait till forward and backward pass are over for all B data-points

This can be a huge burdensome!

Full-Batch Training

Batch \equiv the dataset reserved for *training*

In full-batch training, we calculate the gradient for all data-points in the **batch**: so, we need to wait till forward and backward pass are over for all B data-points

*This can be a **huge** burdensome!*

- + Wait a moment! Don't we use all the dataset for training?
- No! As you may have noticed in the assignments, we reserve a part of it for *testing*
- + And, why should it be a burdensome?
- OK! Let's get more into datasets!

Public Datasets

Let's consider our example of **image recognition**: we want to recognize the *hand-written number* in an *image*. For this, we need to have access to *images* of *hand-written numbers*. This has been done before by people at *National Institute of Standards and Technology* and collected in a **database** called

Modified National Institute of Standards and Technology (MNIST)

that is available for **public** on internet

Public Datasets

Let's consider our example of **image recognition**: we want to recognize the *hand-written number* in an *image*. For this, we need to have access to *images* of *hand-written numbers*. This has been done before by people at *National Institute of Standards and Technology* and collected in a **database** called

Modified National Institute of Standards and Technology (MNIST)

that is available for **public** on internet

There are several of such **public databases**; some well-known examples are

- CIFAR-10 and CIFAR-100 by Canadian Institute For Advanced Research
- ImageNet initiated by *Fei-Fei Li* at Princeton University
- Caltech-101 and Caltech-256 compiled at Caltech
- Fashion MNIST that collects fashion images and labels them

You can find out more about public datasets [online](#)

Public Datasets: *Accessing via PyTorch*

PyTorch provides us a simple tool to access these public datasets, e.g.,

```
>> import torchvision.datasets as DataSets  
  
>> dataset = DataSets.MNIST( ... )  
>> dataset = DataSets.CIFAR10( ... )
```

Public Datasets: Accessing via PyTorch

PyTorch provides us a simple tool to access these public datasets, e.g.,

```
>> import torchvision.datasets as DataSets  
  
>> dataset = DataSets.MNIST( ... )  
>> dataset = DataSets.CIFAR10( ... )
```

In the example of MNIST, we load the dataset which contains the **pixel vectors of the images of size 28×28** . This means that we load a *list of pairs* where each pair contains

a 784-dimensional vector of pixel values and a label that is in $\{0, 1, \dots, 9\}$

Public Datasets: *How Do They Look?*

Public datasets include a large amount of data-points with their labels

*MNIST includes 70,000 images of **hand-written numbers** with their **true labels**: from these 70,000 we use 60,000 for **training** and 10,000 for **test***

This means that once we load the MNIST dataset, we make a **batch** of **60,000 images** to **train** our FNN. Once the **training** is over, we **test** the performance of the trained FNN on the **remaining 10,000 images**

Public Datasets: *How Do They Look?*

Public datasets include a large amount of data-points with their labels

*MNIST includes 70,000 images of **hand-written numbers** with their **true labels**: from these 70,000 we use 60,000 for **training** and 10,000 for **test***

This means that once we load the MNIST dataset, we make a **batch** of **60,000 images** to **train** our FNN. Once the **training** is over, we **test** the performance of the trained FNN on the **remaining 10,000 images**

Back to our problem, this means that our **full-batch training** performs **each iteration** of the **gradient descent** **after**

60,000 forward and backward passes over the FNN

which sounds a lot!

Full-Batch Training: Complexity

Given the example of MNIST, let's see roughly how long it takes to do a full-batch training: *if we need 100 iterations of gradient descent, we need to pass back and forth for 6×10^6 times!*

- + *But do we really need to do this much? This sounds impossible in large NNs!*
- No! We really don't need! We can do the training much faster

Full-Batch Training: Complexity

Given the example of MNIST, let's see roughly how long it takes to do a full-batch training: *if we need 100 iterations of gradient descent, we need to pass back and forth for 6×10^6 times!*

- + *But do we really need to do this much? This sounds impossible in large NNs!*
- *No! We really don't need! We can do the training much faster*

The full-batch training is really not practical: in practice, we use stochastic (mini-batch) gradient descent to train our NN with feasible complexity

Let's take a look at these approaches!

Sample-Level Training

The most primary idea is to apply one step of **gradient descent** after **each forward and backward pass**: *in our FNN this means that we do the following*

Sample-Level Training

The most primary idea is to apply one step of **gradient descent** after **each forward and backward pass**: *in our FNN this means that we do the following*

`SymbLevel_GradientDescent()` :

```

1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}\}$  and set a learning rate  $\eta$ 
2: Start at  $b = 1$ 
3: while weights not converged do
4:   if  $b > B$  then
5:     Update  $b \leftarrow 1$                                 # start over with the dataset
6:   end if
7:    $\text{NN.values} \leftarrow \text{ForwardProp}(\mathbf{x}_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\})$ 
8:    $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow \text{BackProp}(\mathbf{x}_b, v_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}, \text{NN.values})$ 
9:   Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \mathbf{G}_{\ell,b}$           # symbol_level update
10:  Update  $b \leftarrow b + 1$                                 # go for next data-point
11: end while

```

We call this approach *symbol-level training*

Sample-Level Training: *Meaning*

- + *But what does it mean in the sense of empirical risk minimization? Aren't we now doing something different from the standard **gradient descent**?!*
- Yes! We are in fact performing an approximative **gradient descent**

Sample-Level Training: *Meaning*

- + *But what does it mean in the sense of empirical risk minimization? Aren't we now doing something different from the standard **gradient descent**?!*
- Yes! We are in fact performing an approximative **gradient descent**

Consider the an **ideal scenario** in which

$$\mathbf{G}_{\ell,1} = \mathbf{G}_{\ell,2} = \dots = \mathbf{G}_{\ell,B}$$

In this case, we do not need to wait for the **batch** to be fully over, since

$$\mathbf{G}_{\ell,1} = \text{mean} \left(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B} \right)$$

In other words, in this case

symbol-level training \equiv **full-batch** training

Sample-Level Training: *Meaning*

*In practice, at each data-point we calculate a **noisy-version** of **full-batch** gradient^a. In other words, we can think of $G_{\ell,b}$ for each b as*

$$G_{\ell,b} = \text{mean}(G_{\ell,1}, G_{\ell,2}, \dots, G_{\ell,B}) + \text{Noise}$$

^a**Full-batch** gradient means the average of gradients over the whole **batch**.

Sample-Level Training: Meaning

*In practice, at each data-point we calculate a **noisy-version** of **full-batch** gradient^a. In other words, we can think of $G_{\ell,b}$ for each b as*

$$G_{\ell,b} = \text{mean}(G_{\ell,1}, G_{\ell,2}, \dots, G_{\ell,B}) + \text{Noise}$$

^a**Full-batch** gradient means the average of gradients over the whole **batch**.

If this noise is small enough, we can say that

$$G_{\ell,b} \approx \text{mean}(G_{\ell,1}, G_{\ell,2}, \dots, G_{\ell,B})$$

and therefore, we can conclude that

symbol-level training \approx **full-batch** training

*In this case, we say that $G_{\ell,b}$ is an **estimator** of the **full-batch gradient***

Sample-Level Training: *Repetitive Cycle Issue*

Naive sample-level update can trap us into a **repetitive cycle**: *in simple words, we can end up with our initial point at the end of the batch.*

Sample-Level Training: Repetitive Cycle Issue

Naive sample-level update can trap us into a **repetitive cycle**: in simple words, we can end up with our initial point at the end of the batch. For instance, consider the following dummy (but possible) scenario in our three-layer FNN

We start with $\mathbf{W}_\ell^{(0)}$ and get into the **batch for the first time**

- We update $\mathbf{W}_\ell^{(0)}$ after the **first data-point** to $\mathbf{W}_\ell^{(1)}$
- We update $\mathbf{W}_\ell^{(1)}$ after the **second data-point** to $\mathbf{W}_\ell^{(2)}$
- ...
- We update $\mathbf{W}_\ell^{(B-1)}$ after the **last data-point** to $\mathbf{W}_\ell^{(B)}$

Now, assume that $\mathbf{W}_\ell^{(B)} = \mathbf{W}_\ell^{(0)}$ for all layers again!

Sample-Level Training: Repetitive Cycle Issue

Naive sample-level update can trap us into a **repetitive cycle**: in simple words, we can end up with our initial point at the end of the batch. For instance, consider the following dummy (but possible) scenario in our three-layer FNN

We start with $\mathbf{W}_\ell^{(0)}$ and get into the **batch for the first time**

- We update $\mathbf{W}_\ell^{(0)}$ after the **first data-point** to $\mathbf{W}_\ell^{(1)}$
- We update $\mathbf{W}_\ell^{(1)}$ after the **second data-point** to $\mathbf{W}_\ell^{(2)}$
- ...
- We update $\mathbf{W}_\ell^{(B-1)}$ after the **last data-point** to $\mathbf{W}_\ell^{(B)}$

Now, assume that $\mathbf{W}_\ell^{(B)} = \mathbf{W}_\ell^{(0)}$ for all layers again!

In the above dummy example, **further looping over the batch** is **useless**, since we always get back to the initial point: this is the most basic example of the **repetitive cycle issue**

Stochastic Sample-Level Training: SGD

- + How can we avoid such *cyclic behaviors*?
- We can use *Stochastic Gradient Descent (SGD)*

Stochastic Sample-Level Training: SGD

- + How can we avoid such *cyclic behaviors*?
- We can use *Stochastic Gradient Descent (SGD)*

Each time we are to loop over our *training batch*, we *shuffle* the data-points *randomly*: this way we avoid next loop behave like the previous one

This idea is called Stochastic Gradient Descent (SGD)

SGD is the most common algorithm for training of NNs!

Stochastic Sample-Level Training: SGD

- + How can we avoid such **cyclic behaviors**?
- We can use **Stochastic Gradient Descent (SGD)**

Each time we are to loop over our **training batch**, we **shuffle** the data-points **randomly**: this way we avoid next loop behave like the previous one

This idea is called **Stochastic Gradient Descent (SGD)**

SGD is the most common algorithm for training of NNs!

What does **random shuffling** mean?

It means **randomly permuting** the data-points

Stochastic Gradient Descent

SGD() :

```

1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}\}$  and set a learning rate  $\eta$ 
2: Randomly shuffle the batch and start at  $b = 1$ 
3: while weights not converged do
4:   if  $b > B$  then
5:     Randomly shuffle the batch and set  $b \leftarrow 1$            # random shuffling
6:   end if
7:   NN.values  $\leftarrow$  ForwardProp ( $\mathbf{x}_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}$ )
8:    $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow$  BackProp ( $\mathbf{x}_b, v_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}, \text{NN.values}$ )
9:   Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \mathbf{G}_{\ell,b}$            # symbol_level update
10:  Update  $b \leftarrow b + 1$                                    # go for next data-point
11: end while

```

Stochastic Gradient Descent

SGD() :

```

1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}\}$  and set a learning rate  $\eta$ 
2: Randomly shuffle the batch and start at  $b = 1$ 
3: while weights not converged do
4:   if  $b > B$  then
5:     Randomly shuffle the batch and set  $b \leftarrow 1$            # random shuffling
6:   end if
7:   NN.values  $\leftarrow$  ForwardProp ( $x_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}$ )
8:    $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow$  BackProp ( $x_b, v_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}, \text{NN.values}$ )
9:   Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \mathbf{G}_{\ell,b}$            # symbol_level update
10:  Update  $b \leftarrow b + 1$                                    # go for next data-point
11: end while

```

+ But, doesn't sample-level training lead to any drawback?

- Sure! But we accept this drawback as a cost we pay for less complexity

Let's see how this trade-off looks like

Recap: Variance

For random variable x with mean μ , the **variance** is defined as

$$\text{Var}\{x\} = \mathbb{E}\{(x - \mu)^2\} = \mathbb{E}\{x^2\} - \mu^2$$

Clearly, when x is **zero-mean**, we can say $\text{Var}\{x\} = \mathbb{E}\{x^2\}$

Recap: Variance

For random variable x with mean μ , the **variance** is defined as

$$\text{Var}\{x\} = \mathbb{E}\{(x - \mu)^2\} = \mathbb{E}\{x^2\} - \mu^2$$

Clearly, when x is **zero-mean**, we can say $\text{Var}\{x\} = \mathbb{E}\{x^2\}$

Properties of Variance

For any random variable x and constant c , we have

$$\text{Var}\{cx\} = c^2 \text{Var}\{x\}$$

Let x_1, \dots, x_N be N **independent** random variables; then, we have

$$\text{Var}\left\{\sum_{n=1}^N x_n\right\} = \sum_{n=1}^N \text{Var}\{x_n\}$$

Recap: Variance

Now, assume x_1, \dots, x_N are N independent zero-mean random variables all with variance σ^2 : let \bar{x} be the arithmetic average of x_1, \dots, x_N , i.e.,

$$\bar{x} = \text{mean}(x_1, \dots, x_N) = \frac{1}{N} \sum_{n=1}^N x_n$$

We could then say

$$\text{Var}\{\bar{x}\} = \text{Var}\left\{\frac{1}{N} \sum_{n=1}^N x_n\right\}$$

Recap: Variance

Now, assume x_1, \dots, x_N are N independent zero-mean random variables all with variance σ^2 : let \bar{x} be the arithmetic average of x_1, \dots, x_N , i.e.,

$$\bar{x} = \text{mean}(x_1, \dots, x_N) = \frac{1}{N} \sum_{n=1}^N x_n$$

We could then say

$$\text{Var}\{\bar{x}\} = \text{Var}\left\{\frac{1}{N} \sum_{n=1}^N x_n\right\} = \frac{1}{N^2} \text{Var}\left\{\sum_{n=1}^N x_n\right\}$$

Recap: Variance

Now, assume x_1, \dots, x_N are N **independent zero-mean** random variables all with variance σ^2 : let \bar{x} be the **arithmetic average** of x_1, \dots, x_N , i.e.,

$$\bar{x} = \text{mean}(x_1, \dots, x_N) = \frac{1}{N} \sum_{n=1}^N x_n$$

We could then say

$$\begin{aligned} \text{Var}\{\bar{x}\} &= \text{Var}\left\{\frac{1}{N} \sum_{n=1}^N x_n\right\} = \frac{1}{N^2} \text{Var}\left\{\sum_{n=1}^N x_n\right\} \\ &= \frac{1}{N^2} \sum_{n=1}^N \underbrace{\text{Var}\{x_n\}}_{\sigma^2} \end{aligned}$$

Recap: Variance

Now, assume x_1, \dots, x_N are N independent zero-mean random variables all with variance σ^2 : let \bar{x} be the arithmetic average of x_1, \dots, x_N , i.e.,

$$\bar{x} = \text{mean}(x_1, \dots, x_N) = \frac{1}{N} \sum_{n=1}^N x_n$$

We could then say

$$\begin{aligned} \text{Var}\{\bar{x}\} &= \text{Var}\left\{\frac{1}{N} \sum_{n=1}^N x_n\right\} = \frac{1}{N^2} \text{Var}\left\{\sum_{n=1}^N x_n\right\} \\ &= \frac{1}{N^2} \sum_{n=1}^N \underbrace{\text{Var}\{x_n\}}_{\sigma^2} = \frac{1}{N^2} (N\sigma^2) \\ &= \frac{\sigma^2}{N} \end{aligned}$$

Recap: Variance

Now, assume x_1, \dots, x_N are N **independent zero-mean** random variables all with variance σ^2 : let \bar{x} be the **arithmetic average** of x_1, \dots, x_N , i.e.,

$$\bar{x} = \text{mean}(x_1, \dots, x_N) = \frac{1}{N} \sum_{n=1}^N x_n$$

We could then say

$$\begin{aligned} \text{Var}\{\bar{x}\} &= \text{Var}\left\{\frac{1}{N} \sum_{n=1}^N x_n\right\} = \frac{1}{N^2} \text{Var}\left\{\sum_{n=1}^N x_n\right\} \\ &= \frac{1}{N^2} \sum_{n=1}^N \underbrace{\text{Var}\{x_n\}}_{\sigma^2} = \frac{1}{N^2} (N\sigma^2) \\ &= \frac{\sigma^2}{N} \quad \text{variance of average drops by } 1/\# \text{ samples} \end{aligned}$$

Complexity-Accuracy Trade-off of SGD

Now, let's get back to our problem: *when we talked about the meaning of symbol level update, we said*

*In practice, at each data-point we calculate a **noisy-version** of **full-batch** gradient. In other words, we can think of $G_{\ell,b}$ for each b as*

$$G_{\ell,b} = \text{mean}(G_{\ell,1}, G_{\ell,2}, \dots, G_{\ell,B}) + \text{Noise}$$

*and called $G_{\ell,b}$ an **estimator** of the **mean**.*

Complexity-Accuracy Trade-off of SGD

Now, let's get back to our problem: *when we talked about the meaning of symbol level update, we said*

*In practice, at each data-point we calculate a **noisy-version** of **full-batch** gradient. In other words, we can think of $\mathbf{G}_{\ell,b}$ for each b as*

$$\mathbf{G}_{\ell,b} = \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B}) + \text{Noise}$$

and called $\mathbf{G}_{\ell,b}$ an **estimator** of the **mean**.

Let's make the above statement a bit more formal: we assume that Noise for each b is a matrix with **independent zero-mean** entries all with **variance σ^2** , i.e.,

$$\mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b}$$

where we define $\bar{\mathbf{G}}_{\ell}$ to be the gradient of the **true risk**

Complexity-Accuracy Trade-off of SGD

What is the *true risk*? If you remember, when we started with training

Our goal was to *minimize* the *risk* $R(\mathbf{w})$

However, we could not do this: since we *did not know* (1) the *true function*, and (2) the *data distribution*. Thus,

we approximated the *true risk* $R(\mathbf{w})$ with *empirical risk* $\hat{R}(\mathbf{w})$

Complexity-Accuracy Trade-off of SGD

What is the *true risk*? If you remember, when we started with training

Our goal was to *minimize* the *risk* $R(\mathbf{w})$

However, we could not do this: since we *did not know* (1) the *true function*, and (2) the *data distribution*. Thus,

we approximated the *true risk* $R(\mathbf{w})$ with *empirical risk* $\hat{R}(\mathbf{w})$

We assume that $\bar{\mathbf{G}}_\ell$ is the gradient of *true risk* with respect to \mathbf{W}_ℓ , i.e.,

$$\bar{\mathbf{G}}_\ell = \nabla_{\mathbf{W}_\ell} R(\mathbf{w})$$

+ Can we determine this gradient?

- **Of course not!** We can only approximate it with $\nabla_{\mathbf{W}_\ell} \hat{R}(\mathbf{w})$

Complexity-Accuracy Trade-off of SGD

Let's see accurate *the gradient* is *approximated*, when we do *full-batch* training

In *full-batch* training, we determine the gradient as

$$\hat{\mathbf{G}}_{\ell}^{\text{batch}} = \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B}) = \frac{1}{B} \sum_{b=1}^B \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b}$$

Complexity-Accuracy Trade-off of SGD

Let's see accurate *the gradient* is *approximated*, when we do *full-batch* training

In *full-batch* training, we determine the gradient as

$$\begin{aligned}\hat{\mathbf{G}}_{\ell}^{\text{batch}} &= \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B}) = \frac{1}{B} \sum_{b=1}^B \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b} \\ &= \bar{\mathbf{G}}_{\ell} + \underbrace{\frac{1}{B} \sum_{b=1}^B \mathbf{N}_{\ell,b}}_{\hat{\mathbf{N}}_{\ell}^{\text{batch}}}\end{aligned}$$

Complexity-Accuracy Trade-off of SGD

Let's see accurate *the gradient* is *approximated*, when we do *full-batch* training

In *full-batch* training, we determine the gradient as

$$\begin{aligned}\hat{\mathbf{G}}_{\ell}^{\text{batch}} &= \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B}) = \frac{1}{B} \sum_{b=1}^B \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b} \\ &= \bar{\mathbf{G}}_{\ell} + \underbrace{\frac{1}{B} \sum_{b=1}^B \mathbf{N}_{\ell,b}}_{\hat{\mathbf{N}}_{\ell}^{\text{batch}}} = \bar{\mathbf{G}}_{\ell} + \hat{\mathbf{N}}_{\ell}^{\text{batch}}\end{aligned}$$

Complexity-Accuracy Trade-off of SGD

Let's see accurate *the gradient* is *approximated*, when we do *full-batch* training

In *full-batch* training, we determine the gradient as

$$\begin{aligned}\hat{\mathbf{G}}_{\ell}^{\text{batch}} &= \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B}) = \frac{1}{B} \sum_{b=1}^B \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b} \\ &= \bar{\mathbf{G}}_{\ell} + \underbrace{\frac{1}{B} \sum_{b=1}^B \mathbf{N}_{\ell,b}}_{\hat{\mathbf{N}}_{\ell}^{\text{batch}}} = \bar{\mathbf{G}}_{\ell} + \hat{\mathbf{N}}_{\ell}^{\text{batch}}\end{aligned}$$

Recall that by *arithmetic averaging variance drops by 1/# of samples*

In *full-batch* training the *approximated gradient* is different from the *true gradient* by an error whose variance *drops as* σ^2/B

Complexity-Accuracy Trade-off of SGD

In *full-batch* training the *approximated gradient* is different from the *true gradient* by an error whose variance *drops as* σ^2/B

Now, let's compare it to SGD

In SGD, we *approximate* the gradient with a sample gradient, i.e.,

$$\hat{\mathbf{G}}_{\ell}^{\text{SGD}} = \mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b}$$

Complexity-Accuracy Trade-off of SGD

In *full-batch* training the *approximated gradient* is different from the *true gradient* by an error whose variance *drops as* σ^2/B

Now, let's compare it to SGD

In SGD, we *approximate* the gradient with a sample gradient, i.e.,

$$\hat{\mathbf{G}}_{\ell}^{\text{SGD}} = \mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b}$$

As we see, SGD is still *approximating* the *true gradient* but with much *larger variance*: entries of $\mathbf{N}_{\ell,b}$ have all *variance* σ^2

For instance, consider MNIST with 60,000 samples: by *full-batch* training we get gradient values whose difference from the entries of the *true gradient* is approximately 1.67×10^{-5} times smaller than those gradient entries calculated by *SGD*!

Complexity-Accuracy Trade-off of SGD

In the context of ML, we often say: *in the analyses of last slides,*

SGD and full-batch training are both unbiased estimators of \bar{G}_ℓ

We call them unbiased, since $\mathbb{E} \left\{ \hat{G}_\ell^{\text{SGD}} \right\} = \mathbb{E} \left\{ \hat{G}_\ell^{\text{batch}} \right\} = \bar{G}_\ell$

Complexity-Accuracy Trade-off of SGD

In the context of ML, we often say: *in the analyses of last slides,*

SGD and **full-batch** training are both **unbiased estimators of $\bar{\mathbf{G}}_\ell$**

We call them **unbiased**, since $\mathbb{E} \left\{ \hat{\mathbf{G}}_\ell^{\text{SGD}} \right\} = \mathbb{E} \left\{ \hat{\mathbf{G}}_\ell^{\text{batch}} \right\} = \bar{\mathbf{G}}_\ell$

Complexity-Accuracy Trade-off

Assume that a forward and backward pass takes time T and that gradient of the risk at each sample be an **unbiased estimators of the true gradient**; then,

- 1 each step of **SGD** takes time T while each step of **full-batch** training takes BT with B being the **batch size**
- 2 if we denote the variance of **estimation given by SGD** by σ^2 , the variance of **full-batch** estimator is σ^2/B

Training via Mini-Batches

- + But, can't we *play with this trade-off?* For instance, *increase* a bit the *complexity* to *improve* the *accuracy*!
- Yes! This is the idea of *mini-batch training*

Training via Mini-Batches

- + But, can't we *play with this trade-off*? For instance, *increase* a bit the *complexity* to *improve* the *accuracy*!
- Yes! This is the idea of *mini-batch training*

In *mini-batch training*, we divide *the whole batch of data* into *mini-batches*:

- after each *mini-batch* is over, we *average the gradients* over the *mini-batch*
- we apply *one step of gradient descent* using this *averaged gradient*

To avoid *cyclic behavior*, we still *shuffle* the dataset *randomly* each time we *start a new loop* over it. This training approach is hence often called

Mini-Batch Stochastic Gradient Descent = Mini-Batch SGD

Mini-Batch SGD

mBatchSGD() :

```

1: Initiate with some initial values  $\{\mathbf{W}_\ell^{(0)}\}$  and set a learning rate  $\eta$ 
2: Randomly shuffle the batch and divide it into mini-batches of size  $\Omega$ 
3: Denote the number of mini-batches by  $\Xi = \lceil B/\Omega \rceil$  and start at  $\xi = 1$ 
4: while weights not converged do
5:   if  $\xi > \Xi$  then
6:     Randomly shuffle the batch and divide it into mini-batches of size  $\Omega$ 
7:     Set  $\xi \leftarrow 1$                                 # start over with the dataset
8:   end if
9:   for  $\omega = 1, \dots, \Omega$  do
10:    NN.values  $\leftarrow$  ForwardProp ( $\mathbf{x}_\omega, \{\mathbf{W}_\ell^{(t)}\}$ )
11:     $\{\mathbf{G}_{\ell,\omega}\} \leftarrow$  BackProp ( $\mathbf{x}_\omega, v_\omega, \{\mathbf{W}_\ell^{(t)}\}, \text{NN.values}$ )
12:  end for
13:  Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \text{ mean}(\mathbf{G}_{\ell,1}, \dots, \mathbf{G}_{\ell,\Omega})$ 
14:  Update  $\xi \leftarrow \xi + 1$                                 # go for next mini-batch
15: end while

```

Complexity-Accuracy Trade-off

It is easy to see that

- *mini-batch training* reduces to *full-batch training* when we set the size of *mini-batches* to B , i.e., $\Omega = B$
- *mini-batch training* reduces to *SGD* when we set the size of *mini-batches* to 1 , i.e., $\Omega = 1$


Complexity-Accuracy Trade-off

Assume that a forward and backward pass takes time T and that gradient of the risk at each sample be an *unbiased estimators of the true gradient*; then,

- ① each step of *SGD* takes time T while each step of *mini-batch* training takes ΩT with Ω being the *mini-batch size*
- ② if we denote the variance of *estimation given by SGD* by σ^2 , the variance of *mini-batch* estimator is σ^2/Ω


Complexity-Accuracy Trade-off

The complete trade-off can be visualized as




Complexity-Accuracy Trade-off

The complete trade-off can be visualized as




Complexity-Accuracy Trade-off

The complete trade-off can be visualized as




Complexity-Accuracy Trade-off

The complete trade-off can be visualized as



Complexity-Accuracy Trade-off

The complete trade-off can be visualized as



Mini-batch size is what specifies the *trade-off point*

Few Definitions: *Epoch and Iteration*

In the language of deep learning there are few terms that we must know

Batch Size

Through time, the term *mini-batch* has been transformed to *batch*, and the *complete batch* is referred to as *training dataset* or the *full batch*. People hence call the size of each *mini-batch*, i.e., Ω , the *batch size*

Iteration

When we take one step of gradient descent, we take one *iteration*. So, one *iteration* is over when we finish with a *mini-batch*

Epoch

An *epoch* is over when we finish *once* with the *whole training dataset*

Few Definitions: *Epoch and Iteration*

We can annotate these definitions in our algorithm

mBatchSGD() :

```

1: Initiate with some initial values  $\{\mathbf{W}_\ell^{(0)}\}$  and set a learning rate  $\eta$ 
2: Randomly shuffle training dataset and make mini-batches of size  $\Omega \equiv \text{batch-size}$ 
3: Denote the number of mini-batches by  $\Xi = \lceil B/\Omega \rceil$  and start at  $\xi = 1$ 
4: while weights not converged do
5:   if  $\xi > \Xi$  then
6:     Randomly shuffle the batch and divide it into mini-batches of size  $\Omega$ 
7:     Set  $\xi \leftarrow 1$  ← one epoch is over, we start another epoch
8:   end if
9:   for  $\omega = 1, \dots, \Omega$  do
10:    NN.values  $\leftarrow$  ForwardProp ( $\mathbf{x}_\omega, \{\mathbf{W}_\ell^{(t)}\}$ ) going through a min-batch
11:     $\{\mathbf{G}_{\ell, \omega}\} \leftarrow$  BackProp ( $\mathbf{x}_\omega, v_\omega, \{\mathbf{W}_\ell^{(t)}\}, \text{NN.values}$ )
12:  end for
13:  Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \text{ mean}(\mathbf{G}_{\ell, 1}, \dots, \mathbf{G}_{\ell, \Omega})$  ← one iteration
14:  Update  $\xi \leftarrow \xi + 1$  # go for next mini-batch
15: end while

```

Few Definitions: *Epoch and Iteration*

We can consider a simple example: say we train our FNN over MNIST using *mini-batch SGD* with *batch size* $\Omega = 100$. Our *training dataset* has 60,000 data-points; thus, we have

$$\Xi = \frac{60,000}{100} = 600$$

mini-batches. Each time we finish with a *mini-batch*, we do *one iteration* of gradient descent. After *600 iterations*, we finish with a single *epoch*

So, if we have trained the FNN for 10 *epochs*, it means that

we have done $600 \times 10 = 6000$ *iterations of gradient descent*

Testing NNs with New Data-Point

- + Say we are over with the **training**; then, what should we do?
- We need to test it with the data we reserved for **testing**

After **training**, we need to test our **trained** NN: say we get a new data-point \mathbf{x}_{new} with label \mathbf{v}_{new} . We can test our NN for this new **test data-point** by evaluating **classical metrics**

- 1 **Test Risk** also called **Test Loss**: we pass \mathbf{x}_{new} forward through our **trained NN** and get \mathbf{y}_{new} . We then calculate the **test loss** as $\mathcal{L}(\mathbf{y}_{\text{new}}, \mathbf{v}_{\text{new}})$ using the same loss function \mathcal{L} we used for training
- 2 **Test Accuracy**: we use \mathbf{y}_{new} to classify \mathbf{x}_{new} . We then compare it to the **true class** of \mathbf{x}_{new} . If they are the same; then, the test accuracy is 1, if not, it is 0

Testing NNs over Test Dataset

Testing for a single new point is not reliable: *this is why we had reserved the test dataset.*

Given the **test dataset**, we go through every single test **data-point**

- we pass the data-point forward through the **trained NN**
- we compute the **test loss** and **test accuracy**
- we **average them** over the whole **test dataset**

Therefore, we get

- an **average** loss that **approximates** the risk
- a **test accuracy** between 0 and 1 that says how **accurate** our trained NN is

Learning Curves

- + What you said gives us *two numbers!* But, I have seen *curves!*
- Yes! They are *learning curves*

In practice, the SGD can take *very long* to converge, i.e., to stop iterating
it needs too many iterations to get *too close* to the minimum

But, it might be *not really needed* to get *that close!* So,
we test our NN once every epoch


We then plot the *test risk* and *test accuracy* against *number of epochs* in a curve: these curves are often called *learning curves*

if we see that learning curves are saturating, we can stop the training

In practice: *we always perform the training for a fixed number of epochs*

Learning Curves: *Example*

Let's see an example: recall our three-layer FNN. Say, we train it for image classification over MNIST which has 60,000 data-points for and 10,000 for test



In MNIST, we have 10 classes, so $C = 10$. We use *cross-entropy* as *loss function*

Learning Curves: *Example*

We agree to do the following: we use *mini-batch SGD* with *batch size* $\Omega = 100$ and train the FNN for *100 epochs*.

In epoch $\xi = 1, \dots, 100$

Learning Curves: *Example*

We agree to do the following: we use *mini-batch SGD* with *batch size* $\Omega = 100$ and train the FNN for *100 epochs*.

In *epoch* $\xi = 1, \dots, 100$

- 1 we perform 600 iterations of *gradient descent*

Learning Curves: *Example*

We agree to do the following: we use *mini-batch SGD* with *batch size* $\Omega = 100$ and train the FNN for *100 epochs*.

In *epoch* $\xi = 1, \dots, 100$

- 1 we perform 600 iterations of *gradient descent*
- 2 we fix the weights to what we computed at *last iteration* of *the epoch*

Learning Curves: *Example*

We agree to do the following: we use *mini-batch SGD* with *batch size* $\Omega = 100$ and train the FNN for *100 epochs*.

In *epoch* $\xi = 1, \dots, 100$

- ① we perform 600 iterations of *gradient descent*
- ② we fix the weights to what we computed at *last iteration* of *the epoch*
- ③ for each test data-point: we *pass it forward* and determine \mathbf{y}_3
 - ① we compute CE ($\mathbf{y}_3, \mathbf{1}_v$), where v is the *true class* of test data-point
 - ② we find the index of *maximum term in* \mathbf{y}_3 and compare it to v
 - ↳ if they are the same, we set *accuracy* to 1; otherwise, we *set it 0*

Learning Curves: *Example*

We agree to do the following: we use *mini-batch SGD* with *batch size* $\Omega = 100$ and train the FNN for *100 epochs*.

In *epoch* $\xi = 1, \dots, 100$

- ① we perform 600 iterations of *gradient descent*
- ② we fix the weights to what we computed at *last iteration* of *the epoch*
- ③ for each test data-point: we *pass it forward* and determine \mathbf{y}_3
 - ① we compute CE ($\mathbf{y}_3, \mathbf{1}_v$), where v is the *true class* of test data-point
 - ② we find the index of *maximum term in* \mathbf{y}_3 and compare it to v
 - ↳ if they are the same, we set *accuracy* to 1; otherwise, we *set it* 0
- ④ we average test loss and accuracy

Learning Curves: *Example*

We agree to do the following: we use *mini-batch SGD* with *batch size* $\Omega = 100$ and train the FNN for *100 epochs*.

In *epoch* $\xi = 1, \dots, 100$


- ① we perform 600 iterations of *gradient descent*
- ② we fix the weights to what we computed at *last iteration* of *the epoch*
- ③ for each test data-point: we *pass it forward* and determine \mathbf{y}_3
 - ① we compute $\text{CE}(\mathbf{y}_3, \mathbf{1}_v)$, where v is the *true class* of test data-point
 - ② we find the index of *maximum term in* \mathbf{y}_3 and compare it to v
 - ↳ if they are the same, we set *accuracy* to 1; otherwise, we *set it 0*
- ④ we average test loss and accuracy

Now, for each epoch

we have a *test loss* and *test accuracy*: we plot them against ξ


Learning Curves: *Example*

How should the learning curves look? A typical curve for test loss is



Learning Curves: *Example*

How should the learning curves look? A typical curve for test accuracy is



Summary of This Chapter

- To train a NN we need *gradients*
 - ↳ We can *calculate gradient* by *forward* and *backpropagation* over the NN
 - ↳ In FNNs, *forward propagation* uses simple linear and nonlinear operations
 - ↳ *Backpropagation* is readily derived using *computation graph*

Summary of This Chapter

- To train a NN we need *gradients*
 - ↳ We can *calculate gradient* by *forward* and *backpropagation* over the NN
 - ↳ In FNNs, *forward propagation* uses simple linear and nonlinear operations
 - ↳ *Backpropagation* is readily derived using *computation graph*
- We tried Classification via FNNs
 - ↳ Better to work with *probabilities* instead of exact *labels*
 - ↳ For *multiclass* classification, we should use *vector-activated neurons*

Summary of This Chapter

- To train a NN we need **gradients**
 - ↳ We can **calculate gradient** by **forward** and **backpropagation** over the NN
 - ↳ In FNNs, **forward propagation** uses simple linear and nonlinear operations
 - ↳ **Backpropagation** is readily derived using **computation graph**
- We tried Classification via FNNs
 - ↳ Better to work with **probabilities** instead of exact **labels**
 - ↳ For **multiclass** classification, we should use **vector-activated neurons**
- To minimize the **exact** empirical risk, we have to do **full-batch** training
 - ↳ This requires **huge computation complexity**
 - ↳ We can hugely reduce this cost by **SGD** which does **sample-level training**
 - ↳ **SGD** versus **full-batch** describes a **complexity-accuracy trade-off**
 - ↳ We can tune this **trade-off** by **mini-batch SGD**