

ECE 1508S2: Applied Deep Learning

Chapter 6: Recurrent NNs

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Winter 2024

Learning from Sequence Data

In many applications, we have *sequence data*, e.g.,

- *speech data* which is usually a long *time series*
- *text data* which is *sequence of words* and letters
- *financial data* that is typically a *time-dependent sequence of values*

Learning from Sequence Data

In many applications, we have *sequence data*, e.g.,

- *speech data* which is usually a long *time series*
- *text data* which is *sequence of words* and letters
- *financial data* that is typically a *time-dependent sequence of values*

Learning from such data can be *inefficient* via FNNs, i.e., MLPs and CNNs

- On one hand, we have *long* sequence
 - ↳ This can easily make the NN size *infeasible*

Learning from Sequence Data

In many applications, we have *sequence data*, e.g.,

- *speech data* which is usually a long *time series*
- *text data* which is *sequence of words* and letters
- *financial data* that is typically a *time-dependent sequence of values*

Learning from such data can be *inefficient* via FNNs, i.e., MLPs and CNNs

- On one hand, we have *long* sequence
 - ↳ This can easily make the NN size *infeasible*
- On another hand, we do *not* have *so much features*
 - ↳ Just think of a *long* text, where we need to *predict the next word in it*

Learning from Sequence Data

In many applications, we have *sequence data*, e.g.,

- *speech data* which is usually a long *time series*
- *text data* which is *sequence of words* and letters
- *financial data* that is typically a *time-dependent sequence of values*

Learning from such data can be *inefficient* via FNNs, i.e., MLPs and CNNs

- On one hand, we have *long* sequence
 - ↳ This can easily make the NN size *infeasible*
- On another hand, we do *not* have *so much features*
 - ↳ Just think of a *long* text, where we need to *predict the next word in it*

We need to develop some techniques to handle such data

First, let's see some examples!

Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute *talk*: we want to find out whether it is about *sport* or *science*

Learning from Sequence Data: *Example I*


Assume we listen to a 15-minute *talk*: we want to find out whether it is about *sport* or *science*

- *This is a classification problem*
 - ↳ *The data-point, i.e., **talk** is classified either as *sport* or *science**

Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute *talk*: we want to find out whether it is about *sport* or *science*


- This is a classification problem
 - ↳ The data-point, i.e., *talk* is classified either as *sport* or *science*
- What about the data?
 - ↳ We *sample* the audio signal at rate 44.1 kHz and quantize the samples



Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute *talk*: we want to find out whether it is about *sport* or *science*


- This is a classification problem
 - ↳ The data-point, i.e., *talk* is classified either as *sport* or *science*
- What about the data?
 - ↳ We *sample* the audio signal at rate 44.1 kHz and quantize the samples
 - ↳ We put every N successive samples in a *frames* \equiv *vector of samples*



Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**


- This is a classification problem
 - ↳ The data-point, i.e., **talk** is classified either as **sport** or **science**
- What about the data?
 - ↳ We **sample** the audio signal at rate 44.1 kHz and quantize the samples
 - ↳ We put every N successive samples in a **frames** \equiv **vector of samples**
 - ↳ We store the 15-minute talk as a sequence of time **frames**



Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a classification problem
 - ↳ The data-point, i.e., **talk** is classified either as **sport** or **science**
- What about the data?
 - ↳ We **sample** the audio signal at rate 44.1 kHz and quantize the samples
 - ↳ We put every N successive samples in a **frames** \equiv **vector of samples**
 - ↳ We store the 15-minute talk as a sequence of time **frames**
 - ↳ We may also store frequency frames from the Fourier transform



Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute *talk*: we want to find out whether it is about *sport* or *science*

- This is a *binary* classification problem
- What about the data? each data-point is a *sequence* of *vectors*

Learning from Sequence Data: *Example I*

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a **binary** classification problem
- What about the data? each data-point is a **sequence** of **vectors**

Say we make frames of 512 samples; then, we roughly have

$$77,587 \text{ frames} = 39,690,000 \text{ samples}$$

Learning from Sequence Data: Example I

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a **binary** classification problem
- What about the data? each data-point is a **sequence** of **vectors**

Say we make frames of 512 samples; then, we roughly have

$$77,587 \text{ frames} = 39,690,000 \text{ samples}$$

But do we need to pass them all together through an NN?

- It does **not** seem to be!

Learning from Sequence Data: Example I

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a **binary** classification problem
- What about the data? each data-point is a **sequence** of **vectors**

Say we make frames of 512 samples; then, we roughly have

$$77,587 \text{ frames} = 39,690,000 \text{ samples}$$

But do we need to pass them all together through an NN?

- It does **not** seem to be!
 - ↳ We can classify based of **simple words** and **expressions** in the talk
 - ↳ Processing all samples together seems to be an **unnecessary hardness**

Learning from Sequence Data: Example I

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a **binary** classification problem
- What about the data? each data-point is a **sequence** of **vectors**

Say we make frames of 512 samples; then, we roughly have

$$77,587 \text{ frames} = 39,690,000 \text{ samples}$$

But do we need to pass them all together through an NN?

- It does **not** seem to be!
 - ↳ We can classify based of **simple words** and **expressions** in the talk
 - ↳ Processing all samples together seems to be an **unnecessary hardness**
- It does **not generalize** well
 - ↳ We want to classify **shorter and longer talks** as well

Learning from Sequence Data: *Example II*

Now let's consider **another example**: we have a **long text** and want to learn what is the **next word** in the sentence

- This is a **prediction** task
 - ↳ Given **previous text** we **predict** the **next outcome**

Learning from Sequence Data: *Example II*

Now let's consider **another example**: we have a **long text** and want to learn what is the **next word** in the sentence

- This is a **prediction** task
 - ↳ Given **previous text** we **predict** the **next outcome**
- What about the data?
 - ↳ We parse the text into a sequence of words
 - ↳ We represent the letters of each word with their numeric, e.g., ASCII
 - ↳ We save each word into an N -dimensional frame

Learning from Sequence Data: Example II

Now let's consider **another example**: we have a **long text** and want to learn what is the **next word** in the sentence

- This is a **prediction** task
 - ↳ Given **previous text** we **predict** the **next outcome**
- What about the data?
 - ↳ We parse the text into a sequence of words
 - ↳ We represent the letters of each word with their numeric, e.g., ASCII
 - ↳ We save each word into an N -dimensional frame

...therapy. UofT undergraduate students explore the use of AI to treat speech

$x[1]$ $x[2]$ $x[3]$ \dots $x[t]$ $x[t+1]$ \dots $x[T]$ **$y = ?$**

Learning from Sequence Data: *Example III*

Another example: we have a sequence of stock prices and are interested in the future price

- This is again a *prediction* task
 - ↳ Given *previous prices* we *predict* the *future price*

Learning from Sequence Data: *Example III*

Another example: we have a sequence of stock prices and are interested in the future price

- This is again a *prediction* task
 - ↳ Given *previous prices* we *predict* the *future price*
- What about the data?
 - ↳ We put daily prices in form of a *sequence*
 - ↳ We collect every couple of prices along with other indicators into a vector

Learning from Sequence Data: *Example III*

Another example: we have a sequence of stock prices and are interested in the future price

- This is again a *prediction* task
 - ↳ Given *previous prices* we *predict* the *future price*
- What about the data?
 - ↳ We put daily prices in form of a *sequence*
 - ↳ We collect every couple of prices along with other indicators into a vector

In all these problems: we have a sequence of data and we intend to learn from them in a *generalizable way*

Learning from Sequence Data: *Example III*

Another example: we have a sequence of stock prices and are interested in the future price

- This is again a *prediction* task
 - ↳ Given *previous prices* we *predict* the *future price*
- What about the data?
 - ↳ We put daily prices in form of a *sequence*
 - ↳ We collect every couple of prices along with other indicators into a vector

In all these problems: we have a sequence of data and we intend to learn from them in a *generalizable way*

- ↳ We clearly need a *memory* component that can potentially be *infinite*

Learning from Sequence Data: *Example III*

Another example: we have a sequence of stock prices and are interested in the future price

- This is again a *prediction* task
 - ↳ Given *previous prices* we *predict* the *future price*
- What about the data?
 - ↳ We put daily prices in form of a *sequence*
 - ↳ We collect every couple of prices along with other indicators into a vector

In all these problems: we have a sequence of data and we intend to learn from them in a *generalizable way*

- ↳ We clearly need a *memory* component that can potentially be *infinite*
- ↳ We should keep track of this *infinitely large memory* via limited storage

Predicting Next Word

Let's start with a simple example: *we want to train a neural network that gets a sentence and complete the next word*

$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $y = x[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word

Let's start with a simple example: *we want to train a neural network that gets a sentence and complete the next word*

$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $y = x[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

- + *How does the training dataset look like?*
- We are given with several long texts in the same context: *at each entry of sequence in each of these texts the whole sequence is the data-point and the next word is label*

Predicting Next Word

Let's make some specification to clarify the problem


- Each entry is a *vector of dimension N* , i.e., $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with *some hidden layers* to train

Predicting Next Word

Let's make some specification to clarify the problem

- Each entry is a **vector of dimension N** , i.e., $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with **some hidden layers** to train

We show our NN **compactly** with the following diagram




In this diagram

Predicting Next Word

Let's make some specification to clarify the problem

- Each entry is a **vector of dimension N** , i.e., $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with **some hidden layers** to train

We show our NN **compactly** with the following diagram



In this diagram


- **Green box** shows the input layer

Predicting Next Word

Let's make some specification to clarify the problem

- Each entry is a **vector of dimension N** , i.e., $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with **some hidden layers** to train

We show our NN **compactly** with the following diagram



In this diagram


- **Green box** shows the input layer
- **Yellow box** includes hidden layers
 - ↳ It could be several layers

Predicting Next Word

Let's make some specification to clarify the problem

- Each entry is a **vector of dimension N** , i.e., $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with **some hidden layers** to train

We show our NN **compactly** with the following diagram



In this diagram


- **Green box** shows the input layer
- **Yellow box** includes hidden layers
 - ↳ It could be several layers
- **Red box** is the output layer

Predicting Next Word

Let's make some specification to clarify the problem

- Each entry is a **vector of dimension N** , i.e., $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with **some hidden layers** to train

We show our NN **compactly** with the following diagram




In this diagram

- **Green box** shows the input layer
- **Yellow box** includes hidden layers
 - ↳ It could be several layers
- **Red box** is the output layer
- Arrows refer to all links between the layers
 - ↳ They could be **learnable**


Predicting Next Word

For instance, we could think of following equivalence



Predicting Next Word: *MLP*

Let's try solving this problem with a *simple MLP*




We have a fully-connected FNN that

- takes N inputs, i.e., *single entry*
- returns N outputs, i.e., *predicted next word*

We train this MLP


- we go over all text

Predicting Next Word: *MLP*




$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*




$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*




$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*




$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*




$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*




$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*



$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her


Predicting Next Word: *MLP*



$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Does *FNN* *predict* “her”?

Predicting Next Word: *MLP*



$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Does FNN *predict* “her”? No! How can it *remember* we are talking about *Julia*?!

Predicting Next Word: *MLP*

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: *MLP*

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Each time the FNN gets trained with a new sample, it *forgets* previous text

Predicting Next Word: *MLP*

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Each time the FNN gets trained with a new sample, it *forgets previous text*

- At the end, it has a set of weights that are average over all *predictions*
 - ↳ many of these predictions are *irrelevant*, e.g.,
 - ↳ “nominated to” is followed by “receive”: has nothing to say about “her”

Predicting Next Word: MLP

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
 ... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Each time the FNN gets trained with a new sample, it *forgets previous text*

- At the end, it has a set of weights that are average over all *predictions*
 - ↳ many of these predictions are *irrelevant*, e.g.,
 - ↳ “nominated to” is followed by “receive”: has nothing to say about “her”
- By the time we get to $\mathbf{x}[t]$, the FNN gets no input that connects it to **Julia**

Predicting Next Word: *MLP*

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... *Julia* has been nominated to receive Alexander von Humboldt Prize for her

Each time the FNN gets trained with a new sample, it *forgets previous text*

- At the end, it has a set of weights that are average over all *predictions*
 - ↳ many of these predictions are *irrelevant*, e.g.,
 - ↳ “nominated to” is followed by “receive”: has nothing to say about “her”
- By the time we get to $\mathbf{x}[t]$, the FNN gets no input that connects it to *Julia*


This indicates that we need to make a *memory component* for our NN

Predicting Next Word: *Large MLP*

*Maybe we could give **more inputs** to the FNN!*


Predicting Next Word: *Large MLP*

Maybe we could give *more inputs* to the FNN!



Predicting Next Word: *Large MLP*


Maybe we could give *more inputs* to the FNN!



But, what is *Julia* has been mentioned 10 pages ago? Forget about *large* MLPs!

Predicting Next Word: CNN

Let's now think about **CNNs**




We have a fully-connected FNN that

- takes N inputs, i.e., **single entry**
- returns N outputs, i.e., **predicted next word**

Predicting Next Word: CNN

Let's now think about CNNs




We have a fully-connected FNN that

- takes N inputs, i.e., *single entry*
- returns N outputs, i.e., *predicted next word*

We use convolution to


- look into a *larger input* by a *filter of size N*
- extract *features* from a *larger part of text* and pass it to hidden layers

Predicting Next Word: CNN




... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: CNN



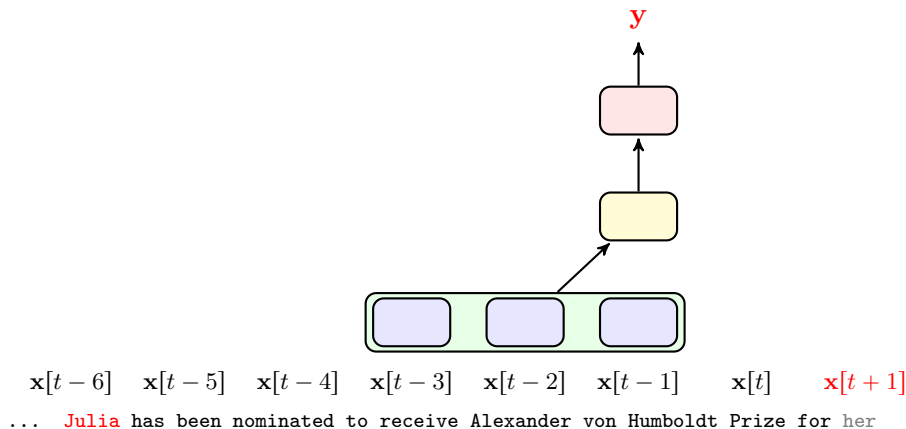
$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: CNN




... Julia has been nominated to receive Alexander von Humboldt Prize for her

Predicting Next Word: CNN



Predicting Next Word: CNN




$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$ $x[t+1]$

... Julia has been nominated to receive Alexander von Humboldt Prize for her

Yet, it doesn't seem to *remember Julia*! Unless we *slide over the whole text*!

Predicting Next Word: *Large CNN*



Though better than MLP, it is still *infeasible* to track *the whole text*

Finite Memory: *Root Problem*

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

The problem with *all architectures we know* is that they have *finite memory*

Finite Memory: *Root Problem*

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

The problem with *all* architectures we know is that they have *finite memory*

- They can only *remember* from *their input*
 - ↳ we always give them *independent inputs* with *similar features*
 - ↳ they gradually learn to connect any of *such inputs* to *its label*

Finite Memory: Root Problem

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

The problem with *all* architectures we know is that they have *finite memory*

- They can only *remember* from *their input*
 - ↳ we always give them *independent inputs* with *similar features*
 - ↳ they gradually learn to connect any of *such inputs* to *its label*
- If we need to *remember* for *long time* we have to give them *huge inputs*

Finite Memory: Root Problem

$\mathbf{x}[t-6]$ $\mathbf{x}[t-5]$ $\mathbf{x}[t-4]$ $\mathbf{x}[t-3]$ $\mathbf{x}[t-2]$ $\mathbf{x}[t-1]$ $\mathbf{x}[t]$ $\mathbf{x}[t+1]$
 ... Julia has been nominated to receive Alexander von Humboldt Prize for her

The problem with *all* architectures we know is that they have *finite memory*

- They can only *remember* from *their input*
 - ↳ we always give them *independent inputs* with *similar features*
 - ↳ they gradually learn to connect any of *such inputs* to *its label*
- If we need to *remember* for *long time* we have to give them *huge inputs*
- But the *memory component* does *not* seem to be so *huge*
 - ↳ We may only remember that *Julia* is a “*single person*” and “*female*”
 - ↳ If text now switches to *Theodore* we should refresh our *memory* that we are talking about a “*single person*” and “*male*”

Finite Memory Component with *Infinite Response*

Component We Miss

We need to extract a *memory component* from our *data* that is *finite in size* but has been *influenced* (at least theoretically) *infinitely*

Finite Memory Component with *Infinite Response*

Component We Miss

We need to extract a *memory component* from our *data* that is *finite in size* but has been *influenced* (at least theoretically) *infinitely*

State-space model can help us building *such memory component*: it is widely used in control theory to describe *evolution* of a system *over time*

Finite Memory Component with Infinite Response

Component We Miss

We need to extract a *memory component* from our *data* that is *finite in size* but has been *influenced* (at least theoretically) *infinitely*

State-space model can help us building *such memory component*: it is widely used in control theory to describe *evolution* of a system *over time*

Assume $\mathbf{x}[t]$ is an *input* to a system at time t : the system returns an *output* $\mathbf{y}[t]$ to this input and a *state variable* $\mathbf{s}[t]$. The output in the next time, i.e., $t + 1$, depends on the *new input* and *current state*, i.e.,

$$\mathbf{y}[t + 1], \mathbf{s}[t + 1] = f(\mathbf{x}[t + 1], \mathbf{s}[t])$$


In the above representation: the *state* is a *finite-size variable* that carries information for *infinitely long time*

State-Space Model for NNs

We can look at a NN as a **state-dependent** system


In this architecture, the NN

- takes $\mathbf{x}[t]$ and previous **state** $\mathbf{s}[t - 1]$ as inputs
- returns $\mathbf{y}[t]$ and new **state** $\mathbf{s}[t]$ as outputs



State-Space Model for NNs

We can look at a NN as a **state-dependent** system



In this architecture, the NN


- takes $x[t]$ and previous **state** $s[t - 1]$ as inputs
- returns $y[t]$ and new **state** $s[t]$ as outputs

This NN captures temporal behavior of data: starting from $t = 1$, the NN

- initiates some $s[0]$

State-Space Model for NNs

We can look at a NN as a **state-dependent** system



In this architecture, the NN


- takes $x[t]$ and previous **state** $s[t - 1]$ as inputs
- returns $y[t]$ and new **state** $s[t]$ as outputs

This NN captures temporal behavior of data: starting from $t = 1$, the NN

- initiates some $s[0]$
- computes $y[1]$ and $s[1]$ from time sample $x[1]$ and $s[0]$

State-Space Model for NNs

We can look at a NN as a **state-dependent** system



In this architecture, the NN


- takes $x[t]$ and previous **state** $s[t - 1]$ as inputs
- returns $y[t]$ and new **state** $s[t]$ as outputs

This NN captures temporal behavior of data: starting from $t = 1$, the NN

- initiates some $s[0]$
- computes $y[1]$ and $s[1]$ from time sample $x[1]$ and $s[0]$
- computes $y[2]$ and $s[2]$ from $x[2]$ and $s[1]$

State-Space Model for NNs

We can look at a NN as a **state-dependent** system



In this architecture, the NN

- takes $x[t]$ and previous **state** $s[t - 1]$ as inputs
- returns $y[t]$ and new **state** $s[t]$ as outputs

This NN captures temporal behavior of data: starting from $t = 1$, the NN


- initiates some $s[0]$
- computes $y[1]$ and $s[1]$ from time sample $x[1]$ and $s[0]$
- computes $y[2]$ and $s[2]$ from $x[2]$ and $s[1]$

...

State-Space Model for NNs

Theoretically, this NN has **infinite** time response

Say NN initiates with some $s[0]$. It takes only sample $x[1]$ and no other time samples is given to it. At time t ,




State-Space Model for NNs

Theoretically, this NN has **infinite** time response

Say NN initiates with some $s[0]$. It takes only sample $x[1]$ and no other time samples is given to it. At time t ,

- $y[t]$ depends on $s[t - 1]$




State-Space Model for NNs

Theoretically, this NN has **infinite** time response

Say NN initiates with some $s[0]$. It takes only sample $x[1]$ and no other time samples is given to it. At time t ,

- $y[t]$ depends on $s[t - 1]$
- $s[t - 1]$ depends on $s[t - 2]$




State-Space Model for NNs

Theoretically, this NN has **infinite** time response


Say NN initiates with some $s[0]$. It takes only sample $x[1]$ and no other time samples is given to it. At time t ,

- $y[t]$ depends on $s[t - 1]$
- $s[t - 1]$ depends on $s[t - 2]$
- ...



State-Space Model for NNs

Theoretically, this NN has **infinite** time response



Say NN initiates with some $s[0]$. It takes only sample $x[1]$ and no other time samples is given to it. At time t ,

- $y[t]$ depends on $s[t - 1]$
- $s[t - 1]$ depends on $s[t - 2]$
- ...
- $s[1]$ depends on $x[1]$

This means that $y[t]$ still **remembers** $x[1]$

State-Dependent NNs

It seems that for our purpose *state-space model helps extracting a good memory components. The challenge is to design a good state-dependent NN*

State-Dependent NNs

It seems that for our purpose *state-space model* helps extracting a good memory components. *The challenge is to design a good state-dependent NN*

- + Why is it a *challenge*? We make an NN with input (\mathbf{x}, \mathbf{s}) and output $(\mathbf{y}, \mathbf{s}')$ and then train it!
- That sounds *easy*, but have some *challenges*

State-Dependent NNs

It seems that for our purpose *state-space model* helps extracting a good memory components. *The challenge is to design a good state-dependent NN*

- + Why is it a **challenge**? We make an NN with input (\mathbf{x}, \mathbf{s}) and output $(\mathbf{y}, \mathbf{s}')$ and then train it!
- That sounds **easy**, but have some **challenges**

Design of state-dependent NNs has two main **challenges**

① Defining the **state variable**

↳ how should we **specify** the **state** as we do not have a clear clue about it

State-Dependent NNs

It seems that for our purpose *state-space model* helps extracting a good memory components. The challenge is to design a good state-dependent NN

- + Why is it a **challenge**? We make an NN with input (\mathbf{x}, \mathbf{s}) and output $(\mathbf{y}, \mathbf{s}')$ and then train it!
- That sounds **easy**, but have some **challenges**

Design of state-dependent NNs has two main **challenges**

- 1 Defining the **state variable**
 - ↳ how should we **specify** the **state** as we do not have a clear clue about it
- 2 Training the NN over time: say we get a time sequence data and compute the **empirical risk** by **averaging** over time samples **up to time t**
 - ↳ **empirical risk** depends on **weights** in the NN
 - ↳ it also depends **previous memory components** which can be **learnable**

State-Dependent NNs

It seems that for our purpose *state-space model* helps extracting a good memory components. The challenge is to design a good state-dependent NN

- + Why is it a **challenge**? We make an NN with input (\mathbf{x}, \mathbf{s}) and output $(\mathbf{y}, \mathbf{s}')$ and then train it!
- That sounds **easy**, but have some **challenges**

Design of state-dependent NNs has two main **challenges**

- 1 Defining the **state variable**
 - ↳ how should we **specify** the **state** as we do not have a clear clue about it
- 2 Training the NN over time: say we get a time sequence data and compute the **empirical risk** by **averaging** over time samples **up to time t**
 - ↳ **empirical risk** depends on **weights** in the NN
 - ↳ it also depends **previous memory components** which can be **learnable**
 - ↳ if we want to train the model **efficiently**, we need to **get back over time** and update all those **memory components**!

State-Dependent NNs

Several attempts have been done: *we look briefly into two of them*

- *Jordan Network*

- ↳ *proposed by Michael Jordan¹ in 1986*

- ↳ *it computes the **state variable** to be a **simple moving average***

¹Professor at CS Department of University of Berkeley; check his [page](#)

²Late Professor at UCSD ('48 - '18); check his [page](#)

State-Dependent NNs

Several attempts have been done: we look *briefly* into two of them

- *Jordan Network*
 - ↳ proposed by Michael Jordan¹ in 1986
 - ↳ it computes the *state variable* to be a *simple moving average*
- *Elman Network*
 - ↳ proposed by Jeffrey Elman² in 1990
 - ↳ it *learns* the *state variable* but *does not track back completely over time*

¹Professor at CS Department of University of Berkeley; check his [page](#)

²Late Professor at UCSD ('48 - '18); check his [page](#)

State-Dependent NNs

Several attempts have been done: we look *briefly* into two of them

- *Jordan Network*
 - ↳ proposed by Michael Jordan¹ in 1986
 - ↳ it computes the *state variable* to be a *simple moving average*
- *Elman Network*
 - ↳ proposed by Jeffrey Elman² in 1990
 - ↳ it *learns* the *state variable* but *does not track back completely over time*

These models are simple forms of what we nowadays know as

Recurrent NNs \equiv *RNNs*

¹Professor at CS Department of University of Berkeley; check his [page](#)

²Late Professor at UCSD ('48 - '18); check his [page](#)

State-Dependent NNs

Several attempts have been done: we look *briefly* into two of them

- *Jordan Network*
 - ↳ proposed by Michael Jordan¹ in 1986
 - ↳ it computes the *state variable* to be a *simple moving average*
- *Elman Network*
 - ↳ proposed by Jeffrey Elman² in 1990
 - ↳ it *learns* the *state variable* but *does not track back completely over time*

These models are simple forms of what we nowadays know as

Recurrent NNs \equiv *RNNs*

RNN


RNN is an *NN* with *state-variable*: the term *recurrent* refers to the connection between *former state* and *new output*

¹Professor at CS Department of University of Berkeley; check his [page](#)

²Late Professor at UCSD ('48 - '18); check his [page](#)

Jordan Model

Jordan Network uses a simple *moving average* of *output* as the state




The proposal was a shallow NN

- It starts with some *initial state* $s[1]$

Jordan Model

Jordan Network uses a simple **moving average** of **output** as the state




The proposal was a shallow NN

- It starts with some **initial state** $s[1]$
- In time t , it updates the **state** $s[t]$ with fixed μ as

$$s[t] = \mu s[t - 1] + y[t - 1]$$

Jordan Model

Jordan Network uses a simple **moving average** of **output** as the state



The proposal was a shallow NN

- It starts with some **initial state** $s[1]$
- In time t , it updates the **state** $s[t]$ with fixed μ as


$$s[t] = \mu s[t-1] + y[t-1]$$

- The hidden layer then computes

$$h[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{s}[t])$$

Jordan Model

Jordan Network uses a simple **moving average** of **output** as the state



The proposal was a shallow NN

- It starts with some **initial state** $s[1]$
- In time t , it updates the **state** $s[t]$ with fixed μ as

$$s[t] = \mu s[t-1] + y[t-1]$$

- The hidden layer then computes

$$h[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{s}[t])$$

- The output is

$$y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$$

Jordan Network

Jordan Network has a **memory** component with **infinte** response: say we set *initial state to zero*, give $\mathbf{x}[1]$, and keep *the input zero* for the rest of time; then,

$$\mathbf{y}[1] = f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))$$

$$\mathbf{y}[2] = f(\mathbf{W}_2 f(\mathbf{W}_m \mathbf{y}[1])) = f(\mathbf{W}_2 f(\mathbf{W}_m f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))))$$

$$\vdots$$

Jordan Network

Jordan Network has a **memory** component with **infinite** response: say we set **initial state to zero**, give **$\mathbf{x}[1]$** , and keep **the input zero** for the rest of time; then,

$$\mathbf{y}[1] = f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))$$

$$\mathbf{y}[2] = f(\mathbf{W}_2 f(\mathbf{W}_m \mathbf{y}[1])) = f(\mathbf{W}_2 f(\mathbf{W}_m f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))))$$

$$\vdots$$

But, the network does **not** learn how to remember


- It takes a fixed **moving average** as **memory component**
- It only learns how to use this **pre-defined memory** for **prediction**
 - ↳ We could say that it **implicitly** learn to **remember** by learning \mathbf{W}_m

Elman Network

Elman Network uses *output of hidden layer* as state: also called *hidden state*


The proposal was a shallow NN

- It starts with some *initial hidden state* $h[0]$



Elman Network

Elman Network uses *output of hidden layer* as state: also called *hidden state*




The proposal was a shallow NN

- It starts with some *initial hidden state* $h[0]$
- In time t , it updates the *state* $s[t]$ as

$$s[t] = h[t - 1]$$

Elman Network

Elman Network uses *output of hidden layer* as state: also called *hidden state*



The proposal was a shallow NN

- It starts with some *initial hidden state* $\mathbf{h}[0]$
- In time t , it updates the *state* $\mathbf{s}[t]$ as


$$\mathbf{s}[t] = \mathbf{h}[t - 1]$$

- The hidden layer then computes

$$\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{s}[t])$$

Elman Network

Elman Network uses *output of hidden layer* as state: also called *hidden state*



The proposal was a shallow NN

- It starts with some *initial hidden state* $\mathbf{h}[0]$
- In time t , it updates the *state* $\mathbf{s}[t]$ as

$$\mathbf{s}[t] = \mathbf{h}[t - 1]$$

- The hidden layer then computes

$$\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{s}[t])$$

- The output is

$$\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$$

Elman Network

Similar to Jordan Network, Elman Network has a **memory** component with **infinte** response: with *zero initial hidden state*, we have

$$\mathbf{y}[1] = f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))$$

$$\mathbf{y}[2] = f(\mathbf{W}_2 f(\mathbf{W}_m \mathbf{h}[1])) = f(\mathbf{W}_2 f(\mathbf{W}_m f(\mathbf{W}_1 \mathbf{x}[1])))$$

$$\vdots$$

Elman network also learns how to remember only implicitly

Challenge of Learning Through Time

Though Jordan and Elman Networks had memory, they did **not** get trained accurately over time, i.e., they **simplified the solution** to the second **challenge**

- + What is really this **challenge**?
- We are going to deal with it in next sections, but let's see it on these **simple networks** first

³We will see that this is not always this easy!

Challenge of Learning Through Time

Though Jordan and Elman Networks had memory, they did **not** get trained accurately over time, i.e., they **simplified the solution** to the second **challenge**

- + What is really this **challenge**?
- We are going to deal with it in next sections, but let's see it on these **simple networks** first

Let's assume a simple setting: we are to train our NN on **single data sequence**

- We have the sequence $\mathbf{x}[1], \dots, \mathbf{x}[T]$ as the **data-point**

³We will see that this is not always this easy!

Challenge of Learning Through Time

Though Jordan and Elman Networks had memory, they did **not** get trained accurately over time, i.e., they **simplified the solution** to the second **challenge**

- + What is really this **challenge**?
- We are going to deal with it in next sections, but let's see it on these **simple networks** first

Let's assume a simple setting: we are to train our NN on **single data sequence**

- We have the sequence $\mathbf{x}[1], \dots, \mathbf{x}[T]$ as the **data-point**
- For each entry of this sequence we have the **true label**
 - ↳ We have sequence $\mathbf{v}[1], \dots, \mathbf{v}[T]$ with $\mathbf{v}[t]$ being label of $\mathbf{x}[t]$

³We will see that this is not always this easy!

Challenge of Learning Through Time

Though Jordan and Elman Networks had memory, they did **not** get trained accurately over time, i.e., they **simplified the solution** to the second **challenge**

- + What is really this **challenge**?
- We are going to deal with it in next sections, but let's see it on these **simple networks** first

Let's assume a simple setting: we are to train our NN on **single data sequence**


- We have the sequence $\mathbf{x}[1], \dots, \mathbf{x}[T]$ as the **data-point**
- For each entry of this sequence we have the **true label**
 - ↳ We have sequence $\mathbf{v}[1], \dots, \mathbf{v}[T]$ with $\mathbf{v}[t]$ being label of $\mathbf{x}[t]$
- We are able to compute the loss between **outputs** and true labels as³

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

³We will see that this is not always this easy!

Recap: *Basic FNN*

For sake if comparison, let's first train a basic FNN on this data sequence



We have a shallow FNN

- *The hidden layer computes*


$$\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t])$$

- *The output is*

$$\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$$

Recap: *Basic FNN*


For sake of comparison, let's first train a basic FNN on this data sequence



How do we train this FNN?

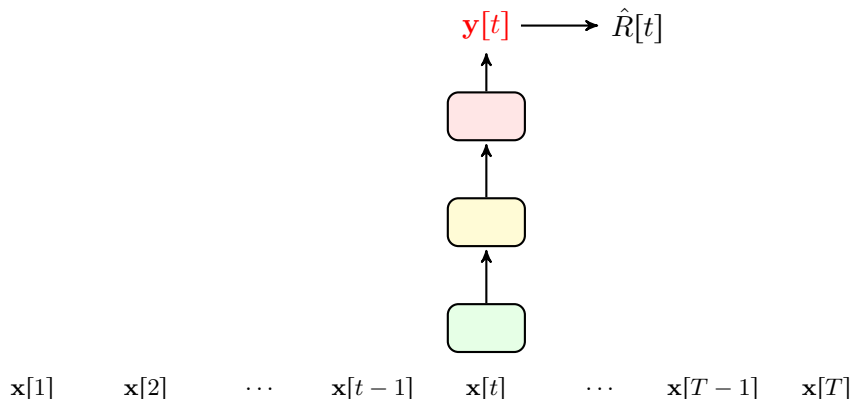
- We compute the gradients $\nabla_{\mathbf{w}_1} \hat{R}$ and $\nabla_{\mathbf{w}_2} \hat{R}$
 - ↳ We do it via backpropagation
- We apply gradient descent

Learning Through Time: *FNNs*



$$\hat{R} = \sum_{t=1}^T \underbrace{\mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])}_{\hat{R}[t]}$$

Learning Through Time: *FNNs*



$$\hat{R} = \sum_{t=1}^T \underbrace{\mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])}_{\hat{R}[t]} = \sum_{t=1}^T \hat{R}[t] \rightsquigarrow \nabla_{\mathbf{w}_i} \hat{R} = \sum_{t=1}^T \nabla_{\mathbf{w}_i} \hat{R}[t]$$

Learning Through Time: *FNNs*

Let's learn \mathbf{W}_1 and to ease computation we use our [cheating notation](#), i.e., use \circ to show *any product*:

Learning Through Time: *FNNs*

Let's learn \mathbf{W}_1 and to ease computation we use our **cheating notation**, i.e., use \circ to show **any product**: to compute the gradient we start with the **output**

$$\begin{aligned}\nabla_{\mathbf{W}_1} \hat{R}[t] &= \nabla_{\mathbf{W}_1} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) \\ &= \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]\end{aligned}$$

Learning Through Time: *FNNs*

Let's learn \mathbf{W}_1 and to ease computation we use our **cheating notation**, i.e., use \circ to show **any product**: to compute the gradient we start with the **output**

$$\begin{aligned}\nabla_{\mathbf{W}_1} \hat{R}[t] &= \nabla_{\mathbf{W}_1} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) \\ &= \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]\end{aligned}$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \left(\dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \right) \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]\end{aligned}$$

Learning Through Time: *FNNs*

Let's learn \mathbf{W}_1 and to ease computation we use our *cheating notation*, i.e., use \circ to show *any product*: to compute the gradient we start with the *output*

$$\begin{aligned}\nabla_{\mathbf{W}_1} \hat{R}[t] &= \nabla_{\mathbf{W}_1} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) \\ &= \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]\end{aligned}$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \left(\dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \right) \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]\end{aligned}$$

What about $\nabla_{\mathbf{W}_1} \mathbf{h}[t]$? We keep on backward!

Learning Through Time: *FNNs*

Up to now, we have

$$\nabla_{\mathbf{w}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{w}_1} \mathbf{h}[t]$$

Learning Through Time: *FNNs*

Up to now, we have

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

We use the fact that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t])$

$$\nabla_{\mathbf{W}_1} \mathbf{h}[t] = \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t]) + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{x}[t]$$

Learning Through Time: FNNs

Up to now, we have

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

We use the fact that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t]) + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{x}[t] \\ &= \dot{f}(\mathbf{W}_1 \mathbf{x}[t]) \circ \mathbf{x}[t] + \left(\dot{f}(\mathbf{W}_1 \mathbf{x}[t]) \circ \mathbf{W}_1 \right) \circ \underbrace{\mathbf{0}}_{\mathbf{x}[t] \text{ is not a function of } \mathbf{W}_1} \end{aligned}$$

Learning Through Time: FNNs

Up to now, we have

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

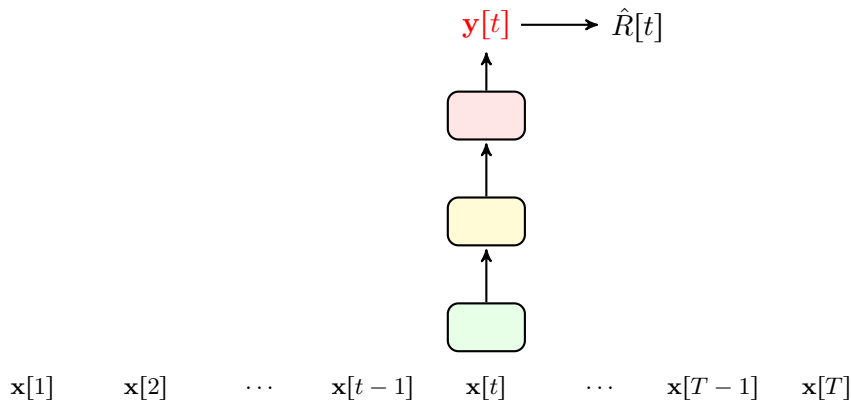
We use the fact that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t]) + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{x}[t] \\ &= \dot{f}(\mathbf{W}_1 \mathbf{x}[t]) \circ \mathbf{x}[t] + \left(\dot{f}(\mathbf{W}_1 \mathbf{x}[t]) \circ \mathbf{W}_1 \right) \circ \underbrace{\mathbf{0}}_{\mathbf{x}[t] \text{ is not a function of } \mathbf{W}_1} \end{aligned}$$

Therefore, we end chain rule here


$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

Learning Through Time: *FNNs*




$$\nabla_{\mathbf{w}_1} \hat{R}[t] =$$

Learning Through Time: *FNNs*




$$\nabla_{\mathbf{w}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t]$$

Learning Through Time: *FNNs*




$$\nabla_{\mathbf{w}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t]$$

Learning Through Time: *FNNs*




$$\nabla_{\mathbf{w}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{w}_1} \mathbf{h}[t]$$

Learning Through Time: *FNNs*



$$\nabla_{\mathbf{w}_2} \hat{R} = \sum_{t=1}^T \nabla_{\mathbf{w}_2} \hat{R}[t]$$


Learning Through Time: *FNNs*



$$\nabla_{\mathbf{w}_2} \hat{R} = \sum_{t=1}^T \nabla_{\mathbf{w}_2} \hat{R}[t]$$

Training a *Basic RNN*

Now, let's train **Elman network** on this sequence



The proposal was a shallow NN


- It starts with some **initial hidden state** $h[0]$
- The hidden layer then computes

$$h[t] = f(W_1 x[t] + W_m h[t-1])$$

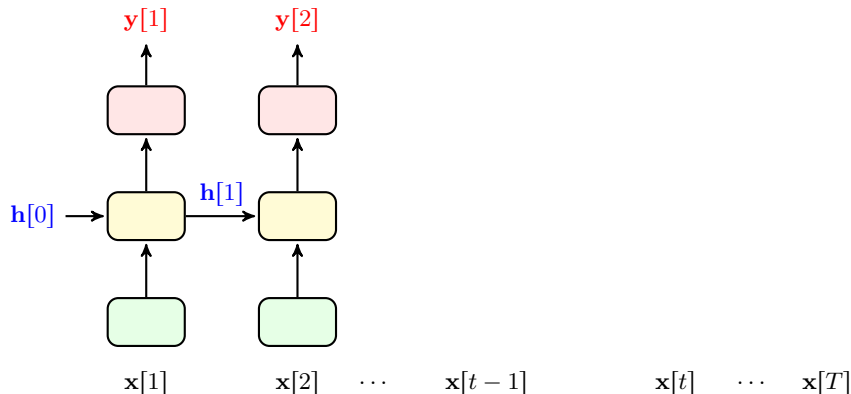
- The output is

$$y[t] = f(W_2 h[t])$$

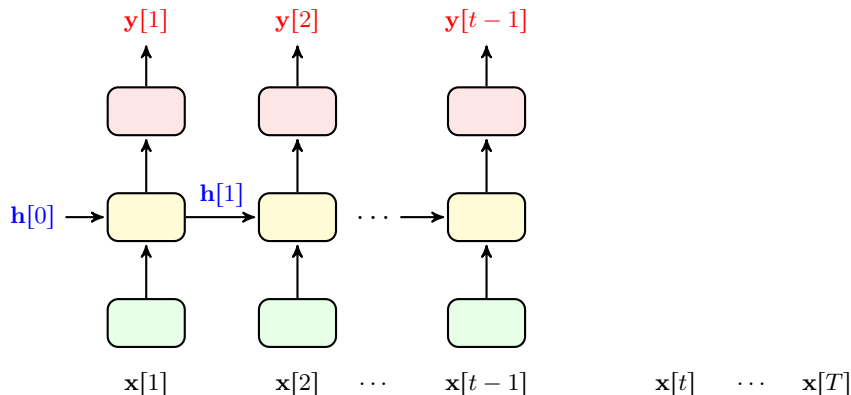
Inferring Through Time: *Elman Network*



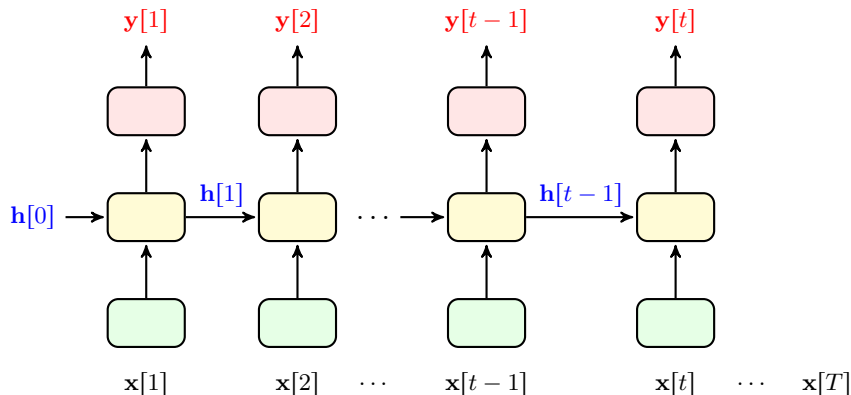
Inferring Through Time: *Elman Network*




Inferring Through Time: *Elman Network*



Inferring Through Time: *Elman Network*




Inferring Through Time: *Elman Network*



$$\hat{R} = \sum_{t=1}^T \underbrace{\mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])}_{\hat{R}[t]} = \sum_{t=1}^T \hat{R}[t]$$

Inferring Through Time: *Elman Network*



$$\hat{R} = \sum_{t=1}^T \underbrace{\mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])}_{\hat{R}[t]} = \sum_{t=1}^T \hat{R}[t]$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : *we start with the output*

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : we start with the *output*

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\nabla_{\mathbf{W}_1} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : we start with the **output**

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \end{aligned}$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : we start with the **output**

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \end{aligned}$$

Next, we note that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$

$$\nabla_{\mathbf{W}_1} \mathbf{h}[t] =$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : we start with the **output**

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \end{aligned}$$

Next, we note that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$

$$\nabla_{\mathbf{W}_1} \mathbf{h}[t] = \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : we start with the **output**

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \end{aligned}$$

Next, we note that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1]) \\ &\quad + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t]}_0 \end{aligned}$$

Learning Through Time: *Elman Network*

Let's again try to learn \mathbf{W}_1 : we start with the **output**

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$, so we can write

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \end{aligned}$$

Next, we note that $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1]) \\ &\quad + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-1]} \mathbf{h}[t] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-1]}_? \end{aligned}$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t - 1] = f(\mathbf{W}_1 \mathbf{x}[t - 1] + \mathbf{W}_m \mathbf{h}[t - 2])$

$$\nabla_{\mathbf{W}_1} \mathbf{h}[t - 1] =$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2])$

$$\nabla_{\mathbf{W}_1} \mathbf{h}[t-1] = \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2])$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1\mathbf{x}[t-1] + \mathbf{W}_m\mathbf{h}[t-2])$

$$\begin{aligned}\nabla_{\mathbf{W}_1}\mathbf{h}[t-1] &= \nabla_{\mathbf{W}_1}f(\mathbf{W}_1\mathbf{x}[t-1] + \mathbf{W}_m\mathbf{h}[t-2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]}\mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1}\mathbf{x}[t-1]}_0\end{aligned}$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t-1] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t-1]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-2]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-2]}_? \end{aligned}$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t-1] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t-1]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-2]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-2]}_? \end{aligned}$$

We are not still done! We know $\mathbf{h}[t-2] = f(\mathbf{W}_1 \mathbf{x}[t-2] + \mathbf{W}_m \mathbf{h}[t-3])$

$$\nabla_{\mathbf{W}_1} \mathbf{h}[t-2] =$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1\mathbf{x}[t-1] + \mathbf{W}_m\mathbf{h}[t-2])$

$$\begin{aligned}\nabla_{\mathbf{W}_1}\mathbf{h}[t-1] &= \nabla_{\mathbf{W}_1}f(\mathbf{W}_1\mathbf{x}[t-1] + \mathbf{W}_m\mathbf{h}[t-2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]}\mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1}\mathbf{x}[t-1]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-2]}\mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1}\mathbf{h}[t-2]}_?\end{aligned}$$

We are not still done! We know $\mathbf{h}[t-2] = f(\mathbf{W}_1\mathbf{x}[t-2] + \mathbf{W}_m\mathbf{h}[t-3])$

$$\nabla_{\mathbf{W}_1}\mathbf{h}[t-2] = \nabla_{\mathbf{W}_1}f(\mathbf{W}_1\mathbf{x}[t-2] + \mathbf{W}_m\mathbf{h}[t-3])$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t-1] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t-1]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-2]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-2]}_? \end{aligned}$$

We are not still done! We know $\mathbf{h}[t-2] = f(\mathbf{W}_1 \mathbf{x}[t-2] + \mathbf{W}_m \mathbf{h}[t-3])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t-2] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-2] + \mathbf{W}_m \mathbf{h}[t-3]) \\ &\quad + \nabla_{\mathbf{x}[t-2]} \mathbf{h}[t-2] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t-2]}_0 \end{aligned}$$

Learning Through Time: *Elman Network*

Well! We know that $\mathbf{h}[t-1] = f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2])$

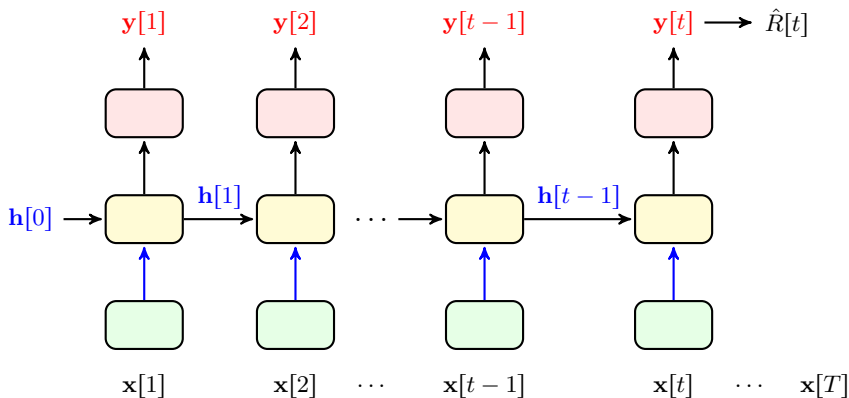
$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{h}[t-1] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-1] + \mathbf{W}_m \mathbf{h}[t-2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t-1]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-2]} \mathbf{h}[t-1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-2]}_?\end{aligned}$$

We are not still done! We know $\mathbf{h}[t-2] = f(\mathbf{W}_1 \mathbf{x}[t-2] + \mathbf{W}_m \mathbf{h}[t-3])$

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{h}[t-2] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t-2] + \mathbf{W}_m \mathbf{h}[t-3]) \\ &\quad + \nabla_{\mathbf{x}[t-2]} \mathbf{h}[t-2] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t-2]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-3]} \mathbf{h}[t-2] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-3]}_?\end{aligned}$$

Learning Through Time: *Elman Network*

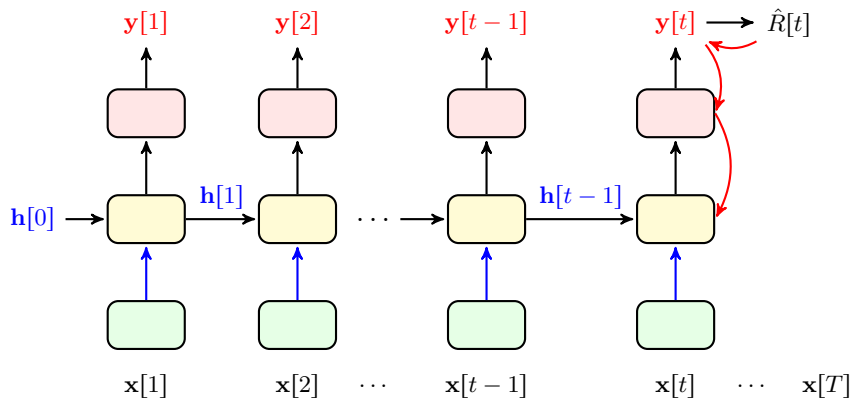
We should in fact pass all the way *back* to the *initial time interval time!*



Note that all *blue edges* are representing \mathbf{W}_1

Learning Through Time: *Elman Network*

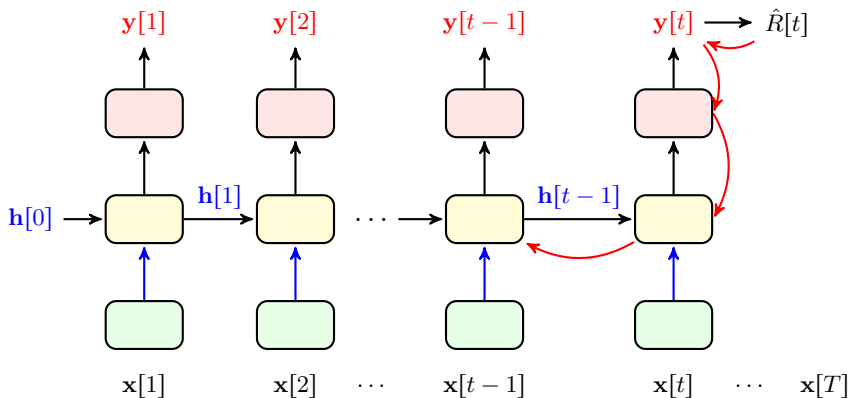
We should in fact pass all the way *back* to the *initial time interval time*!



Note that all *blue edges* are representing W_1

Learning Through Time: *Elman Network*

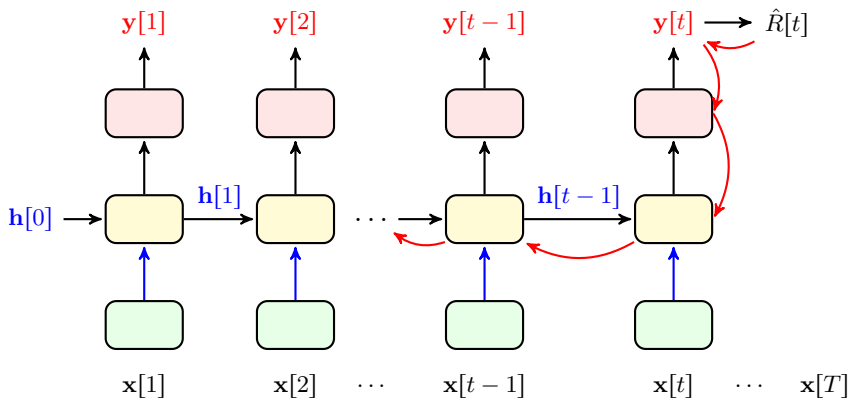
We should in fact pass all the way *back* to the *initial time interval time*!



Note that all *blue edges* are representing \mathbf{W}_1

Learning Through Time: *Elman Network*

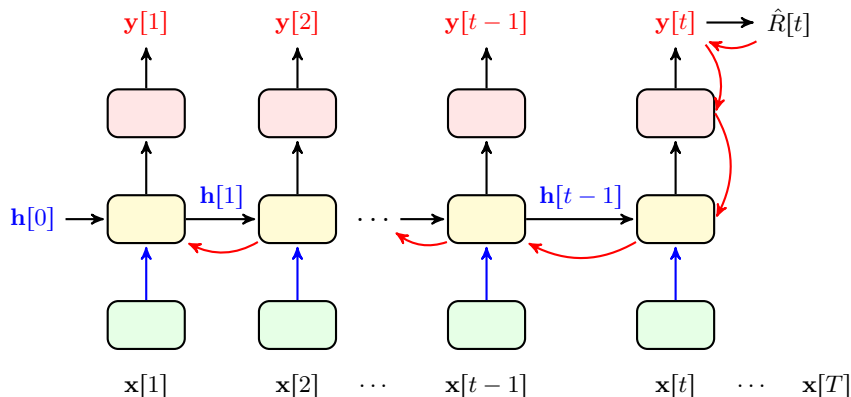
We should in fact pass all the way *back* to the *initial time interval time*!



Note that all *blue edges* are representing W_1

Learning Through Time: *Elman Network*

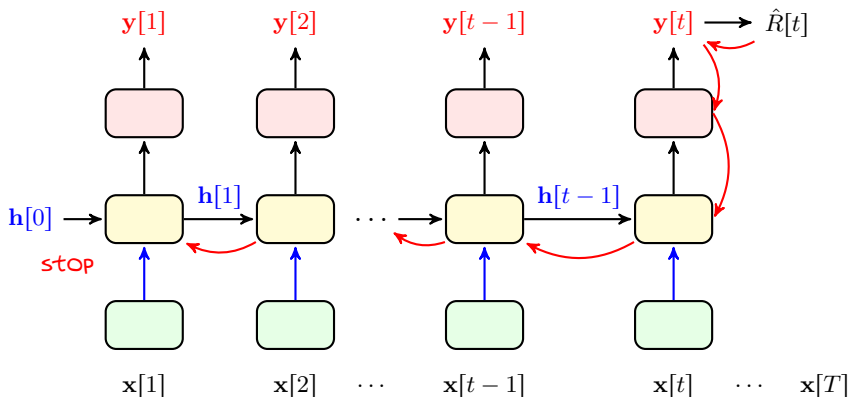
We should in fact pass all the way *back* to the *initial time interval time*!



Note that all *blue edges* are representing \mathbf{W}_1

Learning Through Time: *Elman Network*

We should in fact pass all the way *back* to the *initial time interval time*!



Note that all *blue edges* are representing W_1

Learning Through Time

Moral of Story

To learn how to *remember*, we need to *train* our RNN *through time*: at each time interval, we should move *all the way back to origin* to find out how *exactly* we should change the weights!

Learning Through Time

Moral of Story

To learn how to *remember*, we need to *train* our RNN *through time*: at each time interval, we should move *all the way back to origin* to find out how *exactly* we should change the weights!

- + Did *Elman* did so?
- *Not* really!

Learning Through Time

Moral of Story

To learn how to *remember*, we need to *train* our RNN *through time*: at each time interval, we should move *all the way back to origin* to find out how *exactly* we should change the weights!

- + Did *Elman* did so?
- *Not* really!

For training *Elman* treated the hidden state $\mathbf{h}[t - 1]$ as a fixed variable, i.e., he assumed $\nabla_{\mathbf{w}_1} \mathbf{h}[t - 1] \approx \nabla_{\mathbf{w}_m} \mathbf{h}[t - 1] = \mathbf{0}$! So, he did not need to move backward in time!

Learning Through Time

Moral of Story

To learn how to *remember*, we need to *train* our RNN *through time*: at each time interval, we should move *all the way back to origin* to find out how *exactly* we should change the weights!

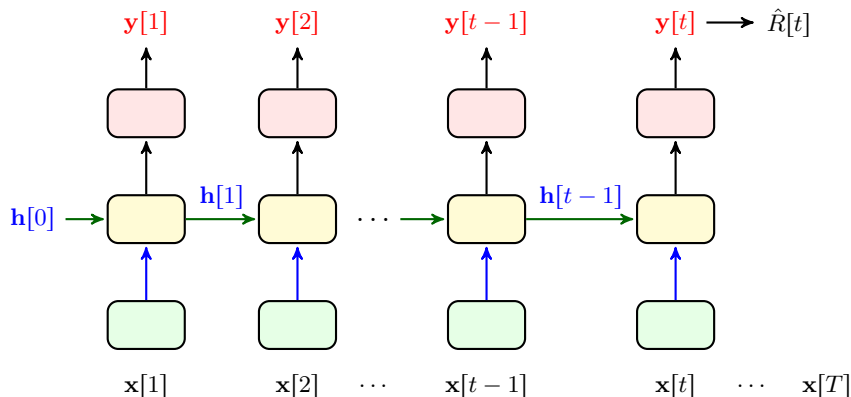
- + Did *Elman* did so?
- *Not* really!

For training *Elman* treated the hidden state $\mathbf{h}[t - 1]$ as a fixed variable, i.e., he assumed $\nabla_{\mathbf{w}_1} \mathbf{h}[t - 1] \approx \nabla_{\mathbf{w}_m} \mathbf{h}[t - 1] = \mathbf{0}$! So, he did not need to move backward in time!

This means that *Elman* did *not* really addressed *the second challenge*!

Learning Through Time: *Elman's Approximation*

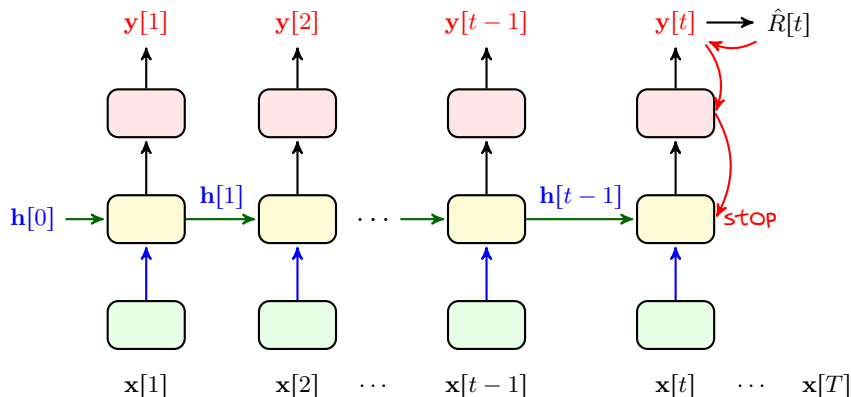
Elman treated it as a simple FNN with only one extra input!



Training \mathbf{W}_1 is exactly as FNN. We just have one extra \mathbf{W}_m here!

Learning Through Time: *Elman's Approximation*

Elman treated it as a simple FNN with only one extra input!



Training \mathbf{W}_1 is exactly as FNN. We just have one extra \mathbf{W}_m here!

RNNs: *Need to Learn Memory*

Though appreciated, Elman and Jordan Networks did not do the job

- ① Their *memory* component is rather *simple*
 - ↳ We should use *deeper models* that enable advanced memory components
- ② They do *not* really *learn* how to *remember*
 - ↳ We should *train* the memory component *over time*

RNNs: Need to Learn Memory

Though appreciated, Elman and Jordan Networks did not do the job

- 1 Their *memory* component is rather *simple*
 - ↳ We should use *deeper models* that enable advanced memory components
- 2 They do *not* really *learn* how to *remember*
 - ↳ We should *train* the memory component *over time*

These led us to development of RNNs!

RNN: Less Generic Definition

An RNN can be designed with *any known architecture* by letting NN also *learn from its past features* and *outputs*. This new enabling is called *recurrence*

RNN: *Dataset and Learning Setting*

We are now looking into a supervised learning problem where we are to learn *label* from a *sequence of data* that has generally a *temporal correlation*

RNN: Dataset and Learning Setting

We are now looking into a **supervised learning** problem where we are to learn **label** from a **sequence of data** that has generally a **temporal correlation**

Let's denote the **sequence** with $\mathbf{x}[1], \dots, \mathbf{x}[T]$

Temporal Correlation

By **temporal correlation** we mean that entries at **other time instances** carry information about one particular entry $\mathbf{x}[t]$

RNN: Dataset and Learning Setting

We are now looking into a supervised learning problem where we are to learn *label* from a *sequence of data* that has generally a *temporal correlation*

Let's denote the *sequence* with $\mathbf{x}[1], \dots, \mathbf{x}[T]$


Temporal Correlation

By *temporal correlation* we mean that entries at *other time instances* carry information about one particular entry $\mathbf{x}[t]$

- + But how is a *label* assigned to this sequence?
- Well, that can be of various forms!

Types of Problems with Sequence Data

We considered a very simple case: *many-to-many type I*




In this case, every entry has a *label*

↳ *speech tagging*: $x[t]$ is a part of speech and $y[t]$ is its tag

Types of Problems with Sequence Data

We considered a very simple case: *many-to-many type II*




In this case, we get a *label* once after multiple entries

↳ *speech recognition: $x[t]$ is a small part of speech and $y[t]$ says what is a every couple of minutes about*

Types of Problems with Sequence Data

We considered a very simple case: *many-to-many type III*




In this case, we start to get *labels* after some delay

↳ language translation: $x[1], \dots, x[t-1]$ is a sentence in German and $y[t-1], \dots, y[T]$ is its translation to English

Types of Problems with Sequence Data

We considered a very simple case: *one-to-many*




In this case, we get only one input data and have a sequence of *labels*

↳ image captioning: $x[1]$ is an image and $y[1], \dots, y[T]$ is a caption describing what is inside the image

Types of Problems with Sequence Data


We considered a very simple case: *many-to-one*



In this case, we get only one *label* for a whole input sequence

↳ *sequence classification*: $x[1], \dots, x[T]$ is a speech and $y[T]$ says if this speech is constructive or destructive


Types of Problems with Sequence Data: *FNNs*



In fact, FNNs are *one-to-one* RNNs

↳ we can think of every data-point as *a sequence of length one*, or

Types of Problems with Sequence Data: *FNNs*



In fact, FNNs are *one-to-one* RNNs

- ↳ we can think of every data-point as *a sequence of length one*, or
- ↳ we may think of dataset as a long sequence with *no temporal correlation*

RNN: General Form

To construct an RNN, we can use any module that we have learned:

- we can use a *fully-connected layer*
- we can use a *convolutional layer*
- we can use a *residual unit*
- we may use an *inception* unit used in *GoogLeNet*
- . . .

RNN: General Form

To construct an RNN, we can use any module that we have learned:

- we can use a *fully-connected layer*
- we can use a *convolutional layer*
- we can use a *residual unit*
- we may use an *inception* unit used in *GoogLeNet*
- . . .

But, classical implementations use *fully-connected layers*

RNN: General Form

To construct an RNN, we can use any module that we have learned:

- we can use a *fully-connected layer*
- we can use a *convolutional layer*
- we can use a *residual unit*
- we may use an *inception* unit used in *GoogLeNet*
- ...


But, classical implementations use *fully-connected layers*

We can use *one* layer to make a *shallow* RNN or *multiple* to make a *deep* RNN

- + Don't we do any *change* to them?
- **Not** a *serious change*
 - ↳ We may change the *activation* to $\tanh(\cdot)$: *we will see later why*
 - ↳ We expand the input dimension: since we need to also give *memory* as input

Shallow RNN

Let's break it down a bit: say the **yellow box** is a **fully-connected layer**




This layer gets input $\mathbf{x} \in \mathbb{R}^N$ and returns **activated** feature $\mathbf{h} \in \mathbb{R}^M$

- We replace its **activation** by $\tanh(\cdot)$
 - ↳ Not necessary, but usually suggested
- We modify it to get a new input in \mathbb{R}^{N+M}
 - ↳ N entries for **data inputs** \mathbf{x} in time t
 - ↳ M entries for **features** \mathbf{h} in time $t - 1$

Shallow RNN

Let's break it down a bit: say the **yellow box** is a **fully-connected** layer




We show it this way now

- We call $h[t]$ usually the **hidden state**
- We pass the **hidden state** through an output layer
↳ Output is **not necessarily** corresponding to label

Shallow RNN

We may show our **shallow RNN** **compactly** via the following diagram




In this diagram

- Each edge is a set of weights, e.g., a weight matrix
- The return edge also has a delay in time

Shallow RNN

We may show our **shallow RNN** **compactly** via the following diagram




In this diagram

- Each edge is a set of weights, e.g., a weight matrix
 - The return edge also has a delay in time
- + But, isn't that simply **Elman Network**?
- If we use a fully-connected layer and sigmoid activation; then, Yes! But, Remember that
- Elman did **not** **train it over time**
 - Elman in its model used **sigmoid activation**

Deep RNN

We can add more layers to make a **deep RNN**




In this diagram

- Each edge is a set of weights, e.g., a weight matrix
- The return edge also has a delay in time
- And, it's no more Elman Network

Deep RNN


It might be easier to think of it as the following diagram



We can use any module that we like

Deep RNN

It might be easier to think of it as the following diagram




We can use any module that we like

- We may add *skip connection*

Deep RNN

It might be easier to think of it as the following diagram




We can use any module that we like

- We may add *skip connection*
- We could make *dense connections*

Deep RNN

It might be easier to think of it as the following diagram




We can use any module that we like

- We may add *skip connection*
- We could make *dense connections*
- We may *skip over time*

Deep RNN

It might be easier to think of it as the following diagram



We can use any module that we like

- We may add **skip connection**
- We could make **dense connections**
- We may **skip over time**
- We may replace hidden layers with **convolutions** or anything else


Shallow RNN: *Elman-Like Network*

Let's start with simple RNN: a shallow RNN with *fully-connected* hidden layer

+ You mean *Elman Network*?

- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again



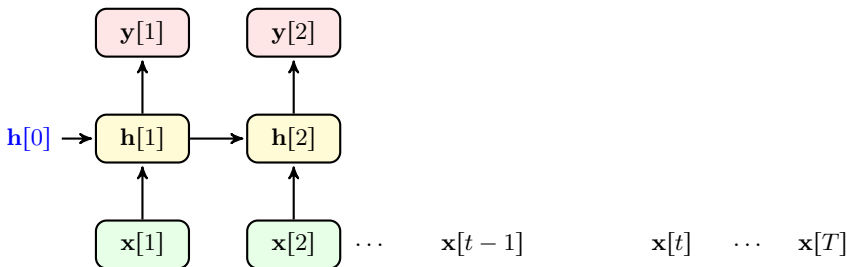
Shallow RNN: *Elman-Like Network*

Let's start with simple RNN: a shallow RNN with *fully-connected* hidden layer

+ You mean *Elman Network*?

- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again




Shallow RNN: *Elman-Like Network*

Let's start with simple RNN: a shallow RNN with *fully-connected* hidden layer

+ You mean *Elman Network*?

- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again




Shallow RNN: *Elman-Like Network*

Let's start with simple RNN: a shallow RNN with *fully-connected* hidden layer

+ You mean *Elman Network*?

- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again




Shallow RNN: *Elman-Like Network*

Let's start with simple RNN: a shallow RNN with *fully-connected* hidden layer

+ You mean *Elman Network*?


- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again



Shallow RNN: *Forward Propagation*


Let's specify the learnable parameters with some colors



Say we set the activation to $f(\cdot)$

Shallow RNN: *Forward Propagation*

Let's specify the learnable parameters with some colors




Say we set the activation to $f(\cdot)$

- 1 We have $y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$: we *can learn* \mathbf{W}_2

Shallow RNN: Forward Propagation

Let's specify the learnable parameters with some colors




Say we set the activation to $f(\cdot)$

- ① We have $y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$: we can learn \mathbf{W}_2
- ② We have $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$: we can learn \mathbf{W}_1 and \mathbf{W}_m

Shallow RNN: Forward Propagation

Let's specify the learnable parameters with some colors




Say we set the activation to $f(\cdot)$

- ① We have $y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$: we *can learn* \mathbf{W}_2
- ② We have $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$: we *can learn* \mathbf{W}_1 and \mathbf{W}_m

Shallow RNN: Forward Propagation

Let's specify the learnable parameters with some colors




Say we set the activation to $f(\cdot)$

- 1 We have $y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$: we *can learn* \mathbf{W}_2
- 2 We have $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$: we *can learn* \mathbf{W}_1 and \mathbf{W}_m
- 3 We can start with any *hidden state*: this means we *can learn* $\mathbf{h}[0]$ too


Training our Shallow RNN

We now want to train this basic RNN



Training our Shallow RNN

We now want to train this basic RNN



Let's consider training for only one sample


We assume that we have a sequence of labels $\mathbf{v}[1], \dots, \mathbf{v}[T]$

- ① We know some $\mathbf{v}[t]$ could be empty, e.g., in many-to-one scenario
- ② So could be some of $\mathbf{x}[t]$'s, e.g., in one-to-many scenario

We however have no problem with that!

Training our Shallow RNN

We now want to train this basic RNN



Let's consider training for only one sample


The loss in general can be written as

$$\hat{R} = \mathcal{L}(\mathbf{y}[1:T], \mathbf{v}[1:T])$$

where we use shorten notation $\mathbf{y}[1:T] = \mathbf{y}[1], \dots, \mathbf{y}[T]$


Training our Shallow RNN

We can think of such a diagram



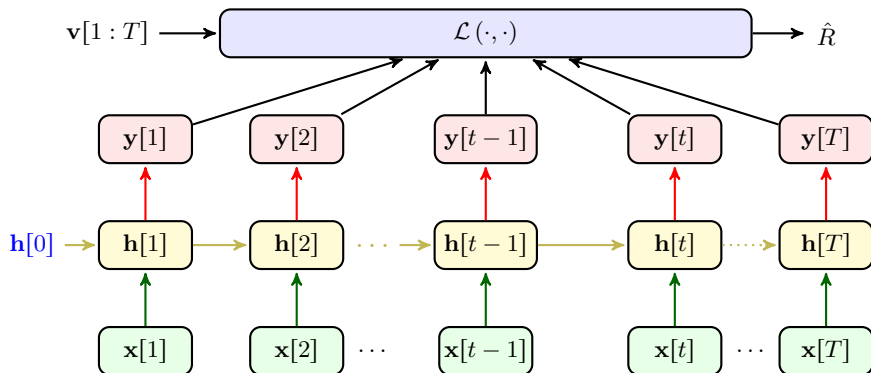
Training our Shallow RNN

We can think of such a diagram



Training our Shallow RNN


We can think of such a diagram



- + It seems to be hard to get back from \hat{R} to each $y[t]$
- Well! That's true, but we have some remedy for it

Training our Shallow RNN


We can think of such a diagram



For the moment, we assume that we can
compute the $\nabla_{y[t]} \hat{R}$ for all t

Training our Shallow RNN

We can think of such a diagram



For the moment, we assume that we can

compute the $\nabla_{y[t]} \hat{R}$ for all $t \equiv$ *move backward from \hat{R} to any $y[t]$*

Backward Pass Through Time

Starting from \hat{R} , say we want to find $\nabla_{\mathbf{w}_m} \hat{R}$

- Since \hat{R} is a function of $\mathbf{y}[1 : T]$, we should write a *vectorized chain rule*

$$\nabla_{\mathbf{w}_m} \hat{R} =$$

Backward Pass Through Time

Starting from \hat{R} , say we want to find $\nabla_{\mathbf{w}_m} \hat{R}$

- Since \hat{R} is a function of $\mathbf{y}[1:T]$, we should write a *vectorized chain rule*

$$\nabla_{\mathbf{w}_m} \hat{R} = \nabla_{\mathbf{y}[1]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[1]$$

Backward Pass Through Time

Starting from \hat{R} , say we want to find $\nabla_{\mathbf{w}_m} \hat{R}$

- Since \hat{R} is a function of $\mathbf{y}[1:T]$, we should write a *vectorized chain rule*

$$\nabla_{\mathbf{w}_m} \hat{R} = \nabla_{\mathbf{y}[1]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[1] + \dots + \nabla_{\mathbf{y}[T]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[T]$$

Backward Pass Through Time

Starting from \hat{R} , say we want to find $\nabla_{\mathbf{w}_m} \hat{R}$

- Since \hat{R} is a function of $\mathbf{y}[1 : T]$, we should write a *vectorized chain rule*

$$\begin{aligned}\nabla_{\mathbf{w}_m} \hat{R} &= \nabla_{\mathbf{y}[1]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[1] + \dots + \nabla_{\mathbf{y}[T]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[T] \\ &= \sum_{t=1}^T \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[t]\end{aligned}$$

Backward Pass Through Time

Starting from \hat{R} , say we want to find $\nabla_{\mathbf{w}_m} \hat{R}$

- Since \hat{R} is a function of $\mathbf{y}[1 : T]$, we should write a *vectorized chain rule*

$$\begin{aligned}\nabla_{\mathbf{w}_m} \hat{R} &= \nabla_{\mathbf{y}[1]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[1] + \dots + \nabla_{\mathbf{y}[T]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[T] \\ &= \sum_{t=1}^T \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[t]\end{aligned}$$

- Since we assumed that we have $\nabla_{\mathbf{y}[t]} \hat{R}$, our main task is to find

$$\nabla_{\mathbf{w}_m} \mathbf{y}[t]$$

for all t : we could hence compute it for a *general t*

↳ This is a tensor-like gradient, i.e., $[\nabla_{\mathbf{w}_m} y_1[t], \dots, \nabla_{\mathbf{w}_m} y_M[t]]$

Backward Pass Through Time

Starting from \hat{R} , say we want to find $\nabla_{\mathbf{w}_m} \hat{R}$

- Since \hat{R} is a function of $\mathbf{y}[1 : T]$, we should write a *vectorized chain rule*

$$\begin{aligned}\nabla_{\mathbf{w}_m} \hat{R} &= \nabla_{\mathbf{y}[1]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[1] + \dots + \nabla_{\mathbf{y}[T]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[T] \\ &= \sum_{t=1}^T \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{w}_m} \mathbf{y}[t]\end{aligned}$$

- Since we assumed that we have $\nabla_{\mathbf{y}[t]} \hat{R}$, our main task is to find

$$\nabla_{\mathbf{w}_m} \mathbf{y}[t]$$

for all t : we could hence compute it for a *general t*

↳ This is a tensor-like gradient, i.e., $[\nabla_{\mathbf{w}_m} y_1[t], \dots, \nabla_{\mathbf{w}_m} y_M[t]]$


- Apparently, we should apply chain rule for several times!

Backward Pass Through Time

+ *But, this is going to be exhausting?*

Backward Pass Through Time

- + *But, this is going to be exhausting?*
- Well, we could again follow Albert Einstein advice!



*"Everything should be made as **simple** as possible, but not **simpler**!"*

Backward Pass Through Time: *Simple Example*

Let's consider a dummy RNN with all variables being *scalar*


- 1 We have $y[t] = f(w_2 h[t])$
- 2 We have $h[t] = f(w_1 x[t] + w_m h[t - 1])$
- 3 We start with *hidden state* $h[0]$

Backward Pass Through Time: *Simple Example*

Let's consider a dummy RNN with all variables being *scalar*

- 1 We have $y[t] = f(w_2 h[t])$
- 2 We have $h[t] = f(w_1 x[t] + w_m h[t-1])$
- 3 We start with *hidden state* $h[0]$

So the diagram gets simplified as below



Backward Pass Through Time: *Simple Example*

Starting from \hat{R} , say we want to find $\nabla_{w_m} \hat{R}$: *our main task is to find*

$$\frac{\partial y[t]}{\partial w_m}$$

Let's start the computation:

Backward Pass Through Time: *Simple Example*

Starting from \hat{R} , say we want to find $\nabla_{w_m} \hat{R}$: *our main task is to find*

$$\frac{\partial y[t]}{\partial w_m}$$

Let's start the computation: say we have passed forward through the RNN

- we now have $y[t]$, $h[t]$ and $x[t]$ for all t*

Backward Pass Through Time: *Simple Example*

Starting from \hat{R} , say we want to find $\nabla_{w_m} \hat{R}$: *our main task is to find*

$$\frac{\partial y[t]}{\partial w_m}$$

Let's start the computation: say we have passed forward through the RNN

- we now have $y[t]$, $h[t]$ and $x[t]$ for all t*

To go backward, we note that $y[t] = f(w_2 h[t])$

Backward Pass Through Time: *Simple Example*

Starting from \hat{R} , say we want to find $\nabla_{w_m} \hat{R}$: *our main task is to find*

$$\frac{\partial y[t]}{\partial w_m}$$

Let's start the computation: say we have passed forward through the RNN

- we now have $y[t]$, $h[t]$ and $x[t]$ for all t

To go backward, we note that $y[t] = f(w_2 h[t])$

- $y[t]$ is a function of $h[t]$: so we write the *chain rule* as

$$\frac{\partial y[t]}{\partial w_m} = \frac{\partial y[t]}{\partial h[t]} \frac{\partial h[t]}{\partial w_m}$$

↳ we can compute the first term as $\partial y[t] / \partial h[t] = w_2 \dot{f}(w_2 h[t])$

↳ we should compute the second term by further chain rule

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(w_1 x[t] + w_m h[t-1])$

- $h[t]$ is a function of w_m and $h[t-1]$.⁴

⁴We ignore w_1 and $x[t]$, as they are obviously not functions of w_m .

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(w_1 x[t] + w_m h[t-1])$

- $h[t]$ is a function of w_m and $h[t-1]$:⁴

↳ let's write $h[t] = g(w_m, h[t-1])$

⁴We ignore w_1 and $x[t]$, as they are obviously not functions of w_m

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(w_1 x[t] + w_m h[t-1])$

- $h[t]$ is a function of w_m and $h[t-1]$.⁴

↳ let's write $h[t] = g(w_m, h[t-1])$

- So we write the *chain rule* as

$$\frac{\partial h[t]}{\partial w_m} =$$

⁴We ignore w_1 and $x[t]$, as they are obviously not functions of w_m

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(w_1 x[t] + w_m h[t-1])$

- $h[t]$ is a function of w_m and $h[t-1]$.⁴
 ↳ let's write $h[t] = g(w_m, h[t-1])$
- So we write the *chain rule* as

$$\frac{\partial h[t]}{\partial w_m} = \frac{\partial g}{\partial w_m} \underbrace{\frac{\partial w_m}{\partial w_m}}_1$$

⁴We ignore w_1 and $x[t]$, as they are obviously not functions of w_m

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(w_1 x[t] + w_m h[t-1])$

- $h[t]$ is a function of w_m and $h[t-1]$.⁴

↳ let's write $h[t] = g(w_m, h[t-1])$

- So we write the *chain rule* as

$$\frac{\partial h[t]}{\partial w_m} = \frac{\partial g}{\partial w_m} \underbrace{\frac{\partial w_m}{\partial w_m}}_1 + \frac{\partial g}{\partial h[t-1]} \frac{\partial h[t-1]}{\partial w_m}$$

⁴We ignore w_1 and $x[t]$, as they are obviously not functions of w_m

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(w_1 x[t] + w_m h[t-1])$

- $h[t]$ is a function of w_m and $h[t-1]$.⁴

↳ let's write $h[t] = g(w_m, h[t-1])$

- So we write the *chain rule* as

$$\begin{aligned} \frac{\partial h[t]}{\partial w_m} &= \frac{\partial g}{\partial w_m} \underbrace{\frac{\partial w_m}{\partial w_m}}_1 + \frac{\partial g}{\partial h[t-1]} \frac{\partial h[t-1]}{\partial w_m} \\ &= \frac{\partial g}{\partial w_m} + \frac{\partial g}{\partial h[t-1]} \frac{\partial h[t-1]}{\partial w_m} \end{aligned}$$

⁴We ignore w_1 and $x[t]$, as they are obviously not functions of w_m

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(z[t])$

- Let's define $z[t] = w_1 x[t] + w_m h[t-1]$

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(z[t])$

- Let's define $z[t] = w_1 x[t] + w_m h[t-1]$
- $h[t] = g(w_m, h[t-1])$
 - ↳ we can compute $\partial g / \partial w_m = h[t-1] \dot{f}(z[t])$
 - ↳ we can compute $\partial g / \partial h[t-1] = w_m \dot{f}(z[t])$

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

We keep going backward, by noting that $h[t] = f(z[t])$

- Let's define $z[t] = w_1 x[t] + w_m h[t-1]$
- $h[t] = g(w_m, h[t-1])$
 - we can compute $\partial g / \partial w_m = h[t-1] \dot{f}(z[t])$
 - we can compute $\partial g / \partial h[t-1] = w_m \dot{f}(z[t])$
- So we can simplify the *chain rule* as

$$\begin{aligned} \frac{\partial h[t]}{\partial w_m} &= h[t-1] \dot{f}(z[t]) + w_m \dot{f}(z[t]) \frac{\partial h[t-1]}{\partial w_m} \\ &= \dot{f}(z[t]) \left(h[t-1] + w_m \frac{\partial h[t-1]}{\partial w_m} \right) \end{aligned}$$

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) \left(h[t-1] + w_m \frac{\partial h[t-1]}{\partial w_m} \right)$$

We keep going backward: $h[t-1] = f(z[t-1])$

↳ Recall that $z[t-1] = w_1 x[t-1] + w_m h[t-2]$

Backward Pass Through Time: *Simple Example*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) \left(h[t-1] + w_m \frac{\partial h[t-1]}{\partial w_m} \right)$$

We keep going backward: $h[t-1] = f(z[t-1])$

↳ Recall that $z[t-1] = w_1 x[t-1] + w_m h[t-2]$

- We just need to replace t with $t-1$ in the last derivation

$$\frac{\partial h[t-1]}{\partial w_m} = \dot{f}(z[t-1]) \left(h[t-2] + w_m \frac{\partial h[t-2]}{\partial w_m} \right)$$

Backward Pass Through Time: *Simple Example*

Backpropagation at time t is hence described by a pair of recursive equations

- *At time t , we compute*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

Backward Pass Through Time: *Simple Example*

Backpropagation at time t is hence described by a pair of recursive equations

- At time t , we compute

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

- For each $i = t, t - 1, \dots, 1$, we use

$$\frac{\partial h[i]}{\partial w_m} = \dot{f}(z[i]) \left(h[i - 1] + w_m \frac{\partial h[i - 1]}{\partial w_m} \right)$$

Backward Pass Through Time: *Simple Example*

Backpropagation at time t is hence described by a pair of recursive equations

- At time t , we compute

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$


- For each $i = t, t - 1, \dots, 1$, we use

$$\frac{\partial h[i]}{\partial w_m} = \dot{f}(z[i]) \left(h[i - 1] + w_m \frac{\partial h[i - 1]}{\partial w_m} \right)$$


- We stop at $i = 1$, where we get

$$\frac{\partial h[1]}{\partial w_m} = \dot{f}(z[1]) \left(h[0] + w_m \underbrace{\frac{\partial h[0]}{\partial w_m}}_0 \right) = \dot{f}(z[1]) h[0]$$


Backward Pass Through Time: *Simple Example*




Backward Pass Through Time: *Simple Example*




Backward Pass Through Time: *Simple Example*




Backward Pass Through Time: *Simple Example*



Backward Pass Through Time: *Simple Example*



Backward Pass Through Time: *Simple Example*



Key Point

Propagating *back in time* is described via a *recursive equation*

Backward Pass Through Time: *Simple Example*

Recursive equation has an interesting feature

- *Say we have backpropagated from time t*

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

↳ We hence have $\partial h[t]/\partial w_m, \partial h[t-1]/\partial w_m, \dots, \partial h[1]/\partial w_m$

Backward Pass Through Time: *Simple Example*

Recursive equation has an interesting feature

- Say we have backpropagated from time t

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

↳ We hence have $\partial h[t]/\partial w_m, \partial h[t-1]/\partial w_m, \dots, \partial h[1]/\partial w_m$

- Now, if we want to backpropagate from $t-1$

↳ We *already have* $\partial h[t-1]/\partial w_m, \partial h[t-2]/\partial w_m, \dots, \partial h[1]/\partial w_m$

↳ We do *not* need to backpropagate through time anymore

Backward Pass Through Time: *Simple Example*

Recursive equation has an interesting feature

- Say we have backpropagated from time t

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \frac{\partial h[t]}{\partial w_m}$$

↳ We hence have $\partial h[t]/\partial w_m, \partial h[t-1]/\partial w_m, \dots, \partial h[1]/\partial w_m$

- Now, if we want to backpropagate from $t-1$

↳ We *already have* $\partial h[t-1]/\partial w_m, \partial h[t-2]/\partial w_m, \dots, \partial h[1]/\partial w_m$

↳ We do *not* need to backpropagate through time anymore

Moral of Story

Pass *forward* till *end of sequence* and *backpropagate* to the beginning *just once*

Backpropagation Through Time (BPTT)

BPTT() :

- ① Start at t with $\nabla_{\mathbf{W}_m} \mathbf{y}[t] = \mathbf{W}_2 \circ \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \nabla_{\mathbf{W}_m} \mathbf{h}[t]$
- ② Go back in time as $\nabla_{\mathbf{W}_m} \mathbf{h}[i] = \dot{f}(\mathbf{z}[i]) \circ (\mathbf{h}[i-1] + \mathbf{W}_m \circ \nabla_{\mathbf{W}_m} \mathbf{h}[i-1])$
- ③ Stop at $i = 1$ with $\nabla_{\mathbf{W}_m} \mathbf{h}[1] = \dot{f}(\mathbf{z}[1]) \circ \mathbf{h}[0]$

Backpropagation Through Time (BPTT)

BPTT() :

- ① Start at t with $\nabla_{\mathbf{W}_m} \mathbf{y}[t] = \mathbf{W}_2 \circ \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \nabla_{\mathbf{W}_m} \mathbf{h}[t]$
- ② Go back in time as $\nabla_{\mathbf{W}_m} \mathbf{h}[i] = \dot{f}(\mathbf{z}[i]) \circ (\mathbf{h}[i-1] + \mathbf{W}_m \circ \nabla_{\mathbf{W}_m} \mathbf{h}[i-1])$
- ③ Stop at $i = 1$ with $\nabla_{\mathbf{W}_m} \mathbf{h}[1] = \dot{f}(\mathbf{z}[1]) \circ \mathbf{h}[0]$

- + It looks very similar to backpropagation in deep NNs!
- Exactly! Even simple RNN is very *deep* through time
- + Don't we experience *vanishing* or *exploding* gradient then!
- Yes! Let's check it out

BPTT: *Vanishing Gradient*

Let's expand the gradient in our example

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) h[t-1]$$

BPTT: *Vanishing Gradient*

Let's expand the gradient in our example

$$\frac{\partial y[t]}{\partial w_m} = w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) h[t-1] \\ + w_2 \dot{f}(w_2 h[t]) w_m \dot{f}(z[t]) \dot{f}(z[t-1]) h[t-2]$$

BPTT: *Vanishing Gradient*

Let's expand the gradient in our example

$$\begin{aligned}
 \frac{\partial y[t]}{\partial w_m} = & w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) h[t-1] \\
 & + w_2 \dot{f}(w_2 h[t]) w_m \dot{f}(z[t]) \dot{f}(z[t-1]) h[t-2] \\
 & + w_2 \dot{f}(w_2 h[t]) w_m^2 \dot{f}(z[t]) \dot{f}(z[t-1]) \dot{f}(z[t-2]) h[t-3] \\
 & + \dots
 \end{aligned}$$

BPTT: *Vanishing Gradient*

Let's expand the gradient in our example

$$\begin{aligned}
 \frac{\partial y[t]}{\partial w_m} = & w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) h[t-1] \\
 & + w_2 \dot{f}(w_2 h[t]) w_m \dot{f}(z[t]) \dot{f}(z[t-1]) h[t-2] \\
 & + w_2 \dot{f}(w_2 h[t]) w_m^2 \dot{f}(z[t]) \dot{f}(z[t-1]) \dot{f}(z[t-2]) h[t-3] \\
 & + \dots \\
 & + w_2 \dot{f}(w_2 h[t]) w_m^{i-1} \left(\prod_{j=0}^{i-1} \dot{f}(z[t-j]) \right) h[t-i] \\
 & + \dots
 \end{aligned}$$

BPTT: Vanishing Gradient

Let's expand the gradient in our example

$$\begin{aligned}
 \frac{\partial y[t]}{\partial w_m} = & w_2 \dot{f}(w_2 h[t]) \dot{f}(z[t]) h[t-1] \\
 & + w_2 \dot{f}(w_2 h[t]) w_m \dot{f}(z[t]) \dot{f}(z[t-1]) h[t-2] \\
 & + w_2 \dot{f}(w_2 h[t]) w_m^2 \dot{f}(z[t]) \dot{f}(z[t-1]) \dot{f}(z[t-2]) h[t-3] \\
 & + \dots \\
 & + w_2 \dot{f}(w_2 h[t]) w_m^{i-1} \left(\prod_{j=0}^{i-1} \dot{f}(z[t-j]) \right) h[t-i] \\
 & + \dots \\
 & + w_2 \dot{f}(w_2 h[t]) w_m^{t-1} \left(\prod_{j=0}^{t-1} \dot{f}(z[t-j]) \right) h[0]
 \end{aligned}$$

BPTT: *Vanishing Gradient*

This says that gradient of hidden state in i steps back in time is multiplied by

$$w_2 w_m^{i-1} \dot{f}(w_2 h[t]) \left(\prod_{j=0}^{i-1} \dot{f}(z[t-j]) \right)$$

BPTT: Vanishing Gradient

This says that gradient of hidden state in i steps back in time is multiplied by

$$w_2 w_m^{i-1} \dot{f}(w_2 h[t]) \left(\prod_{j=0}^{i-1} \dot{f}(z[t-j]) \right)$$

Recall **deep** NNs: we could see two cases

- If $\dot{f}(\cdot) > 1$ most of the time
 - ↳ **very old** hidden states can **explode** the gradient
 - ↳ this **usually** does **not** happen, because most activations are not like that

BPTT: Vanishing Gradient

This says that gradient of hidden state in i steps back in time is multiplied by

$$w_2 w_m^{i-1} \dot{f}(w_2 h[t]) \left(\prod_{j=0}^{i-1} \dot{f}(z[t-j]) \right)$$

Recall **deep** NNs: we could see two cases

- If $\dot{f}(\cdot) > 1$ most of the time
 - ↳ **very old** hidden states can **explode** the gradient
 - ↳ this **usually** does **not** happen, because most activations are not like that
- If $\dot{f}(\cdot) < 1$ most of the time
 - ↳ **very old** hidden states **have pretty much no impact on the gradient**
 - ↳ this means that the RNN can **train only up to a finite memory**
 - ↳ this **typically** happens and we call it **vanishing gradient through time**

Handling *Vanishing Gradient Through Time*

- + *Cant we use the same remedy as in deep NNs?*

Handling *Vanishing Gradient Through Time*

- + *Cant we use the same remedy as in deep NNs?*
- Yes and No!

Handling *Vanishing Gradient Through Time*

- + *Cant we use the same remedy as in deep NNs?*
- Yes and No!

It is important to note that *vanishing gradient through time* is a bit *different*: we are still updating weights with *large enough gradients*, but these gradients have *no information* of *long-term memory*

Handling Vanishing Gradient Through Time

- + Cant we use the same remedy as in deep NNs?
- Yes and No!

It is important to note that *vanishing gradient through time* is a bit *different*: we are still updating weights with *large enough gradients*, but these gradients have *no information* of *long-term memory*

Solutions to *vanishing gradient through time* are mainly two approaches

- Use $\tanh(\cdot)$ *activation*
 - ↳ $\tanh(\cdot)$ is able to keep memory for a longer period
- Use *truncated BPTT*
 - ↳ Repeat short-term BPTTs every couple of time intervals

Handling Vanishing Gradient Through Time

- + Cant we use the same remedy as in deep NNs?
- Yes and No!

It is important to note that *vanishing gradient through time* is a bit *different*: we are still updating weights with *large enough gradients*, but these gradients have *no information* of *long-term memory*

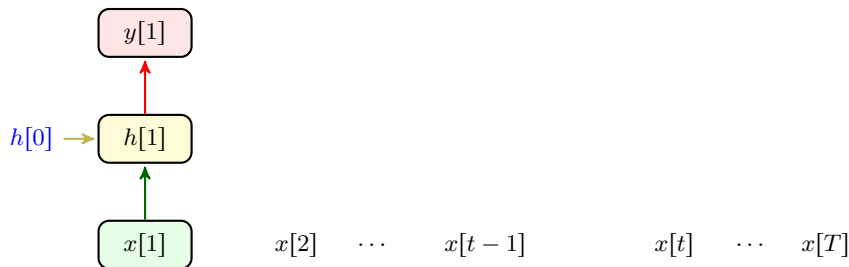
Solutions to *vanishing gradient through time* are mainly two approaches

- Use $\tanh(\cdot)$ *activation*
 - ↳ $\tanh(\cdot)$ is able to keep memory for a longer period
- Use *truncated BPTT*
 - ↳ Repeat short-term BPTTs every couple of time intervals
- Invoke *gating approach*
 - ↳ This is the most *sophisticated approach*
 - ↳ It was known for a long time, but received attention much later!

Principle of Gating

To understand the idea of *Gating*, let's get back to our *basic RNN*

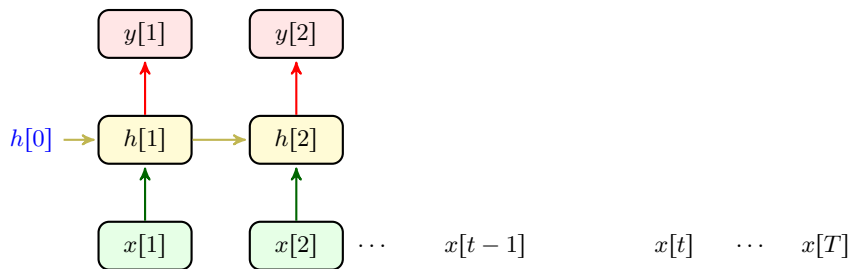
- 1 We start with *hidden state* $h[0]$
- 2 We have $h[t] = f(w_1x[t] + w_mh[t - 1])$
- 3 We have $y[t] = f(w_2h[t])$



Principle of Gating

To understand the idea of *Gating*, let's get back to our *basic RNN*


- 1 We start with *hidden state* $h[0]$
- 2 We have $h[t] = f(w_1x[t] + w_mh[t-1])$
- 3 We have $y[t] = f(w_2h[t])$



Principle of Gating

To understand the idea of *Gating*, let's get back to our *basic RNN*


- 1 We start with *hidden state* $h[0]$
- 2 We have $h[t] = f(w_1x[t] + w_mh[t-1])$
- 3 We have $y[t] = f(w_2h[t])$



Principle of Gating

To understand the idea of *Gating*, let's get back to our *basic RNN*


- 1 We start with *hidden state* $h[0]$
- 2 We have $h[t] = f(w_1x[t] + w_mh[t-1])$
- 3 We have $y[t] = f(w_2h[t])$



Principle of Gating

To understand the idea of *Gating*, let's get back to our *basic RNN*

- 1 We start with *hidden state* $h[0]$
- 2 We have $h[t] = f(w_1x[t] + w_mh[t-1])$
- 3 We have $y[t] = f(w_2h[t])$



Looking at $h[t]$ as *memory*, we can say *we are always updating the memory*

Principle of Gating

Recall our motivating example: we wanted to predict the *next word*

$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Principle of Gating

Recall our motivating example: we wanted to predict the *next word*

$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Should we updating the memory *all the way* from Julia?

Principle of Gating

Recall our motivating example: we wanted to predict the *next word*

$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$
... Julia has been nominated to receive Alexander von Humboldt Prize for her

Should we updating the memory *all the way* from Julia?

- Obviously No!
 - ↳ We should stop updating at Julia, since it is something we should remember

Principle of Gating

Recall our motivating example: we wanted to predict the *next word*

$x[t-6]$ $x[t-5]$ $x[t-4]$ $x[t-3]$ $x[t-2]$ $x[t-1]$ $x[t]$
... *Julia* has been nominated to receive Alexander von Humboldt Prize for her

Should *we updating the memory all the way from Julia*?

- Obviously No!
 - ↳ We should stop updating at *Julia*, since it is something we should remember
- How can we do it? *Let's try a thought experiment*

Principle of Gating

Say we access to a sequence $u[t] \in [0, 1]$: assume the following happens

- ↳ We arrive at “*Julia*” at time t_0 , i.e., $x[t_0] \propto$ “*Julia*”
- ↳ At time t_0 , we have $u[t_0] = 1$

Principle of Gating

Say we access to a sequence $u[t] \in [0, 1]$: assume the following happens

- ↳ We arrive at *“Julia”* at time t_0 , i.e., $x[t_0] \propto$ *“Julia”*
- ↳ At time t_0 , we have $u[t_0] = 1$
- ↳ We want to predict at $t + 1$
- ↳ From $t_0 + 1$ to t , we have $u[t_0 + 1] = \dots = u[t] = 0$

Principle of Gating

Say we access to a sequence $u[t] \in [0, 1]$: assume the following happens

- ↳ We arrive at “*Julia*” at time t_0 , i.e., $x[t_0] \propto$ “*Julia*”
- ↳ At time t_0 , we have $u[t_0] = 1$
- ↳ We want to predict at $t + 1$
- ↳ From $t_0 + 1$ to t , we have $u[t_0 + 1] = \dots = u[t] = 0$

Now, we update our memory like this

- 1 We start with hidden state $h[0]$
- 2 We compute $\tilde{h}[t] = f(w_1 x[t] + w_m h[t - 1])$
- 3
- 4 At each time, we have $y[t] = f(w_2 h[t])$

Principle of Gating

Say we access to a sequence $u[t] \in [0, 1]$: assume the following happens

- ↳ We arrive at “*Julia*” at time t_0 , i.e., $x[t_0] \propto$ “*Julia*”
- ↳ At time t_0 , we have $u[t_0] = 1$
- ↳ We want to predict at $t + 1$
- ↳ From $t_0 + 1$ to t , we have $u[t_0 + 1] = \dots = u[t] = 0$

Now, we update our memory like this

- 1 We start with hidden state $h[0]$
- 2 We compute $\tilde{h}[t] = f(w_1 x[t] + w_m h[t - 1])$
- 3 We update $h[t] = u[t] \tilde{h}[t] + (1 - u[t]) h[t - 1]$
- 4 At each time, we have $y[t] = f(w_2 h[t])$

Let's see what happens to the *memory*

Principle of Gating

At time t_0 , we can say

- We have $x[t_0] \propto$ “*Julia*” and compute $\tilde{h}[t_0] = f(w_1 x[t_0] + w_m h[t_0 - 1])$
 - ↳ $\tilde{h}[t_0]$ has fresh memory about “*Julia*”
 - ↳ Since $u[t_0] = 1$, we update as $h[t_0] = 1 \times \tilde{h}[t_0] + 0 \times h[t_0 - 1] = \tilde{h}[t_0]$
 - ↳ RNN has a fresh memory about “*Julia*”

Principle of Gating

At time t_0 , we can say

- We have $x[t_0] \propto$ “*Julia*” and compute $\tilde{h}[t_0] = f(w_1 x[t_0] + w_m h[t_0 - 1])$
 - ↳ $\tilde{h}[t_0]$ has fresh memory about “*Julia*”
 - ↳ Since $u[t_0] = 1$, we update as $h[t_0] = 1 \times \tilde{h}[t_0] + 0 \times h[t_0 - 1] = \tilde{h}[t_0]$
 - ↳ RNN has a fresh memory about “*Julia*”

At the next time, i.e., $t_0 + 1$, we have

- No matter what $\tilde{h}[t_0 + 1]$ is we have $u[t_0 + 1] = 0$
 - ↳ We update as $h[t_0 + 1] = 0 \times \tilde{h}[t_0 + 1] + 1 \times h[t_0] = h[t_0] = \tilde{h}[t_0]$
 - ↳ We still have a fresh memory about “*Julia*”

Principle of Gating

At time t_0 , we can say

- We have $x[t_0] \propto$ “Julia” and compute $\tilde{h}[t_0] = f(w_1 x[t_0] + w_m h[t_0 - 1])$
 - ↳ $\tilde{h}[t_0]$ has fresh memory about “Julia”
 - ↳ Since $u[t_0] = 1$, we update as $h[t_0] = 1 \times \tilde{h}[t_0] + 0 \times h[t_0 - 1] = \tilde{h}[t_0]$
 - ↳ RNN has a fresh memory about “Julia”

At the next time, i.e., $t_0 + 1$, we have

- No matter what $\tilde{h}[t_0 + 1]$ is we have $u[t_0 + 1] = 0$
 - ↳ We update as $h[t_0 + 1] = 0 \times \tilde{h}[t_0 + 1] + 1 \times h[t_0] = h[t_0] = \tilde{h}[t_0]$
 - ↳ We still have a fresh memory about “Julia”

This repeats from $t_0 + 1$ till t , so at time t

- No matter what $\tilde{h}[t]$, we have $u[t] = 0$
 - ↳ We update as $h[t] = 0 \times \tilde{h}[t] + 1 \times h[t - 1] = h[t - 1] = \dots = \tilde{h}[t_0]$
 - ↳ We still have a fresh memory about “Julia”

Principle of Gating: *Updating via Gates*

$u[t]$ *gates the memory*: it decides how much *memory* we should *pass* and *forget*

- We update $h[t] = u[t]\tilde{h}[t] + (1 - u[t])h[t - 1]$

Principle of Gating: *Updating via Gates*

$u[t]$ *gates the memory*: it decides how much *memory* we should *pass* and *forget*

- We update $h[t] = u[t]\tilde{h}[t] + (1 - u[t])h[t - 1]$
- + How does it help with vanishing gradient?
- Well! It is *implicitly* making *skip connections through time*

Principle of Gating: *Updating via Gates*

$u[t]$ *gates the memory*: it decides how much *memory* we should *pass* and *forget*

- We update $h[t] = u[t]\tilde{h}[t] + (1 - u[t])h[t - 1]$
- + How does it help with vanishing gradient?
- Well! It is *implicitly* making *skip connections through time*

Recall that with standard BPTT, we have for $i = t, t - 1, \dots, t_0$

$$\frac{\partial h[i]}{\partial w_m} = f'(z[i]) \left(h[i - 1] + w_m \frac{\partial h[i - 1]}{\partial w_m} \right)$$

Principle of Gating: *Updating via Gates*

$u[t]$ *gates the memory*: it decides how much *memory* we should *pass* and *forget*

- We update $h[t] = u[t]\tilde{h}[t] + (1 - u[t])h[t - 1]$
- + How does it help with vanishing gradient?
- Well! It is *implicitly* making *skip connections through time*

Recall that with standard BPTT, we have for $i = t, t - 1, \dots, t_0$

$$\frac{\partial h[i]}{\partial w_m} = f'(z[i]) \left(h[i - 1] + w_m \frac{\partial h[i - 1]}{\partial w_m} \right)$$

But, now we skip multiple time slots, as we have for $i = t, t - 1, \dots, t_0$

$$\frac{\partial h[i]}{\partial w_m} = \frac{\partial h[i - 1]}{\partial w_m}$$

Principle of Gating: A Generic Gate

+ *Sounds inspiring! But, how could you get $u[t]$?*

⁵We are going to drop **bias** and you all know why!

Principle of Gating: A Generic Gate

- + *Sounds inspiring! But, how could you get $u[t]$?*
- *Well! Like what we did the whole time: we **learn** it!*

⁵We are going to drop **bias** and you all know why!

Principle of Gating: A Generic Gate

- + Sounds inspiring! But, how could you get $u[t]$?
- Well! Like what we did the whole time: we *learn* it!

Gate

Let $\mathbf{x}[t]$ be input and $\mathbf{h}[t - 1]$ be last hidden state: a gate $\Gamma[t]$ is computed as

$$\Gamma[t] = \sigma (\mathbf{W}_{\Gamma,\text{in}}\mathbf{x}[t] + \mathbf{W}_{\Gamma,\text{m}}\mathbf{h}[t - 1] + \mathbf{b}_{\Gamma})$$

where $\sigma(\cdot)$ is sigmoid function, and $\mathbf{W}_{\Gamma,\text{in}}$, $\mathbf{W}_{\Gamma,\text{m}}$ and \mathbf{b}_{Γ} are *learnable*⁵

⁵We are going to drop *bias* and you all know why!

Principle of Gating: A Generic Gate

- + Sounds inspiring! But, how could you get $u[t]$?
- Well! Like what we did the whole time: we *learn* it!

Gate

Let $\mathbf{x}[t]$ be input and $\mathbf{h}[t - 1]$ be last hidden state: a gate $\Gamma[t]$ is computed as

$$\Gamma[t] = \sigma(\mathbf{W}_{\Gamma, \text{in}} \mathbf{x}[t] + \mathbf{W}_{\Gamma, \text{m}} \mathbf{h}[t - 1] + \mathbf{b}_{\Gamma})$$

where $\sigma(\cdot)$ is sigmoid function, and $\mathbf{W}_{\Gamma, \text{in}}$, $\mathbf{W}_{\Gamma, \text{m}}$ and \mathbf{b}_{Γ} are *learnable*⁵

- + Why do we use sigmoid function?
 - Simply because *it is between 0 and 1*
- + What should we set the dimension of $\Gamma[t]$?
 - Same as the variable (*memory* component) that *we want to gate*

⁵We are going to drop *bias* and you all know why!

Practical Gated Architectures

There are various *gated* architectures: *we look into two of them*

- 1 *Gated Recurrent Unit (GRU)*
- 2 *Long Short-Term Memory (LSTM)*

Practical Gated Architectures

There are various *gated* architectures: we look into two of them

- 1 Gated Recurrent Unit (GRU)
- 2 Long Short-Term Memory (LSTM)

Before we start, let's recall their basic RNN counterpart


Basic RNN Counterpart

Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- 1 Start with an initial *hidden state*
 \hookrightarrow we *can learn* $\mathbf{h}[0]$
- 2 Compute memory as $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$
 \hookrightarrow we *can learn* \mathbf{W}_1 and \mathbf{W}_m
- 3 Compute output $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t]) \rightsquigarrow f_{\text{out}}$ and f could be *different*
 \hookrightarrow we *can learn* \mathbf{W}_2


Classical Diagram: *Hidden Layer as Unit*

When we study a **gated architecture**: it is **common** to look at the **hidden layer as a unit** which takes **some inputs** and returns **some outputs**



Classical Diagram: *Hidden Layer as Unit*


When we study a **gated architecture**: it is **common** to look at the **hidden layer as a unit** which takes **some inputs** and returns **some outputs**



We are mainly interested on **this block**: we want to know that given **last state** and **new input**

Classical Diagram: *Hidden Layer as Unit*

When we study a **gated architecture**: it is **common** to look at the **hidden layer as a unit** which takes **some inputs** and returns **some outputs**




We are mainly interested on **this block**: we want to know that given **last state** and **new input**

- 1 How does **this unit** **update hidden state**?
- 2 What components are passed to the **next time interval**?

Classical Diagram: *Hidden Layer as Unit*


When we study a **gated architecture**: it is **common** to look at the **hidden layer as a unit** which takes **some inputs** and returns **some outputs**




We are mainly interested on **this block**: we want to know that given **last state** and **new input**

- 1 How does **this unit** **update hidden state**?
- 2 What components are passed to the **next time interval**?
 - ↳ Here, we have only **$h[t]$**
 - ↳ But, we may have **other components**
 - ↳ We will see it in **LSTM**

Classical Diagram: *Basic RNN*



Classical Diagram: *Basic RNN*



Practical Gated Architectures: *Gated Recurrent Unit*

Gated Recurrent Unit (GRU)

Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- 1 Start with an initial *hidden state*
- 2 Compute *update gate* $\Gamma_u[t] = \sigma(\mathbf{W}_{u,\text{in}}\mathbf{x}[t] + \mathbf{W}_{u,\text{m}}\mathbf{h}[t-1])$
- 3 Compute *reset gate* $\Gamma_r[t] = \sigma(\mathbf{W}_{r,\text{in}}\mathbf{x}[t] + \mathbf{W}_{r,\text{m}}\mathbf{h}[t-1])$

Practical Gated Architectures: *Gated Recurrent Unit*

Gated Recurrent Unit (GRU)

Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- ① Start with an initial *hidden state*
- ② Compute *update gate* $\Gamma_u[t] = \sigma(\mathbf{W}_{u,\text{in}}\mathbf{x}[t] + \mathbf{W}_{u,\text{m}}\mathbf{h}[t-1])$
- ③ Compute *reset gate* $\Gamma_r[t] = \sigma(\mathbf{W}_{r,\text{in}}\mathbf{x}[t] + \mathbf{W}_{r,\text{m}}\mathbf{h}[t-1])$
- ④ Compute *actual memory* $\tilde{\mathbf{h}}[t] = f(\mathbf{W}_1\mathbf{x}[t] + \mathbf{W}_m\Gamma_r[t] \odot \mathbf{h}[t-1])$

Practical Gated Architectures: *Gated Recurrent Unit*

Gated Recurrent Unit (GRU)

Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- ① Start with an initial *hidden state*
- ② Compute *update gate* $\mathbf{\Gamma}_u[t] = \sigma(\mathbf{W}_{u,\text{in}}\mathbf{x}[t] + \mathbf{W}_{u,\text{m}}\mathbf{h}[t-1])$
- ③ Compute *reset gate* $\mathbf{\Gamma}_r[t] = \sigma(\mathbf{W}_{r,\text{in}}\mathbf{x}[t] + \mathbf{W}_{r,\text{m}}\mathbf{h}[t-1])$
- ④ Compute *actual memory* $\tilde{\mathbf{h}}[t] = f(\mathbf{W}_1\mathbf{x}[t] + \mathbf{W}_m\mathbf{\Gamma}_r[t] \odot \mathbf{h}[t-1])$
- ⑤ Update *hidden state* as $\mathbf{h}[t] = (1 - \mathbf{\Gamma}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{\Gamma}_u[t] \odot \tilde{\mathbf{h}}[t]$

Practical Gated Architectures: *Gated Recurrent Unit*

Gated Recurrent Unit (GRU)


Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- ① Start with an initial *hidden state*
- ② Compute *update gate* $\Gamma_u[t] = \sigma(\mathbf{W}_{u,\text{in}}\mathbf{x}[t] + \mathbf{W}_{u,\text{m}}\mathbf{h}[t-1])$
- ③ Compute *reset gate* $\Gamma_r[t] = \sigma(\mathbf{W}_{r,\text{in}}\mathbf{x}[t] + \mathbf{W}_{r,\text{m}}\mathbf{h}[t-1])$
- ④ Compute *actual memory* $\tilde{\mathbf{h}}[t] = f(\mathbf{W}_1\mathbf{x}[t] + \mathbf{W}_m\Gamma_r[t] \odot \mathbf{h}[t-1])$
- ⑤ Update *hidden state* as $\mathbf{h}[t] = (1 - \Gamma_u[t]) \odot \mathbf{h}[t-1] + \Gamma_u[t] \odot \tilde{\mathbf{h}}[t]$
- ⑥ Compute output $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2\mathbf{h}[t]) \rightsquigarrow f_{\text{out}}$ and f could be *different*

Or we could give $\mathbf{h}[t]$ to a *new layer*: for instance a *new GRU* whose *input is* $\mathbf{h}[t]$ and has its own *state*


Practical Gated Architectures: GRU

This is what's going on in a *GRU cell*




Practical Gated Architectures: GRU

Compute *update gate* $\Gamma_u[t] = \sigma(\mathbf{W}_{u,\text{in}}\mathbf{x}[t] + \mathbf{W}_{u,\text{m}}\mathbf{h}[t-1])$




Practical Gated Architectures: GRU

Compute *reset gate* $\Gamma_r[t] = \sigma(\mathbf{W}_{r,\text{in}}\mathbf{x}[t] + \mathbf{W}_{r,\text{m}}\mathbf{h}[t-1])$



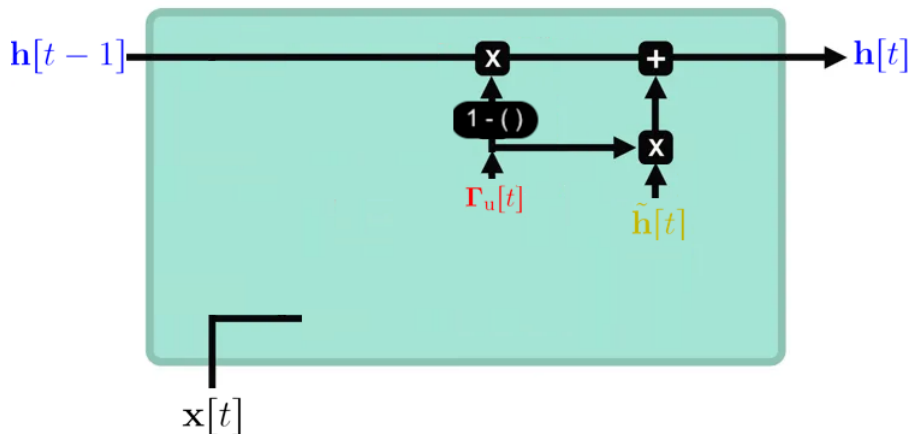
Practical Gated Architectures: GRU

Compute *actual memory* $\tilde{\mathbf{h}}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{\Gamma}_r[t] \odot \mathbf{h}[t - 1])$




Practical Gated Architectures: GRU

Update *hidden state* as $\mathbf{h}[t] = (1 - \mathbf{\Gamma}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{\Gamma}_u[t] \odot \tilde{\mathbf{h}}[t]$




GRU: Forward Pass

Starting from an *initial state*: GRU applies the first 5 steps *each time*




GRU: Forward Pass

Starting from an *initial state*: GRU applies the first 5 steps *each time*



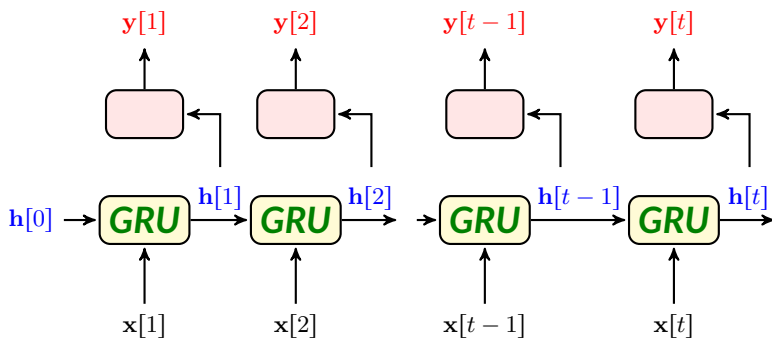
GRU: Forward Pass

Starting from an *initial state*: GRU applies the first 5 steps *each time*



GRU: Forward Pass

Starting from an *initial state*: GRU applies the first 5 steps *each time*



GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$:

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

- ① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

- ② We know that $\mathbf{h}[t] = (1 - \mathbf{r}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{r}_u[t] \odot \tilde{\mathbf{h}}[t]$

$$\nabla_{\mathbf{W}} \mathbf{h}[t] =$$

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

- ① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

- ② We know that $\mathbf{h}[t] = (1 - \mathbf{r}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{r}_u[t] \odot \tilde{\mathbf{h}}[t]$

$$\nabla_{\mathbf{W}} \mathbf{h}[t] = \nabla_{\mathbf{r}_u[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{r}_u[t]$$

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

② We know that $\mathbf{h}[t] = (1 - \mathbf{r}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{r}_u[t] \odot \tilde{\mathbf{h}}[t]$

$$\nabla_{\mathbf{W}} \mathbf{h}[t] = \nabla_{\mathbf{r}_u[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{r}_u[t] + \nabla_{\mathbf{h}[t-1]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t-1]$$

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

- ① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

- ② We know that $\mathbf{h}[t] = (1 - \mathbf{r}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{r}_u[t] \odot \tilde{\mathbf{h}}[t]$

$$\begin{aligned} \nabla_{\mathbf{W}} \mathbf{h}[t] = & \nabla_{\mathbf{r}_u[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{r}_u[t] + \nabla_{\mathbf{h}[t-1]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t-1] \\ & + \nabla_{\tilde{\mathbf{h}}[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \tilde{\mathbf{h}}[t] \end{aligned}$$

GRU: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside GRU*, e.g., $\mathbf{W}_{u,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

- ① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

- ② We know that $\mathbf{h}[t] = (1 - \mathbf{r}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{r}_u[t] \odot \tilde{\mathbf{h}}[t]$

$$\begin{aligned} \nabla_{\mathbf{W}} \mathbf{h}[t] = & \nabla_{\mathbf{r}_u[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{r}_u[t] + \nabla_{\mathbf{h}[t-1]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t-1] \\ & + \nabla_{\tilde{\mathbf{h}}[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}} \tilde{\mathbf{h}}[t] \end{aligned}$$

- ③ ...

Practical Gated Architectures: Long Short-Term Memory

Long Short-Term Memory (LSTM)

Say we set activation to $f(\cdot)$ \rightsquigarrow we usually set it to $\tanh(\cdot)$

- 1 Start with initial *hidden state* and *cell state*
- 2 Compute *forget gate* $\Gamma_f[t] = \sigma(\mathbf{W}_{f,\text{in}}\mathbf{x}[t] + \mathbf{W}_{f,\text{m}}\mathbf{h}[t-1])$
- 3 Compute *input gate* $\Gamma_i[t] = \sigma(\mathbf{W}_{i,\text{in}}\mathbf{x}[t] + \mathbf{W}_{i,\text{m}}\mathbf{h}[t-1])$
- 4 Compute *output gate* $\Gamma_o[t] = \sigma(\mathbf{W}_{o,\text{in}}\mathbf{x}[t] + \mathbf{W}_{o,\text{m}}\mathbf{h}[t-1])$

Practical Gated Architectures: Long Short-Term Memory

Long Short-Term Memory (LSTM)

Say we set activation to $f(\cdot)$ \rightsquigarrow we usually set it to $\tanh(\cdot)$

- ① Start with initial *hidden state* and *cell state*
- ② Compute *forget gate* $\Gamma_f[t] = \sigma(\mathbf{W}_{f,\text{in}}\mathbf{x}[t] + \mathbf{W}_{f,\text{m}}\mathbf{h}[t-1])$
- ③ Compute *input gate* $\Gamma_i[t] = \sigma(\mathbf{W}_{i,\text{in}}\mathbf{x}[t] + \mathbf{W}_{i,\text{m}}\mathbf{h}[t-1])$
- ④ Compute *output gate* $\Gamma_o[t] = \sigma(\mathbf{W}_{o,\text{in}}\mathbf{x}[t] + \mathbf{W}_{o,\text{m}}\mathbf{h}[t-1])$
- ⑤ Compute *actual cell state* $\tilde{\mathbf{c}}[t] = f(\mathbf{W}_1\mathbf{x}[t] + \mathbf{W}_m\mathbf{h}[t-1])$
- ⑥ Update *cell state* as $\mathbf{c}[t] = \Gamma_f[t]\mathbf{c}[t-1] + \Gamma_i[t] \odot \tilde{\mathbf{c}}[t]$

Practical Gated Architectures: Long Short-Term Memory

Long Short-Term Memory (LSTM)

Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- ① Start with initial *hidden state* and *cell state*
- ② Compute *forget gate* $\mathbf{\Gamma}_f[t] = \sigma(\mathbf{W}_{f,\text{in}}\mathbf{x}[t] + \mathbf{W}_{f,\text{m}}\mathbf{h}[t-1])$
- ③ Compute *input gate* $\mathbf{\Gamma}_i[t] = \sigma(\mathbf{W}_{i,\text{in}}\mathbf{x}[t] + \mathbf{W}_{i,\text{m}}\mathbf{h}[t-1])$
- ④ Compute *output gate* $\mathbf{\Gamma}_o[t] = \sigma(\mathbf{W}_{o,\text{in}}\mathbf{x}[t] + \mathbf{W}_{o,\text{m}}\mathbf{h}[t-1])$
- ⑤ Compute *actual cell state* $\tilde{\mathbf{c}}[t] = f(\mathbf{W}_1\mathbf{x}[t] + \mathbf{W}_m\mathbf{h}[t-1])$
- ⑥ Update *cell state* as $\mathbf{c}[t] = \mathbf{\Gamma}_f[t]\mathbf{c}[t-1] + \mathbf{\Gamma}_i[t] \odot \tilde{\mathbf{c}}[t]$
- ⑦ Update *hidden state* as $\mathbf{h}[t] = \mathbf{\Gamma}_o[t] \odot f(\mathbf{c}[t])$

Practical Gated Architectures: Long Short-Term Memory

Long Short-Term Memory (LSTM)


Say we set activation to $f(\cdot) \rightsquigarrow$ we usually set it to $\tanh(\cdot)$

- ① Start with initial *hidden state* and *cell state*
- ② Compute *forget gate* $\Gamma_f[t] = \sigma(\mathbf{W}_{f,\text{in}}\mathbf{x}[t] + \mathbf{W}_{f,\text{m}}\mathbf{h}[t-1])$
- ③ Compute *input gate* $\Gamma_i[t] = \sigma(\mathbf{W}_{i,\text{in}}\mathbf{x}[t] + \mathbf{W}_{i,\text{m}}\mathbf{h}[t-1])$
- ④ Compute *output gate* $\Gamma_o[t] = \sigma(\mathbf{W}_{o,\text{in}}\mathbf{x}[t] + \mathbf{W}_{o,\text{m}}\mathbf{h}[t-1])$
- ⑤ Compute *actual cell state* $\tilde{\mathbf{c}}[t] = f(\mathbf{W}_1\mathbf{x}[t] + \mathbf{W}_m\mathbf{h}[t-1])$
- ⑥ Update *cell state* as $\mathbf{c}[t] = \Gamma_f[t]\mathbf{c}[t-1] + \Gamma_i[t] \odot \tilde{\mathbf{c}}[t]$
- ⑦ Update *hidden state* as $\mathbf{h}[t] = \Gamma_o[t] \odot f(\mathbf{c}[t])$
- ⑧ Compute output $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2\mathbf{h}[t]) \rightsquigarrow f_{\text{out}}$ and f could be *different*

Or we could give $\mathbf{h}[t]$ to a *new layer*: for instance a *new LSTM* whose *input* is $\mathbf{h}[t]$ and has its own *hidden and cell states*


Practical Gated Architectures: LSTM

This is how inside an **LSTM unit** looks like




Practical Gated Architectures: *LSTM*

Actual cell state $\tilde{\mathbf{c}}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$




Practical Gated Architectures: *LSTM*

We use *forget gate* and *update gate* to update *cell state*



Practical Gated Architectures: LSTM


We use *output gate* to control fellow of memory to the *hidden state*



Practical Gated Architectures: LSTM

Intuitively, the gates in **LSTM** impact the *flow of information* as follows


- *Forget gate* controls how much we forget from *last state*
 - ↳ Assume $\Gamma_f[t] = 0$: then, we remember nothing of $c[t-1]$



Practical Gated Architectures: LSTM

Intuitively, the gates in **LSTM** impact the *flow of information* as follows


- **Forget gate** controls how much we forget from *last state*
 - ↳ Assume $\Gamma_f[t] = 0$: then, we remember nothing of $c[t-1]$
- **Input gate** controls how much we remember from *new cell state*
 - ↳ Assume $\Gamma_i[t] = 0$: then, we remember nothing of $\tilde{c}[t]$



Practical Gated Architectures: LSTM


Intuitively, the gates in LSTM impact the *flow of information* as follows

- **Output gate** controls how much we let from *updated state* to *go out*
 - ↳ Assume $\mathbf{\Gamma}_0[t] = \mathbf{0}$: then, we send nothing of $\mathbf{c}[t]$ out



LSTM: Forward Pass

Starting from *initial hidden and cell state*: LSTM passes forward as




Pay Attention

Note that unlike other architectures, LSTM does **not** keep all memory inside *hidden state* but it carries it also in *cell state*. *This state* is only for *memory* and is not directly used by higher layers, e.g., *output layer of the NN*

LSTM: Forward Pass

Starting from *initial hidden and cell state*: LSTM passes forward as




Pay Attention

Note that unlike other architectures, LSTM does **not** keep all memory inside *hidden state* but it carries it also in *cell state*. This state is only for *memory* and is not directly used by higher layers, e.g., *output layer of the NN*

LSTM: Forward Pass

Starting from *initial hidden and cell state*: LSTM passes forward as

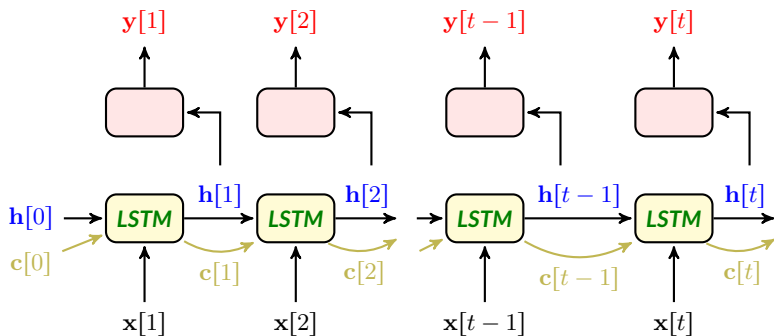


Pay Attention

Note that unlike other architectures, LSTM does **not** keep all memory inside *hidden state* but it carries it also in *cell state*. This state is only for *memory* and is not directly used by higher layers, e.g., *output layer of the NN*

LSTM: Forward Pass

Starting from *initial hidden and cell state*: LSTM passes forward as



Pay Attention

Note that unlike other architectures, LSTM does **not** keep all memory inside *hidden state* but it carries it also in *cell state*. This state is only for *memory* and is not directly used by higher layers, e.g., *output layer of the NN*

LSTM: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside LSTM*, e.g., $\mathbf{W}_{i,m}$:

LSTM: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside LSTM*, e.g., $\mathbf{W}_{i,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W} \mathbf{y}[t]}$$

LSTM: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside LSTM*, e.g., $\mathbf{W}_{i,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

LSTM: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside LSTM*, e.g., $\mathbf{W}_{i,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

② ...

LSTM: Backward Pass

Say we finished *forward pass* at time t . We now want to find $\nabla_{\mathbf{W}} \hat{R}$ for some \mathbf{W} that is *inside LSTM*, e.g., $\mathbf{W}_{i,m}$: we start *backpropagating* from $\nabla_{\mathbf{y}[t]} \hat{R}$

$$\nabla_{\mathbf{W}} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R} \circ \nabla_{\mathbf{W}} \mathbf{y}[t]$$

① We know $\mathbf{y}[t] = f_{\text{out}}(\mathbf{W}_2 \mathbf{h}[t])$

$$\nabla_{\mathbf{W}} \mathbf{y}[t] = \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}} \mathbf{h}[t]$$

② ...


Suggestion

Try writing it once to see the impact of gates!

Bidirectional RNNs

We have up to now considered *unidirectional* RNNs


we *start from beginning* of the sequence and *move in one direction*



Bidirectional RNNs

We have up to now considered *unidirectional* RNNs

we *start from beginning* of the sequence and *move in one direction*



But, can't we *learn from future* input as well?

Bidirectional RNNs

Future entries can have information about **past**: say our RNN wants to fill the empty field

... the color that many people assume is the color of **sun** ...

Obviously, **future** input in the sequence is helping in this example!

Bidirectional RNNs

Future entries can have information about **past**: say our RNN wants to fill the empty field

... the color that many people assume is the color of **sun** ...

Obviously, **future** input in the sequence is helping in this example!

- + But, how can we **get information** from future?
- Well, we have **the whole sequence**: we could move **once from left to right** and **once from right to left**


Bidirectional RNNs

Bidirectional RNNs


A *bidirectional* RNN (BRNN) consists of two RNNs

- one that starts with an *initial hidden state* at $t = 0$ and computes $\mathbf{h}[t]$ from $\mathbf{h}[t - 1]$ and $\mathbf{x}[t]$
- another that starts with an *initial hidden state* at $t = T + 1$ and computes $\mathbf{h}'[t]$ from $\mathbf{h}'[t + 1]$ and $\mathbf{x}[t]$


Output at time t is determined from *merged* version of the two hidden states



Bidirectional RNNs




Bidirectional RNNs




- + What exactly is this *merge* block?
- It gets the *two states* and returns a vector that matches *output layer*

Bidirectional GRU




- + Should we use *any RNN* here?
- Sure! We may use *GRU*

Bidirectional GRU




- + Should we use *any* RNN here?
- Sure! We may use GRU

Bidirectional LSTM




- + Should we use *any RNN* here?
- Sure! We may use *LSTM*

Bidirectional LSTM



- + Should we use *any RNN* here?
- Sure! We may use *LSTM*

RNNs in PyTorch



- + Any suggestion for *merging* the *hidden states*?
- Sure! Let's see some code

RNNs in PyTorch

In PyTorch, we can access a *basic RNN* in `torch.nn` module

```
torch.nn.RNN()
```

We can make it *deep* by simply choosing `num_layers` *more than one* and *bidirectional* by setting `bidirectional` to `True`.

RNNs in PyTorch

In PyTorch, we can access a *basic RNN* in `torch.nn` module

```
torch.nn.RNN()
```

We can make it *deep* by simply choosing `num_layers` *more than one* and *bidirectional* by setting `bidirectional` to `True`. Same with *GRU* and *LSTM*

```
torch.nn.LSTM()
```

```
torch.nn.GRU()
```

RNNs in PyTorch

In PyTorch, we can access a *basic RNN* in `torch.nn` module

```
torch.nn.RNN()
```

We can make it *deep* by simply choosing `num_layers` *more than one* and *bidirectional* by setting `bidirectional` to `True`. Same with *GRU* and *LSTM*

```
torch.nn.LSTM()
```

```
torch.nn.GRU()
```

In *bidirectional* case, we get access to *both states*. To *merge* them, we could

- add the two states
- average them
- concatenate them, i.e., $\mathbf{h}_c[t] = (\mathbf{h}[t], \mathbf{h}'[t])$

or do *any other operation* that we find useful

Computing Loss: *Challenge*

We mentioned several times in this chapter that we *assume*

we *can compute* the loss *between RNN's output sequence* and *label sequence*

However, it is in general a challenge!

Computing Loss: Challenge

We mentioned several times in this chapter that we *assume*

we *can compute* the loss *between RNN's output sequence* and *label sequence*

However, it is in general a challenge!

- + Why is it a challenge? We did it *easily* in *FNN and CNN* chapters!
- Because the *problem* there was *already properly segmented!*
- + What do you mean by *segmented*?
- Let's break it down!

Computing Loss: *Motivating Example*

Let's consider a simple example: *we have an image that includes a sequence of handwritten digits, e.g.,*

- *The sequence includes **five digits***
- *Each digit is **either 1, 2, 3, or 4***

Computing Loss: Motivating Example

Let's consider a simple example: we have an image that includes a sequence of handwritten digits, e.g.,

- The sequence includes **five digits**
- Each digit is **either 1, 2, 3, or 4**

Our task is to **recognize** this **sequence**, i.e., return the **five** digits in **correct order**

- This is a **classification** task
- How can we do it? We use NNs
 - ↳ We train an NN over **lots of images**: we have lots of images of **sequence of digits**
 - ↳ We then use it to **recognize** new images

23 241 \rightsquigarrow 23241

Let's say we are going to use a CNN

Computing Loss: Motivating Example

To use a CNN, we need to specify our input size

- We segment an input image into a sequence of **five images**
↳ These images are all as large as CNN's input size

23 241 \rightsquigarrow 2, 3, 2, 4, 1

Computing Loss: Motivating Example

To use a CNN, we need to specify our input size

- We segment an input image into a sequence of **five images**
↳ These images are all as large as CNN's input size

23 241 \rightsquigarrow 2, 3, 2, 4, 1

- We label each image with its **label**, e.g., 2 is **labeled as 2**

Computing Loss: Motivating Example

To use a CNN, we need to specify our input size

- We segment an input image into a sequence of **five images**
 - ↳ These images are all as large as CNN's input size

23 241 \rightsquigarrow 2, 3, 2, 4, 1

- We label each image with its **label**, e.g., 2 is **labeled as 2**
- We give these **five images** to our CNN and get **five outputs**
 - ↳ Assume we use **softmax** at the output layer
 - ↳ For each image, we get a **vector of size 4** as output
 - ↳ Each entry represents **probability** of **image** being one of digits **1, 2, 3, and 4**

Computing Loss: Motivating Example

To use a CNN, we need to specify our input size

- We segment an input image into a sequence of **five images**
 - ↳ These images are all as large as CNN's input size

23 241 \rightsquigarrow 2, 3, 2, 4, 1

- We label each image with its **label**, e.g., 2 is **labeled as 2**
- We give these **five images** to our CNN and get **five outputs**
 - ↳ Assume we use **softmax** at the output layer
 - ↳ For each image, we get a **vector of size 4** as output
 - ↳ Each entry represents **probability** of **image** being one of digits **1, 2, 3, and 4**

Computing Loss: Motivating Example

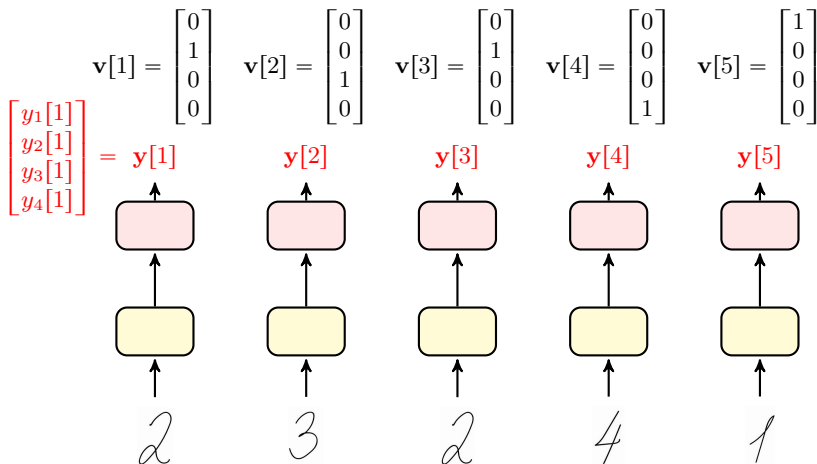
To use a CNN, we need to specify our input size

- We segment an input image into a sequence of **five images**
 - ↳ These images are all as large as CNN's input size

23 241 \rightsquigarrow 2, 3, 2, 4, 1

- We label each image with its **label**, e.g., 2 is **labeled as 2**
- We give these **five images** to our CNN and get **five outputs**
 - ↳ Assume we use **softmax** at the output layer
 - ↳ For each image, we get a **vector of size 4** as output
 - ↳ Each entry represents **probability** of **image** being one of digits **1, 2, 3, and 4**
- To compute loss, we compare each **output** with its **corresponding label**

Computing Loss: *Motivating Example*



↳ $y_j[t]$ is the *probability* of digit in *time* t being j

Computing Loss: *Motivating Example*

Here, we already have the data *segmented* into

a *sequence* that for each *time step* has a *label*

Computing Loss: Motivating Example

Here, we already have the data *segmented* into

a *sequence* that for each *time step* has a *label*

So, computing loss is easy as pie!

$$\begin{aligned}\hat{R} &= \mathcal{L}(\mathbf{y}[1], \dots, \mathbf{y}[5], \mathbf{v}[1], \dots, \mathbf{v}[5]) \\ &= \sum_{t=1}^5 \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) = \sum_{t=1}^5 \hat{R}[t]\end{aligned}$$

Computing Loss: Motivating Example

Here, we already have the data *segmented* into

a *sequence* that for each *time step* has a *label*

So, computing loss is easy as pie!

$$\begin{aligned}\hat{R} &= \mathcal{L}(\mathbf{y}[1], \dots, \mathbf{y}[5], \mathbf{v}[1], \dots, \mathbf{v}[5]) \\ &= \sum_{t=1}^5 \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) = \sum_{t=1}^5 \hat{R}[t]\end{aligned}$$

When we compute gradients, we note that *only* $\hat{R}[t]$ depends on $\mathbf{y}[t]$: so, for a given output at time $t = i$ we can simply write

$$\nabla_{\mathbf{y}[i]} \hat{R} = \sum_{t=1}^5 \nabla_{\mathbf{y}[i]} \hat{R}[t] = \nabla_{\mathbf{y}[i]} \hat{R}[i] = \nabla_{\mathbf{y}[i]} \mathcal{L}(\mathbf{y}[i], \mathbf{v}[i])$$

Computing Loss: One-to-One Correspondence

- + *But is it practical to do segmentation **by hand**?*
- **No!** This is why we built RNNs!

Computing Loss: One-to-One Correspondence

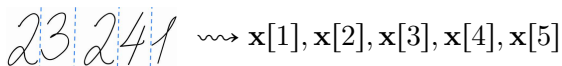
- + *But is it practical to do segmentation **by hand**?*
- **No!** This is why we built RNNs!

With RNNs, we address this learning task as bellow

- We look at the **complete image** as a sequence of data
 - ↳ We divide input into **multiple equal-size frames**
- We go over each frame separately
 - ↳ We give the frame as the input along with **previous state**
 - ↳ We compute a **new state** which can potentially give us the output

Computing Loss: One-to-One Correspondence

If we are extremely lucky; then, our segmentation looks like this


$$23241 \rightsquigarrow \mathbf{x}[1], \mathbf{x}[2], \mathbf{x}[3], \mathbf{x}[4], \mathbf{x}[5]$$


and we have a label for each time step

Computing Loss: One-to-One Correspondence

If we are extremely lucky; then, our segmentation looks like this

23241 \rightsquigarrow $x[1], x[2], x[3], x[4], x[5]$

and we have a label for each time step




Computing Loss: One-to-One Correspondence

But, that's too good to happen! Usually we have

23 241 $\rightsquigarrow \mathbf{x}[1], \mathbf{x}[2], \dots, \mathbf{x}[T]$


and we have a label in **some** time steps



Computing Loss: One-to-One Correspondence

In this typical case, two questions seem non-trivial


- 1 Where should we **put** each **label**? \equiv Where should we **read** each **label**?



Computing Loss: One-to-One Correspondence

In this typical case, two questions seem non-trivial

- 1 Where should we **put** each **label**? \equiv Where should we **read** each **label**?
- 2 What should we do with **non-labeled outputs**, e.g., $y[1]$?



Computing Loss: *One-to-One Correspondence*

The key challenge in computing the loss is that we do **not** have necessarily *one-to-one correspondence* with *sequence data*

Computing Loss: One-to-One Correspondence

The key challenge in computing the loss is that we do **not** have necessarily *one-to-one correspondence* with *sequence data*

Correspondence Problem

With sequence data, we could have a data-sequence of time T that is labeled by a sequence of size $K < T$ where

***no time index** is specified for any label in the K -long label sequence*

Computing Loss: One-to-One Correspondence

The key challenge in computing the loss is that we do **not** have necessarily *one-to-one correspondence* with *sequence data*

Correspondence Problem

With *sequence data*, we could have a data-sequence of time T that is labeled by a sequence of size $K < T$ where

no time index is specified for any label in the K -long label sequence

Correspondence problem exists pretty much in *all practical* sequence data

- In speech recognition, *multiple time frames* correspond to a *single word*
- In text recognition, *multiple image frames* correspond to a *single letter*
- ...

Correspondence Problem: *Formulation*

Let's formulate the problem clearly: *Say we have*

A sequence of data

$$\mathbf{x}[1 : T] = \mathbf{x}[1], \dots, \mathbf{x}[T]$$

*that is **labeled** with the sequence of K true labels*

$$\mathbf{v}[1 : K] = \mathbf{v}[1], \dots, \mathbf{v}[K]$$

*where K and T can be **different***

Correspondence Problem: *Formulation*

Let's formulate the problem clearly: *Say we have*

A sequence of data

$$\mathbf{x}[1 : T] = \mathbf{x}[1], \dots, \mathbf{x}[T]$$

*that is **labeled** with the sequence of K true labels*


$$\mathbf{v}[1 : K] = \mathbf{v}[1], \dots, \mathbf{v}[K]$$

*where K and T can be **different***

*For this setting, we want to train an RNN with this data sequence: starting with an initial state, this RNN returns an **output** sequence*

$$\mathbf{y}[1 : T] = \mathbf{y}[1], \dots, \mathbf{y}[T]$$

Correspondence Problem: *Formulation*




To be able to train this RNN, we need to

- 1 define a loss function that computes $\hat{R} = \mathcal{L}(\mathbf{y}[1:T], \mathbf{v}[1:K])$

↳ We need this loss function to be differentiable with respect to all outputs

$$\nabla_{\mathbf{y}[1]} \hat{R}, \dots, \nabla_{\mathbf{y}[T]} \hat{R}$$

Correspondence Problem: *Formulation*



To use this RNN after training, i.e., for inferring, we need to

- ② know how to map **outputs** to **predicted labels**

↳ We need to extract K labels from $y[1 : T]$, i.e.,

$$y[1], \dots, y[T] \mapsto \hat{v}[1], \dots, \hat{v}[K]$$


Let's look into different settings

Setting I: *Perfectly Segmented*

In some problems, we have our data *perfectly segmented*

- There is a separate label for each time step, i.e., $K = T$

↳ *many-to-many type I and one-to-many*




Setting I: Perfectly Segmented

In some problems, we have our data perfectly segmented

- There is a separate label for each time step, i.e., $K = T$

↳ many-to-many type I and one-to-many



Attention

We can always treat a non-existing input entry as an empty

↳ We are good as long as we have a label at each time t

Setting I: *Defining Loss*

In such settings, we define the loss to be *aggregated loss* over time

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

for some loss function $\mathcal{L}(\cdot, \cdot)$

Setting I: *Defining Loss*

In such settings, we define the loss to be *aggregated loss* over time

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

for some loss function $\mathcal{L}(\cdot, \cdot)$

The gradients are then trivially computed

Gradient with respect to particular output $\mathbf{y}[t]$ is

$$\nabla_{\mathbf{y}[t]} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

Setting I: Inference

Inference in such setting is performed *by one-to-one mapping: at time t , we predict based on $\mathbf{y}[t]$*

$$\mathbf{y}[1] \mapsto \hat{\mathbf{v}}[1], \dots, \mathbf{y}[T] \mapsto \hat{\mathbf{v}}[T]$$

Setting I: Inference

Inference in such setting is performed by *one-to-one mapping*: at time t , we predict based on $\mathbf{y}[t]$

$$\mathbf{y}[1] \mapsto \hat{\mathbf{v}}[1], \dots, \mathbf{y}[T] \mapsto \hat{\mathbf{v}}[T]$$

For instance, assume $\mathbf{y}[t]$ is output of a softmax activation; then, we set

$$\hat{\mathbf{v}}[t] = \operatorname{argmax} \mathbf{y}[t]$$

where argmax returns the index of the largest entry, e.g.,

$$\operatorname{argmax} \begin{bmatrix} 0.1 \\ 0.7 \\ 0.2 \\ 0 \end{bmatrix} = 2$$

Setting II: Known Segments


In some problems, we have only one label for the whole sequence, i.e., $K = 1$

↳ It corresponds to **many-to-one** type of problems

↳ This can be that we have really only one label, e.g., content classification

↳ It can be that we know the time index t at which each label is assigned

↳ We split data to sub-sequences with each sub-sequence having only one label



Setting II: *Defining Loss for Dumb NN*

A **naive** approach to define loss is to set it be the loss between **last output** and **label**, i.e.,

$$\hat{R} = \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1])$$

Setting II: *Defining Loss for Dumb NN*

A **naive** approach to define loss is to set it be the loss between **last output** and **label**, i.e.,

$$\hat{R} = \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1])$$

With this loss, *gradient with respect to particular output* $\mathbf{y}[t]$ is

$$\nabla_{\mathbf{y}[t]} \hat{R} = \begin{cases} \nabla_{\mathbf{y}[T]} \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1]) & t = T \\ 0 & t \neq T \end{cases}$$

Setting II: Defining Loss for Dumb NN

A **naive** approach to define loss is to set it be the loss between **last output** and **label**, i.e.,

$$\hat{R} = \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1])$$


With this loss, *gradient with respect to particular output* $\mathbf{y}[t]$ is

$$\nabla_{\mathbf{y}[t]} \hat{R} = \begin{cases} \nabla_{\mathbf{y}[T]} \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1]) & t = T \\ 0 & t \neq T \end{cases}$$

- + But does it make sense to **ignore all other outputs**?
- Not at all! We are training a **dumb** NN that can respond only when **it's over with the whole sequence**!

Setting II: Loss for Smarter Training

An *extremely smart* NN is the one who knows the label *before the input speaks!*




For this NN, the loss is

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

Setting II: Loss for Smarter Training

An *extremely smart* NN is the one who knows the label *before the input speaks!*



For this NN, the loss is

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

But, we should be *careful!* We should not expect NN to know everything from potentially *irrelevant* input!

Setting II: Defining Proper Loss

A **realistic** approach is to define the loss via a **weighted sum**, i.e.,

$$\hat{R} = \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

where w_t is the weight at time t

- initially w_t is **small**
 - ↳ we do not expect the NN to know everything from very beginning
- it gradually **increases** up to its maximum w_T
 - ↳ by time T the NN should know the label

Setting II: Defining Proper Loss

A **realistic** approach is to define the loss via a **weighted sum**, i.e.,

$$\hat{R} = \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

where w_t is the weight at time t

- initially w_t is **small**
 - ↳ we do not expect the NN to know everything from very beginning
- it gradually **increases** up to its maximum w_T
 - ↳ by time T the NN should know the label

With this loss, gradient with respect to particular output $\mathbf{y}[t]$ is

$$\nabla_{\mathbf{y}[t]} \hat{R} = w_t \nabla_{\mathbf{y}[t]} \hat{R}[t] = w_t \nabla_{\mathbf{y}[t]} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

Setting II: *Inference*

Inference in such setting is performed *by many-to-one mapping*: we only predict based on $\mathbf{y}[1 : T]$

$$\mathbf{y}[1 : T] \mapsto \hat{\mathbf{v}}[1]$$

Setting II: Inference

Inference in such setting is performed by *many-to-one mapping*: we only predict based on $\mathbf{y}[1 : T]$

$$\mathbf{y}[1 : T] \mapsto \hat{\mathbf{v}}[1]$$

For instance, assume $\mathbf{y}[t]$ is output of a softmax activation; then, we set


$$\hat{\mathbf{v}}[t] = \operatorname{argmax} \mathbf{y}[t]$$

and then take a (potentially *weighted*) majority vote: $\hat{\mathbf{v}}[1]$ is the class that *most often estimated* with occurrence at each time being *weighted* by some *weight*

Setting III: Unknown Segments

Most common case is that we have a label sequence *shorter* than our *data*


- ↳ Each label in this sequence is corresponding to a *segment of input*
 - ↳ We *do not know* where this *segment* begins and where it ends
 - ↳ There might be even no clear answer to that!



Setting III: Unknown Segments

Most common case is that we have a label sequence *shorter* than our *data*

- ↳ Each label in this sequence is corresponding to a *segment of input*
 - ↳ We *do not know* where this *segment* begins and where it ends
 - ↳ There might be even no clear answer to that!




Note that we are dealing with a *sequence to sequence* model: we want to learn relation between *sequence* $x[1 : T]$ and *sequence* $v[1 : K]$!

Setting II: Example

Assume we have image *121* that is divided into a sequence of *five pixel vectors*


- Since it is a training data, it is *labeled as 121*
 - ↳ We *do not know* after which output we should expect RNN to know *first, second* or *third digit*!



Setting II: Example

Assume we have image *121* that is divided into a sequence of *five pixel vectors*


- Since it is a training data, it is *labeled as 121*
 - ↳ We *do not know* after which output we should expect RNN to know *first, second or third digit!*



Setting II: Example

Assume we have image *121* that is divided into a sequence of *five pixel vectors*

- Since it is a training data, it is *labeled as 121*
 - ↳ We *do not know* after which output we should expect RNN to know *first, second* or *third digit*!




Setting II: Example

Assume we have image *121* that is divided into a sequence of *five pixel vectors*

- Since it is a training data, it is *labeled as 121*
 - ↳ We *do not know* after which output we should expect RNN to know *first*, *second* or *third digit*!


label



- + Sounds *impossible*!
- Only *impossible is impossible*! Let's carry on and see what we can do!


Setting II: Genie-Defined Loss

Assume a genie has told us end of each **segment**



Setting II: Genie-Defined Loss

Assume a genie has told us end of each **segment**




We can fill the **empty labels** with **repetition**, and then define the loss as

$$\hat{R} = \sum_{k=1}^K \sum_{t=i_{k-1}+1}^{i_k} w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[k])$$


where i_k is where label \mathbf{v}_k ends, e.g., in above diagram $i_1 = 2$

Setting II: Defining Loss



We don't have the *genie*:

Setting II: Defining Loss



We don't have the *genie*: we could assume that i_k is something to *learn*!

$$\hat{R}(\mathbf{i}) = \sum_{k=1}^K \sum_{t=i_{k-1}+1}^{i_k} w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[k])$$

where $\mathbf{i} = [0, i_1, \dots, i_K]$ is something we need to learn

Setting II: Optimal Segmentation

- + How could we learn \mathbf{i} ? Should we compute also $\nabla_{\mathbf{i}} \hat{R}$?
- Well! You may try! But, obviously i_k is an *integer*!

Setting II: Optimal Segmentation

- + How could we learn \mathbf{i} ? Should we compute also $\nabla_{\mathbf{i}} \hat{R}$?
- Well! You may try! But, obviously i_k is an *integer*!

Optimal Segmentation

Optimal approach for finding \mathbf{i} is to *train* the NN for *all possible choice for \mathbf{i}* and then find the final training loss $\hat{R}(\mathbf{i})$. The *optimal segmentation* is then given by

$$\mathbf{i}^* = \underset{\mathbf{i}}{\operatorname{argmin}} \hat{R}(\mathbf{i})$$

Setting II: Optimal Segmentation

- + How could we learn \mathbf{i} ? Should we compute also $\nabla_{\mathbf{i}} \hat{R}$?
- Well! You may try! But, obviously i_k is an *integer*!

Optimal Segmentation

Optimal approach for finding \mathbf{i} is to *train* the NN for *all possible choice for \mathbf{i}* and then find the final training loss $\hat{R}(\mathbf{i})$. The *optimal segmentation* is then given by

$$\mathbf{i}^* = \underset{\mathbf{i}}{\operatorname{argmin}} \hat{R}(\mathbf{i})$$

- + Is it *computationally feasible*?
- No! The number of *possible choice for \mathbf{i}* grows *exponentially with T* ! We need to go for sub-optimal approaches

Setting II: Number of Possible Segmentations

- + How is it **exponentially large**?
- Let's look at our example

In our last example, we should assign **label sequence 121** to a **sequence of length 5**: each entry of **output sequence** in this case can be labeled by **1 (the first one)**, **2** or **1 (the last one)**. This means that we have **3 choices** of label for each time interval; thus, the total number of possible segmentations is around **3^5** .

In general number of segmentations grows **exponentially** with **T**


- + But wait a moment! We have also counted the case of **labeling all outputs with 1!** This cannot be the case!
- This is right! It is in general **much less than 3^5** but it's still **exponential**

Let's see the exact possible segmentations!

Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the *label sequence is 121*



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the *label sequence is 121*
 - First output is definitely in the *first segment*: it's *label* is definitely 1



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label

- We know that the *label sequence is 121*
 - First output is definitely in the *first segment*: it's *label* is definitely 1
 - Second output could be *still in the first segment* or in the *second segment*



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label

- We know that the *label sequence is 121*
 - First output is definitely in the *first segment*: it's *label* is definitely 1
 - Second output could be *still in the first segment* or in the *second segment*
 - Third output could be in the *first*, *second*, or *third segment*

| | $y[1]$ | $y[2]$ | $y[3]$ | $y[4]$ | $y[5]$ |
|-------|--------|--------|--------|--------|--------|
| label | 1 | 1 | 1 | | |
| | 2 | 2 | 2 | | |
| | 1 | 1 | 1 | | |

Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label

- We know that the *label sequence is 121*
 - First output is definitely in the *first segment*: it's *label* is definitely 1
 - Second output could be *still in the first segment* or in the *second segment*
 - Third output could be in the *first*, *second*, or *third segment*
 - Our labels should finish by the end of output sequence: fourth output *cannot* be in *first segment*

| | $y[1]$ | $y[2]$ | $y[3]$ | $y[4]$ | $y[5]$ |
|-------|--------|--------|--------|--------|--------|
| label | 1 | 1 | 1 | 1 | |
| | 2 | 2 | 2 | 2 | |
| | 1 | 1 | 1 | 1 | |

Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**

| | $y[1]$ | $y[2]$ | $y[3]$ | $y[4]$ | $y[5]$ |
|-------|--------|--------|--------|--------|--------|
| label | 1 | 1 | 1 | 1 | 1 |
| | 2 | 2 | 2 | 2 | 2 |
| | 1 | 1 | 1 | 1 | 1 |

Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: its **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**



Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label


- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**




Setting II: Number of Possible Segmentations

We intend to compare each of $y[1], \dots, y[5]$ with a label

- We know that the **label sequence is 121**
 - First output is definitely in the **first segment**: it's **label** is definitely **1**
 - Second output could be **still in the first segment** or in the **second segment**
 - Third output could be in the **first**, **second**, or **third segment**
 - Our labels should finish by the end of output sequence: fourth output **cannot** be in **first segment**
 - Last output could be only in the **third segment**




Setting II: Showing Segmentations on Graph



Though it's **exponentially** large: we see that each segmentation corresponds to one path on this graph


Setting II: Showing Segmentations on Graph



Though it's **exponentially** large: we see that each segmentation corresponds to one path on this graph


Blue path corresponds to $i_1 = 3$, $i_2 = 4$, and $i_3 = 5$, i.e., $\mathbf{i} = [0, 3, 4, 5]$

Setting II: Loss on Segmentation Graph



We can *compute the loss* for each segmentation *directly on this graph*: let's say that we have L *different paths* on the graph. For *each path*, we can write an *expanded* label sequence, e.g.,


Setting II: Loss on Segmentation Graph



We can **compute the loss** for each segmentation **directly on this graph**: let's say that we have L **different paths** on the graph. For **each path**, we can write an **expanded** label sequence, e.g.,

Expanded label sequence of blue path is $\{1, 1, 1, 2, 1\}$

Setting II: Loss on Segmentation Graph




We can *compute the loss* for each segmentation *directly on this graph*: let's say that we have L *different paths* on the graph. For *each path*, we can write an *expanded* label sequence, e.g.,

Expanded label sequence of *blue path* is $\{1, 1, 1, 2, 1\}$

This sequence is of length T and we show it for *path* ℓ with $\tilde{\mathbf{v}}_{\ell}[t]$


Setting II: Loss on Segmentation Graph



For each *path* $\ell = 1, \dots, L$, the loss is computed by aggregating the losses between outputs and *extended labels*

$$\hat{R}_{\ell} = \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_{\ell}[t]) =$$

Setting II: Loss on Segmentation Graph



For each *path* $\ell = 1, \dots, L$, the loss is computed by aggregating the losses between outputs and *extended labels*

$$\hat{R}_{\ell} = \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_{\ell}[t]) = \sum_{t=1}^T \hat{R}_{\ell}[t]$$

It again decomposes into sum of T terms with only *one* being function of $\mathbf{y}[t]$

Setting II: Optimal Segmentation on Graph

We can represent the optimal segmentation on the graph as below

OptimalSegmentTraining():

- 1: Initiate with $\hat{R} = +\infty$ and some random $\ell^* = \emptyset$
- 2: **for** $\ell = 1, \dots, L$ **do**
- 3: Let the loss be \hat{R}_ℓ
- 4: Train for **sufficient epochs**
- 5: **if** After training $\hat{R}_\ell < \hat{R}$ **then**
- 6: $\hat{R} \leftarrow \hat{R}_\ell$ and $\ell^* \leftarrow \ell$
- 7: **end if**
- 8: **end for**
- 9: Return **learnable parameters** and ℓ^*


- + Say we could be over with this **infeasible** training! How do we use the trained RNN for inference?
- In this case, we have ℓ^* which gives us optimal segmentation: we infer label of each segment based on its **corresponding outputs**

Setting II: Maximum-Likelihood Segmentation

Since optimal segmentation is **infeasible**, people use **maximum-likelihood approach** that is well-known in detection and coding theory

Maximum-Likelihood Segmentation

Start with an **initial guess** for **optimal path** on segmentation graph and do one step of training; then, **improve the guess** based on the outputs of next forward pass and go for **next step of training**



Let's look at its pseudo-code

Setting II: Maximum-Likelihood Segmentation

MaxLikelihoodTraining():


- 1: **for** Iteration $i = 1, \dots, I$ **do**
- 2: Pass forward through time: Compute **output sequence** $\mathbf{y}[1 : T]$
- 3: Compute $p(\tilde{\mathbf{v}}_{\ell}[1 : T]|\ell)$ for each **path** ℓ on segmentation graph
- 4: Update $\ell^* = \operatorname{argmax}_{\ell} p(\tilde{\mathbf{v}}_{\ell}[1 : T]|\ell)$
- 5: Set loss to \hat{R}_{ℓ^*} and backpropagate over RNN
- 6: Update **learnable parameters**
- 7: **end for**
- 8: Return **learnable parameters** and ℓ^*

⊗

⊗

- + Why we call it **maximum likelihood**?
- Because we guess **path** by **maximizing** the **likelihood** $p(\tilde{\mathbf{v}}_{\ell}[1 : T]|\ell)$
- + But how can find **likelihood** of a **path**?
- We can use **output sequence** $\mathbf{y}[1 : T]$


Setting II: Finding Likelihood on Segmentation Graph



Assume that each label could be 1, 2, 3, or 4: at each time t the RNN returns a 4-dimensional vector whose entries are probability of each class

we can *multiply the probabilities* of classes on *the path*


Setting II: Finding Likelihood on Segmentation Graph



For instance, the *yellow path* has a likelihood

$$p(\tilde{\mathbf{v}}_\ell[1:T]|\ell) = \prod_{t=1}^T p(\tilde{\mathbf{v}}_\ell[t]|\ell) = y_1[1]y_1[2]y_1[3]y_2[4]y_1[5]$$

Setting II: Finding Likelihood on Segmentation Graph



Or better to say: we just put **output entries** in **graph** and move on the **path**

$$p(\tilde{\mathbf{v}}_\ell[1:T]|\ell) = \prod_{t=1}^T y_{\tilde{v}_\ell[t]}[t] = y_1[1]y_1[2]y_1[3]y_2[4]y_1[5]$$

Setting II: Maximum-Likelihood Segmentation

- + OK! We can find the *likelihood*, but how can we *maximize it*? It's again an *exponentially* large search!

$$\ell^* = \operatorname{argmax}_{\ell} p(\tilde{\mathbf{v}}_{\ell}[1:T]|\ell)$$

- Well! If we only need the *maximum*, it turns not to be *exponential*

Setting II: Maximum-Likelihood Segmentation

- + OK! We can find the *likelihood*, but how can we *maximize it*? It's again an *exponentially* large search!

$$\ell^* = \underset{\ell}{\operatorname{argmax}} p(\tilde{\mathbf{v}}_{\ell}[1:T]|\ell)$$

- Well! If we only need the *maximum*, it turns not to be *exponential*

We can readily show that finding maximum likelihood on the graph is a *dynamic programming* problem and can be solved by the *Viterbi algorithm*

Maximum likelihood training can be *implemented efficiently*

Setting II: Maximum-Likelihood Inference

MaxLikelihoodTraining():

- 1: **for** Iteration $i = 1, \dots, I$ **do**
- 2: Pass forward through time: Compute **output sequence** $\mathbf{y}[1 : T]$
- 3: Compute $p(\tilde{\mathbf{v}}_\ell[1 : T]|\ell)$ for each **path** ℓ on segmentation graph
- 4: Update $\ell^* = \operatorname{argmax}_\ell p(\tilde{\mathbf{v}}_\ell[1 : T]|\ell)$
- 5: Set loss to \hat{R}_{ℓ^*} and backpropagate over RNN
- 6: Update **learnable parameters**
- 7: **end for**
- 8: Return **learnable parameters** and ℓ^*



- + How can we use our RNN for inference after training via **maximum likelihood** segmentation?
- We have access to ℓ^* : we **predict the label** of each segment based on its **corresponding** outputs

Setting II: Connectionist Temporal Classification

It turns out that *maximum-likelihood* could stick to a bad local minimum, i.e., it quickly converges to a *path* ℓ^* that is *much different* from ℓ^*

- + Is there any solution to this?
- Yes! We can use *connectionist temporal classification (CTC) loss*

CTC Loss

Instead of searching for a best segmentation and then minimizing its loss, we learn directly from *unsegmented data* by minimizing the average loss over *all possible segmentations*, i.e., we define loss to be

$$\hat{R} = \mathbb{E}_{\ell} \left\{ \hat{R}_{\ell} \right\} = \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_{\ell}[1:T]) \hat{R}_{\ell}$$

and train the RNN by finding *learnable parameters* that minimize this loss

Setting II: CTC Loss

- + But, why should it be a *better choice* of loss?
- Because we are *sure* that *optimal segmentation* is *contributing* to our loss

$$\hat{R} = \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_\ell[1:T]) \hat{R}_\ell = p(\ell^* | \tilde{\mathbf{v}}_\ell[1:T]) \hat{R}_{\ell^*} + \sum_{\ell \neq \ell^*} p(\ell | \tilde{\mathbf{v}}_\ell[1:T]) \hat{R}_\ell$$

- + Agreed! Now, how should we determine $p(\ell^* | \tilde{\mathbf{v}}_\ell[1:T])$?
- Just use the Bayes rule!
- + What about the expectation? It is at the end sum of *exponentially* large number of terms!
- We can again go on the *graph* and determine it via *dynamic programming*

Setting II: CTC Loss

The CTC loss can be written as

$$\begin{aligned}
 \hat{R} &= \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T]) \hat{R}_{\ell} = \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T]) \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_{\ell}[t]) \\
 &= \sum_{t=1}^T w_t \underbrace{\sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T]) \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_{\ell}[t])}_{\check{R}[t]} = \sum_{t=1}^T w_t \check{R}[t]
 \end{aligned}$$

This has been shown that $\check{R}[t]$ can be *recursively computed*⁶:

by some *approximation* we are able to readily compute $\nabla_{\mathbf{y}[t]} \check{R}[t]$

and we set $\nabla_{\mathbf{y}[t']} \check{R}[t] \approx \mathbf{0}$ for $t' \neq t$

⁶Check out the [original paper](#)

Setting II: Training with CTC Loss

```
CTC_Training():  
1: for iteration  $i = 1, \dots, I$  do  
2:   Pass forward through time: Compute output sequence  $\mathbf{y}[1 : T]$   
3:   Compute CTC loss  $\hat{R}$  and  $\nabla_{\mathbf{y}[t]} \hat{R}$  by recursion  
4:   Backpropagate through time and update learnable parameters  
5: end for  
6: Return learnable parameters
```

This looks like *standard training loop* now


the *loss* is only replaced with *CTC loss*

- + What about inference?
- Well! We should figure it out, since the training loop *does not* compute any *segmentation path*!

Setting II: Inference with CTC-Trained RNN

Let's get back to our **simple example**: assume that after training with **CTC loss** we give an image of handwritten *121* to the **RNN**


- RNN divides it into **5 frames** and is able to track optimal **segmentation**
 - ↳ The first three frames belong to the first segment
 - ↳ The remaining frames belong to the second and third segments



Setting II: Inference with CTC-Trained RNN

Let's get back to our **simple example**: assume that after training with **CTC loss** we give an image of handwritten *121* to the **RNN**

- RNN divides it into **5 frames** and is able to track optimal **segmentation**
 - The first three frames belong to the first segment
 - The remaining frames belong to the second and third segments
- RNN infers from output sequence $\hat{v}[1 : 5]$ but does **not** return **optimal path**



Setting II: Inference with CTC-Trained RNN

We can conclude from $\hat{v}[1 : 5]$ that the sequence is $\{1,2,1\}$ if we are sure the sequence *has no repetition*

Label Encoding and Decoding

CTC uses this fact and constructs following *encoding* and *decoding* method: it introduces a *new label* called “*blank:-*” which *does not belong* to *set of classes*

Setting II: Inference with CTC-Trained RNN

We can conclude from $\hat{v}[1 : 5]$ that the sequence is $\{1,2,1\}$ if we are sure the sequence *has no repetition*

Label Encoding and Decoding

CTC uses this fact and constructs following *encoding* and *decoding* method: it introduces a *new label* called “*blank:-*” which *does not belong* to *set of classes*

- While *training*, it adds *blank* between any two *repetitions*
 - For instance, we encode $112 \mapsto 1-12$, or $111 \mapsto 1-1-1$

Setting II: Inference with CTC-Trained RNN

We can conclude from $\hat{v}[1 : 5]$ that the sequence is $\{1,2,1\}$ if we are sure the sequence *has no repetition*

Label Encoding and Decoding

CTC uses this fact and constructs following *encoding* and *decoding* method: it introduces a *new label* called “*blank:-*” which *does not belong* to *set of classes*

- While *training*, it adds *blank* between any two *repetitions*
 - ↳ For instance, we encode $112 \mapsto 1-12$, or $111 \mapsto 1-1-1$
- For *inference*, it *removes any repetition* in *inferred sequence* $\hat{v}[1 : T]$ and *then drops blanks*
 - ↳ For instance, we decode $1-11-312 \mapsto 11312$, or $3333-3121 \mapsto 33121$

Setting II: *Training and Inference with CTC*

CTC_Training():

- 1: **for** iteration $i = 1, \dots, I$ **do**
- 2: Add **blanks** to the label sequences with repetition
- 3: Pass forward through time: Compute **output sequence** $\mathbf{y}[1 : T]$
- 4: Compute CTC loss \hat{R} and $\nabla_{\mathbf{y}[t]} \hat{R}$ by **recursion**
- 5: Backpropagate through time and update **learnable parameters**
- 6: **end for**
- 7: Return **learnable parameters**

Setting II: Training and Inference with CTC

CTC_Training():

- 1: **for** iteration $i = 1, \dots, I$ **do**
- 2: Add **blanks** to the label sequences with repetition
- 3: Pass forward through time: Compute **output sequence** $\mathbf{y}[1 : T]$
- 4: Compute CTC loss \hat{R} and $\nabla_{\mathbf{y}[t]} \hat{R}$ by **recursion**
- 5: Backpropagate through time and update **learnable parameters**
- 6: **end for**
- 7: Return **learnable parameters**

CTC_Inference():

- 1: Pass forward through time the input and compute **output** $\mathbf{y}[1 : T]$
- 2: Infer encoded sequence $\hat{\mathbf{v}}[1 : T]$ from $\mathbf{y}[1 : T]$
- 3: Remove repetitions from $\hat{\mathbf{v}}[1 : T] \mapsto \hat{\mathbf{v}}[1 : T']$
- 4: Remove **blanks** from $\hat{\mathbf{v}}[1 : T'] \mapsto \hat{\mathbf{v}}[1 : K]$
- 5: Return $\hat{\mathbf{v}}[1 : K]$

In PyTorch: CTC Loss

We can access CTC loss in `torch.nn` module as

```
torch.nn.CTCLoss()
```

Few notes about CTC loss implementation

- We need to specify the index of **blank** label
 - ↳ It should be **out of our set of classes**
 - ↳ By default, it is set to `blank = 0`
- When we define our model, we should always take **blank** label into account
 - ↳ If we do classification with C **classes**, model should return $C + 1$ **classes** with blank being one of them
- PyTorch considers cross-entropy loss function, i.e., $\mathcal{L}(\mathbf{y}, \tilde{\mathbf{v}}) = \text{CE}(\mathbf{y}, \tilde{\mathbf{v}})$
- As input to CTC loss: \mathbf{y} should be **logarithm of probabilities**
 - ↳ We can activate the output layer with **logarithmic softmax**